

# 데이터사이언스응용 (Capstone design)

김응희

[ehkim@sunmoon.ac.kr](mailto:ehkim@sunmoon.ac.kr)

Week 02


# Working with Real Data

- The best way to learn about machine learning
  - Experiment with real-world data

Category	Repository
Popular open data repositories	UCI Irvine Machine Learning Repository <a href="http://archive.ics.uci.edu/ml">http://archive.ics.uci.edu/ml</a>
	Kaggle datasets <a href="https://www.kaggle.com/datasets">https://www.kaggle.com/datasets</a>
	Amazon's AWS datasets <a href="https://registry.opendata.aws/">https://registry.opendata.aws/</a>
	공공데이터포털 <a href="https://www.data.go.kr/">https://www.data.go.kr/</a>
Meta portals	<a href="http://dataportals.org">http://dataportals.org</a>
	<a href="http://opendatamonitor.eu">http://opendatamonitor.eu</a>
	<a href="http://quandl.com">http://quandl.com</a>
Others	<a href="https://homl.info/9">https://homl.info/9</a>
	<a href="https://reddit.com/r/datasets">https://reddit.com/r/datasets</a>

# Kaggle dataset

<https://www.kaggle.com/datasets>



## Datasets


DocumentationNew Dataset

Search 20,809 datasets

FeedbackFilter

PUBLIC

Sort by: Hottest




New York City Airbnb Open Data

Dgomonov

17 days2 MB10.02 Files (other, CSV)

44




India - Trade Data

Lakshya Agarwal

13 days1 MB10.02 Files (CSV)

47




Crimes in Boston

Analyze Boston

a year10 MB8.22 Files (CSV)

306



PGA Tour Golf Data

Brad Klassen

9 days92 MB10.02 Files (CSV)

189

3

# 공공데이터포털

<https://www.data.go.kr/>

**DATA** 공공데이터포털  
. GO . KR

데이터찾기   국가데이터맵   데이터요청   데이터활용   정보공유

데이터찾기>   데이터목록   국가중점데이터   이슈데이터

어떤 공공데이터를 찾으시나요?

인기검색어  
2. 날씨

검색조건

분류체계

서비스유형

확장자

검색도움말

# End-to-End Machine Learning Project

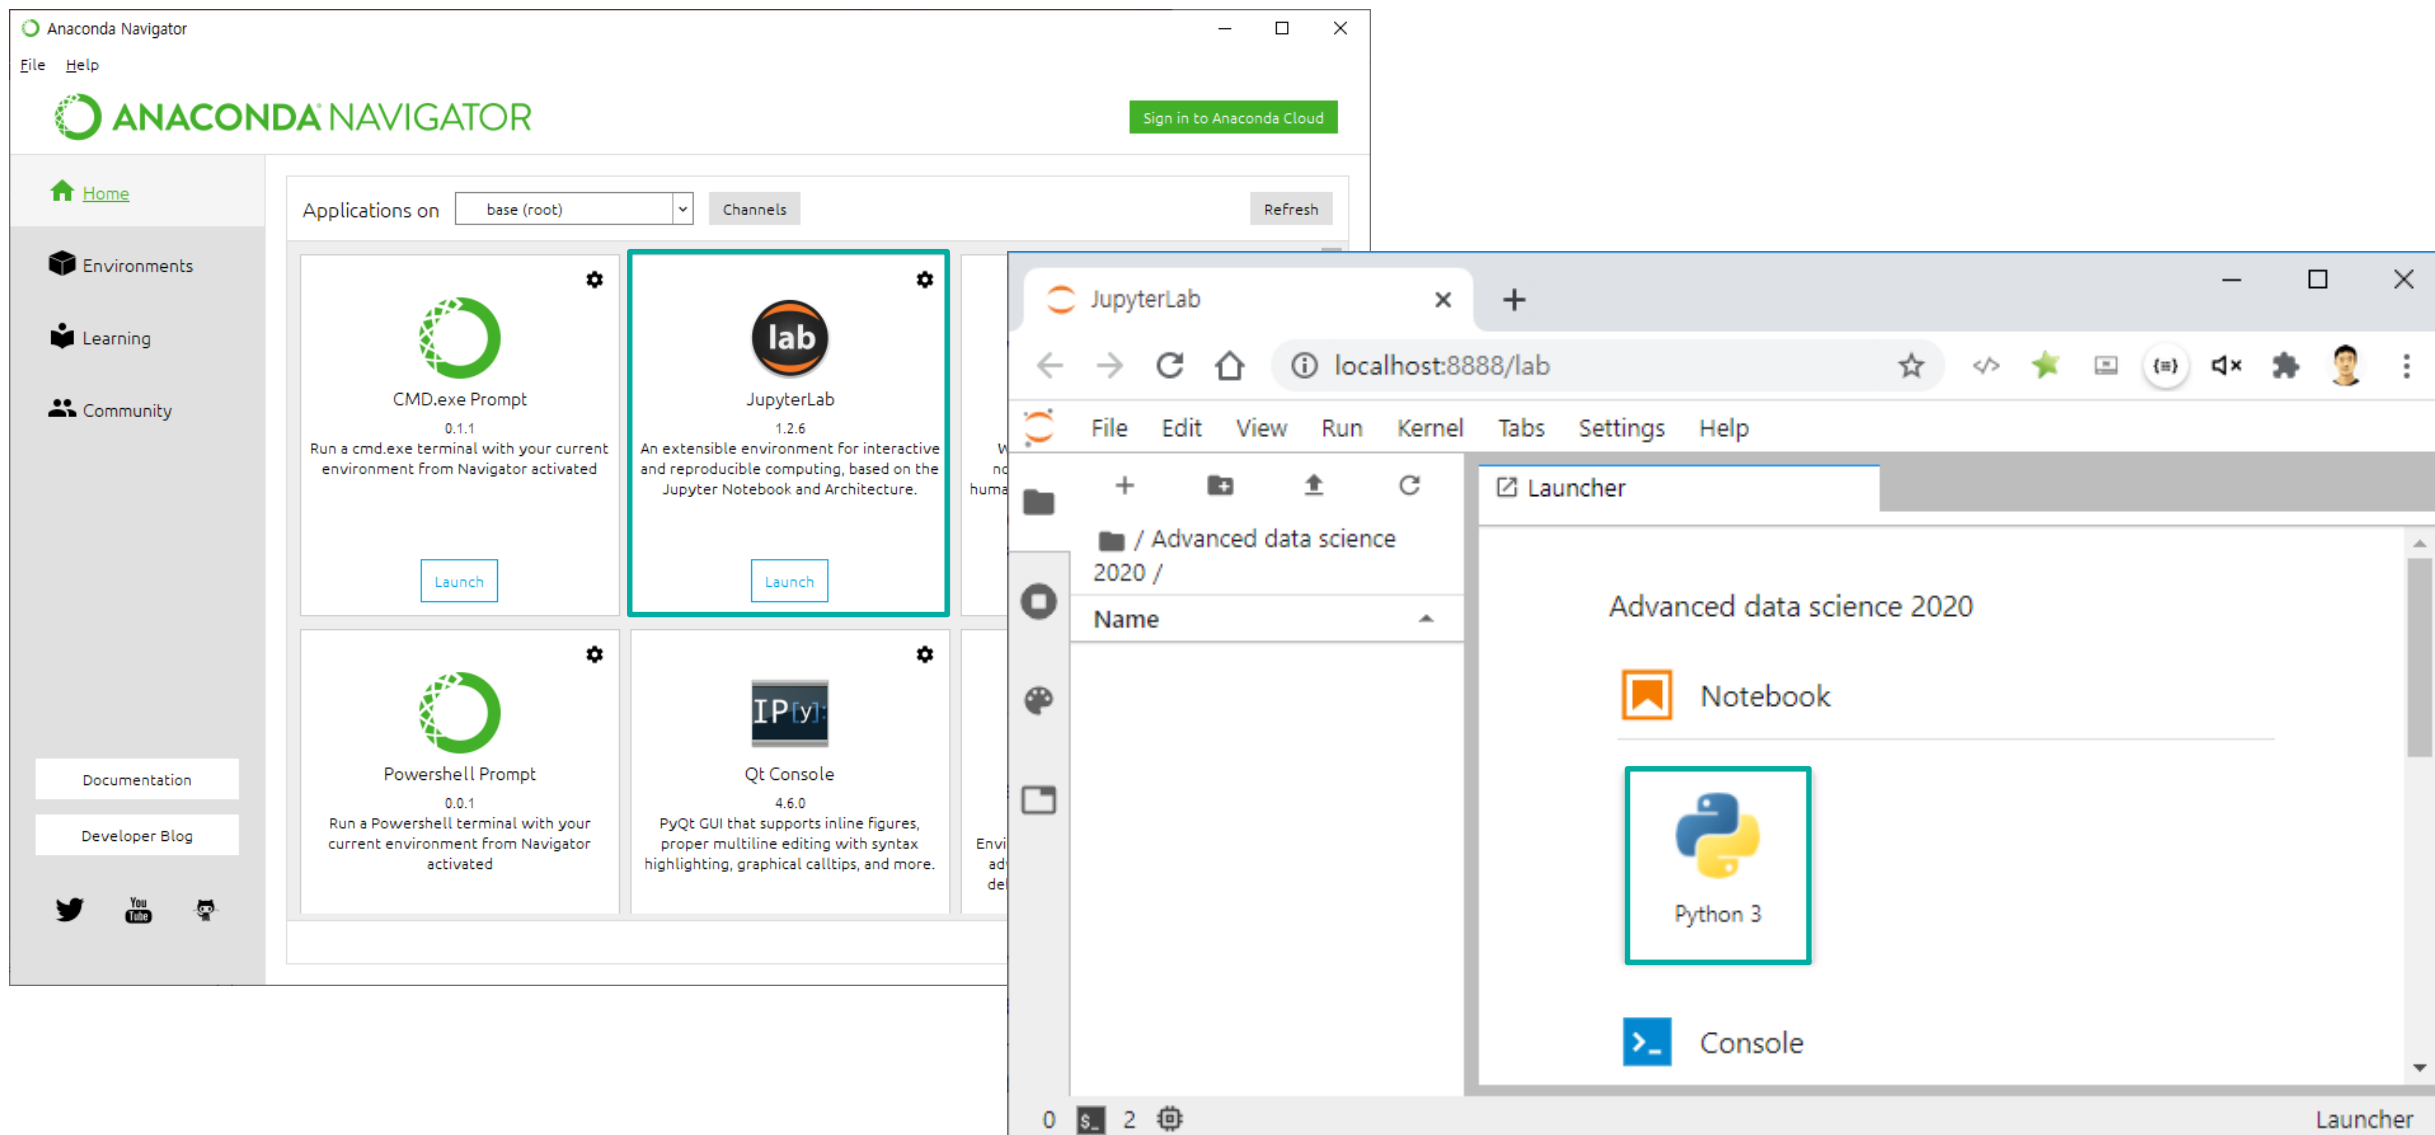
0. Look at the big picture.
1. Get the data.
2. Discover and visualize the data to gain insights.
3. Prepare the data for Machine Learning algorithms.
4. Select a model and train it.
5. Fine-tune your model.
6. Present your solution
7. Launch, monitor, and maintain your system.

# End-to-End Machine Learning Project

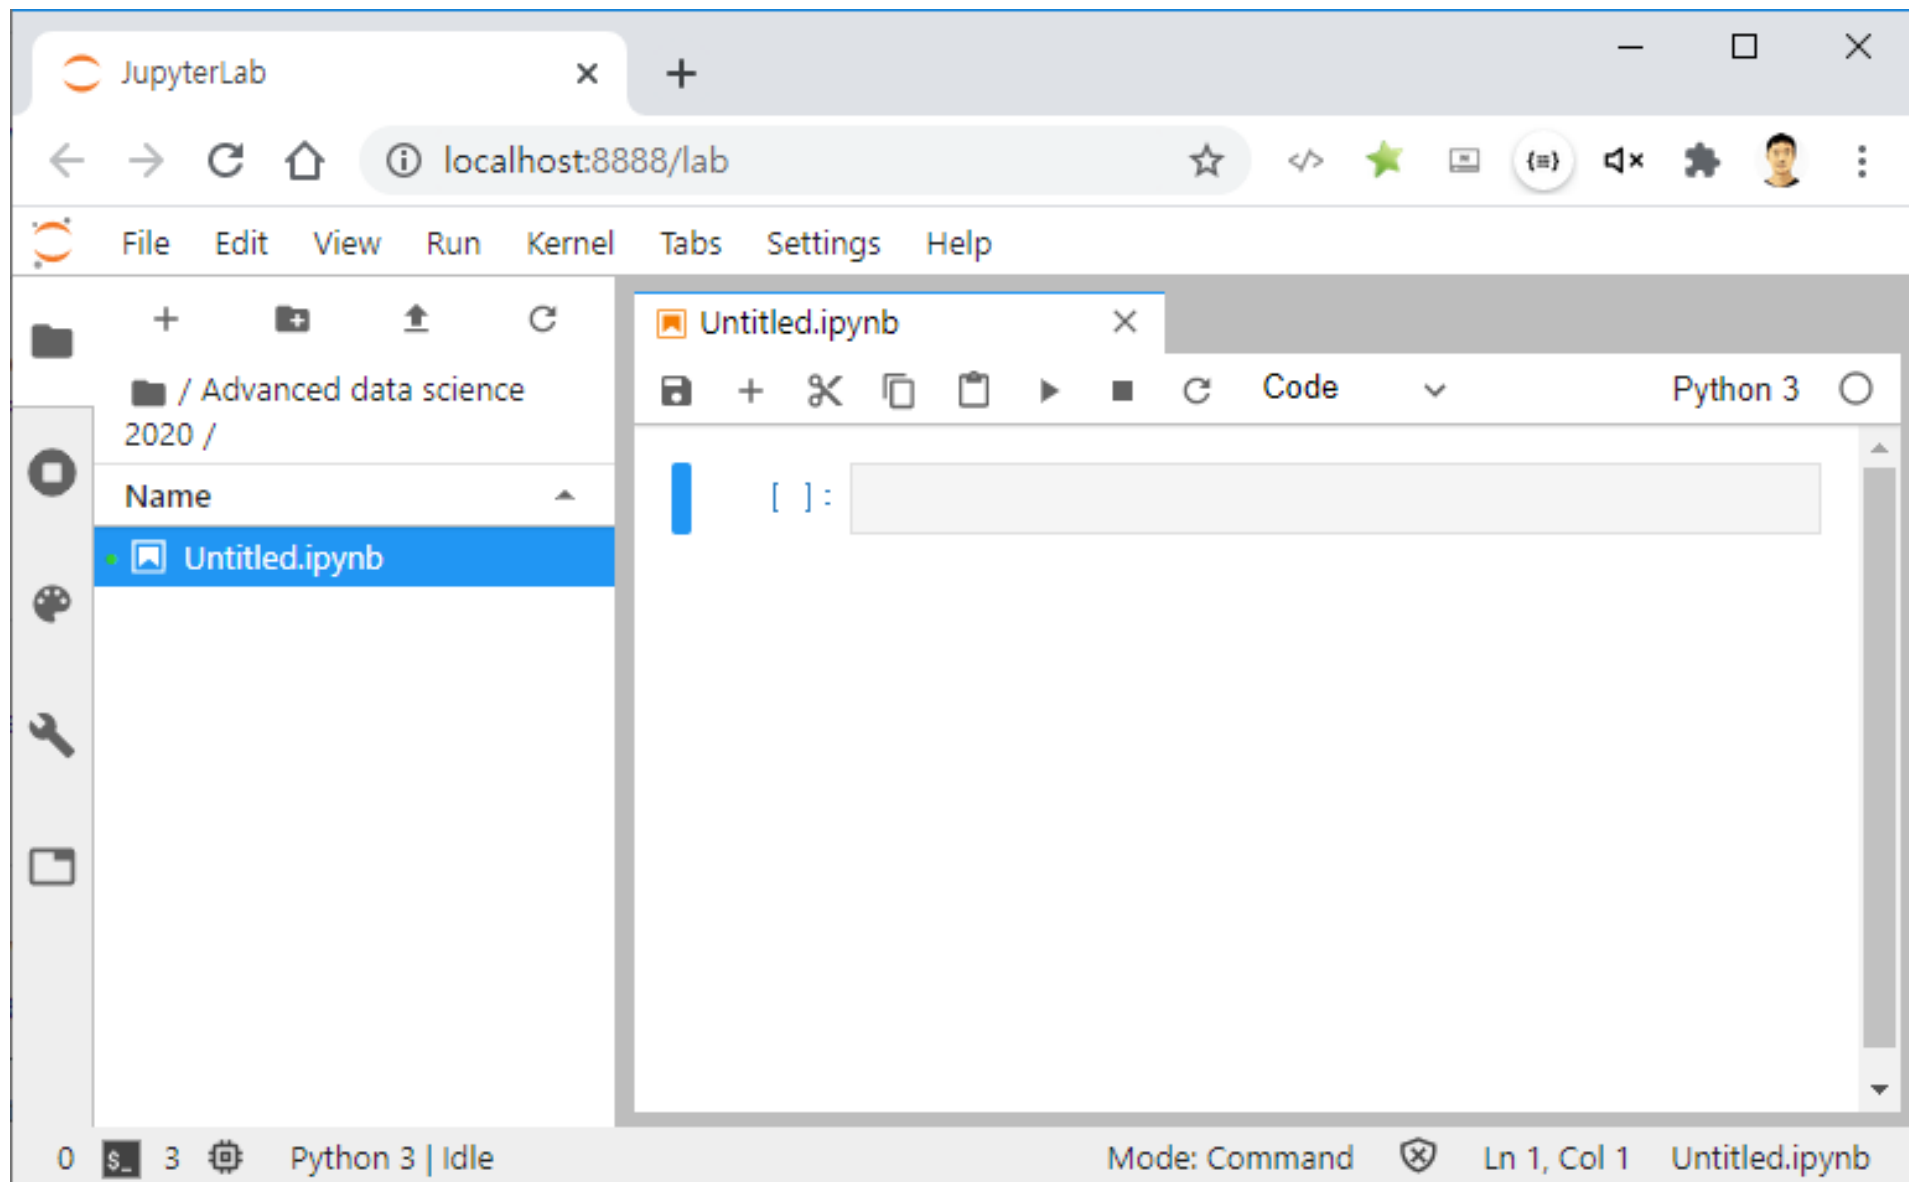
No.	Action	Package/library
0	Look at the big picture	—
1	Get the data	tarfile, urllib, pandas
2	Discover and visualize the data to gain insights	pandas, matplotlib
3	Prepare the data for Machine Learning algorithms	pandas, scikit-learn, numpy
4	Select a model and train it	
5	Fine-tune your model	
6	Present your solution	
7	Launch, monitor and maintain your system	joblib, flask

# Get ready

- Run Anaconda Navigator > JupyterLab > New Notebook



# Get ready





# End-to-End Machine Learning Project

No.	Action	Package/library
0	Look at the big picture	—
1	Get the data	tarfile, urllib, pandas
2	Discover and visualize the data to gain insights	pandas, matplotlib
3	Prepare the data for Machine Learning algorithms	pandas, scikit-learn, numpy
4	Select a model and train it	
5	Fine-tune your model	
6	Present your solution	
7	Launch, monitor and maintain your system	joblib, flask

# You are a recently hired data scientist.

- 1<sup>st</sup> task
  - Goal: building a **model of housing prices** in California
  - Data: California census data
    - **population**, **median income**, ..., **median housing price** for each **block group** (district)
      - 20,640 districts

housing.csv - Excel

파일 홈 삽입 페이지 레이아웃 수식 데이터 검토 보기 Acrobat 수행할 작업을 알려 주세요. Eung-Hee Kim 공유

붙여넣기 글꼴 맞춤 표시 형식 스타일 셀 편집

A20 : -122.26


	A	B	C	D	E	F	G	H	I	J
1	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value	ocean_proximity
2	-122.23	37.88	41	880	129	322	126	8.3252	452600	NEAR BAY
3	-122.22	37.86	21	7099	1106	2401	1138	8.3014	358500	NEAR BAY
4	-122.24	37.85	52	1467	190	496	177	7.2574	352100	NEAR BAY
5	-122.25	37.85	52	1274	235	558	219	5.6431	341300	NEAR BAY
6	-122.25	37.85	52	1627	280	565	259	3.8462	342200	NEAR BAY
7	-122.25	37.85	52	919	213	413	193	4.0368	269700	NEAR BAY
8	-122.25	37.84	52	2535	489	1094	514	3.6591	299200	NEAR BAY
9	-122.25	37.84	52	3104	687	1157	647	3.12	241400	NEAR BAY
10	-122.26	37.84	42	2555	665	1206	505	2.0004	226700	NEAR BAY

housing

준비 평균: 18085.61901 개수: 10 합계: 162770.5711 100%

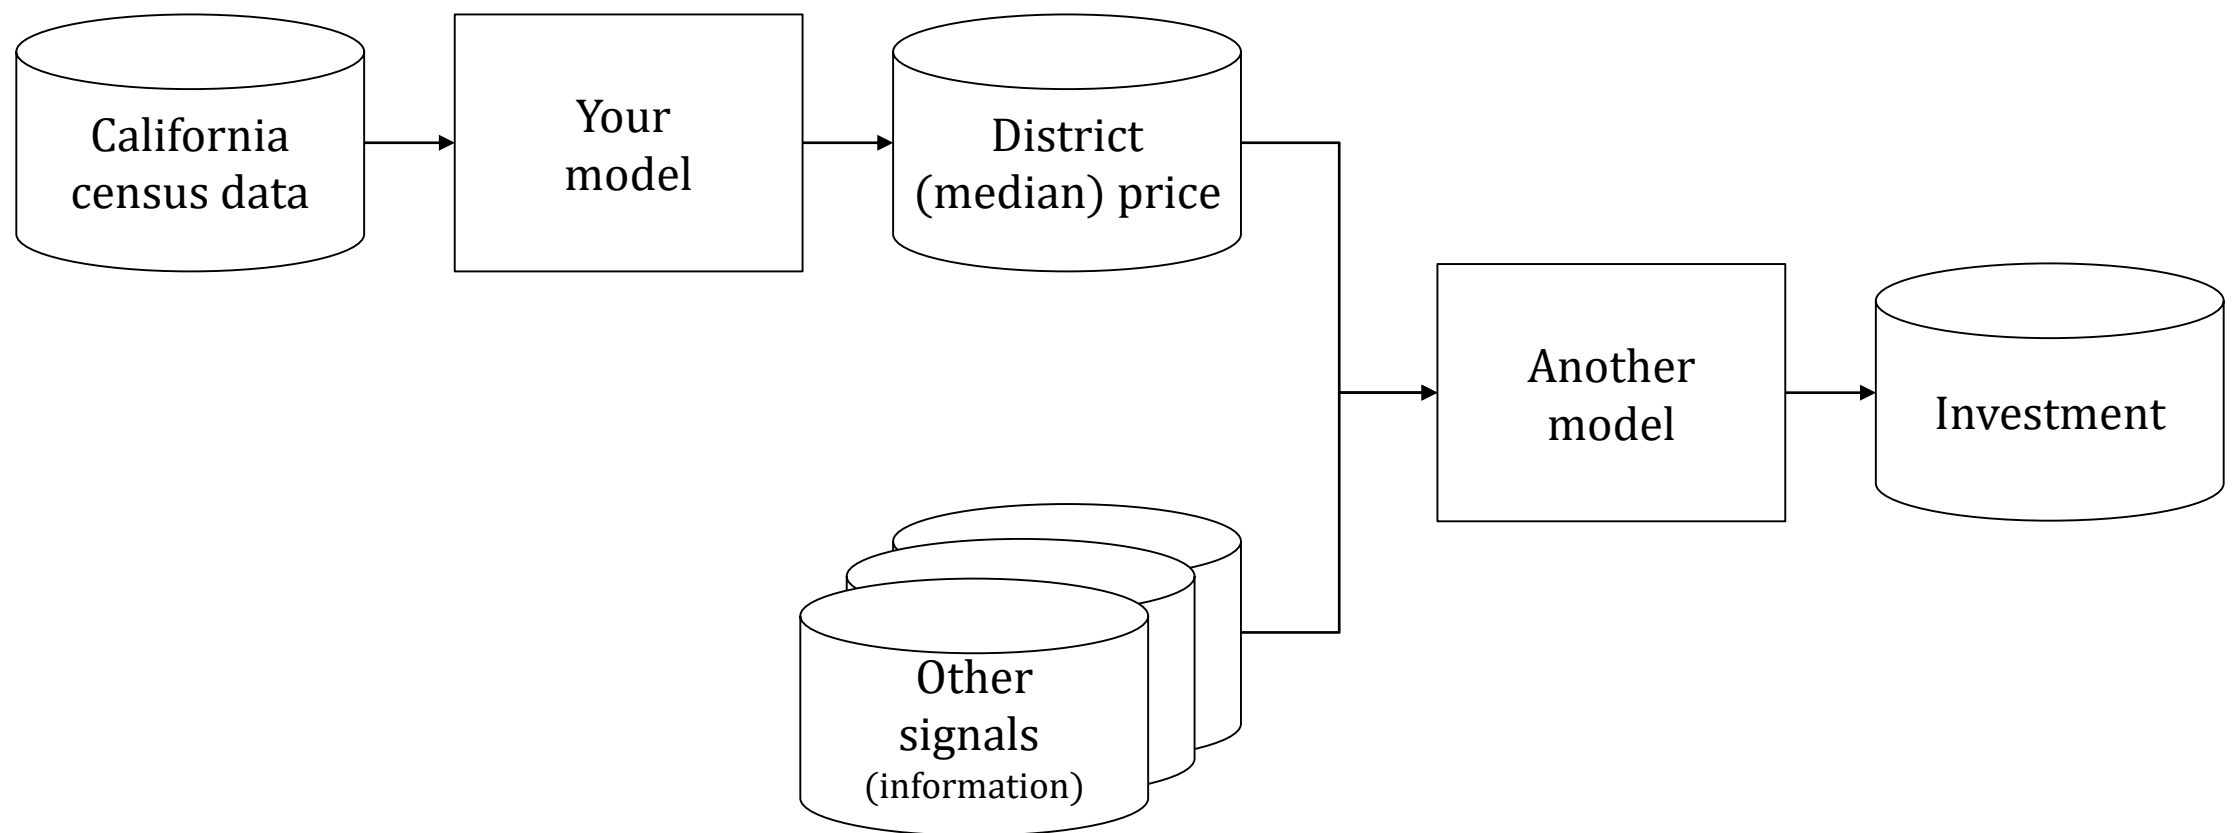
## 0.1 Frame the Problem

- Make it clear!
  - How does the company expect to use and benefit from your model?



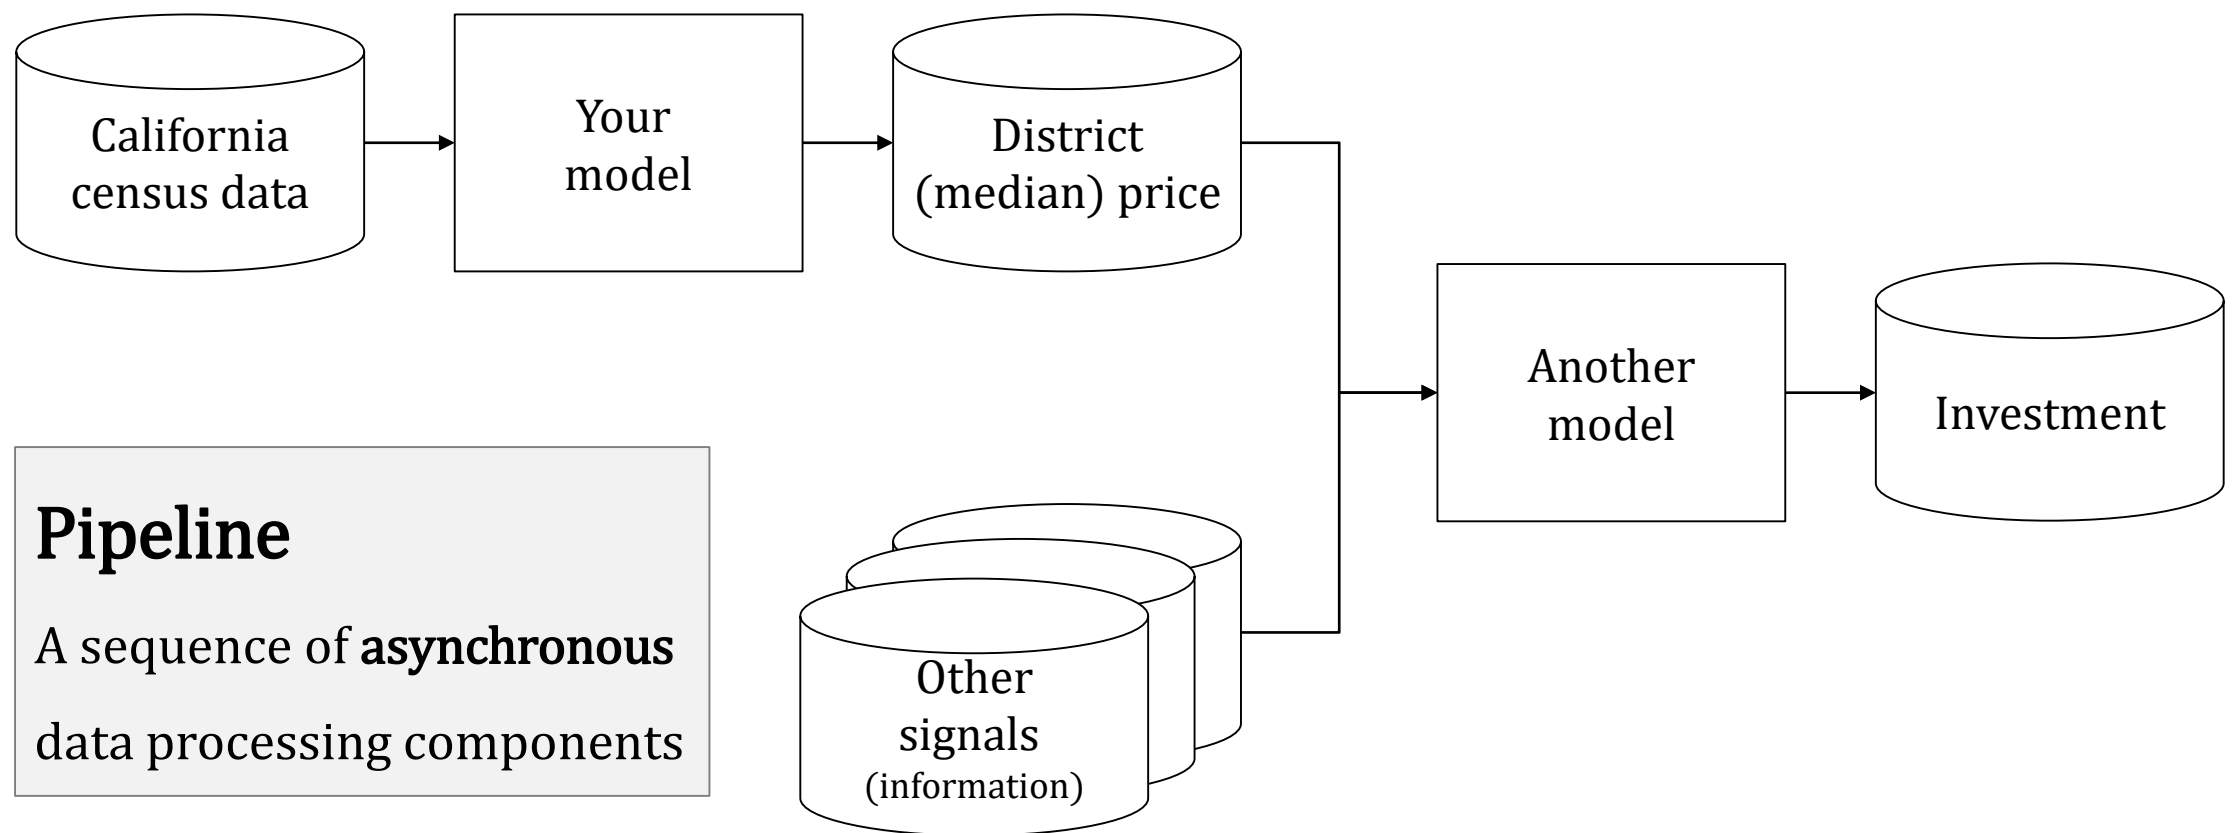
## 0.1 Frame the Problem

- Make it clear!
  - How does the company expect to use and benefit from your model?



## 0.1 Frame the Problem

- Make it clear!
  - How does the company expect to use and benefit from your model?



## 0.1 Frame the Problem

Supervised Learning	Unsupervised Learning	Reinforcement Learning
---------------------	-----------------------	------------------------

Classification	Regression	Something else
----------------	------------	----------------

Univariate regression	Multivariate regression
-----------------------	-------------------------

Batch learning	Online learning
----------------	-----------------

## 0.2 Select a Performance Measure

- Common machine learning notations
  - $m$ : the number of instances in the dataset
  - $\mathbf{x}^{(i)}$ : a vector of all the feature values (excluding the label) of the  $i^{th}$  instance
  - $y^{(i)}$ : the label (answer) of the  $i^{th}$  instance

## 0.2 Select a Performance Measure

- Common machine learning notations
  - $m$ : the number of instances in the dataset
  - $\mathbf{x}^{(i)}$ : a vector of all the feature values (excluding the label) of the  $i^{th}$  instance
  - $y^{(i)}$ : the label (answer) of the  $i^{th}$  instance

$$\mathbf{x}^{(1)} = \begin{pmatrix} -122.23 \\ 37.88 \\ 41 \\ 880 \\ \dots \end{pmatrix} \quad y^{(1)} = 452,600$$

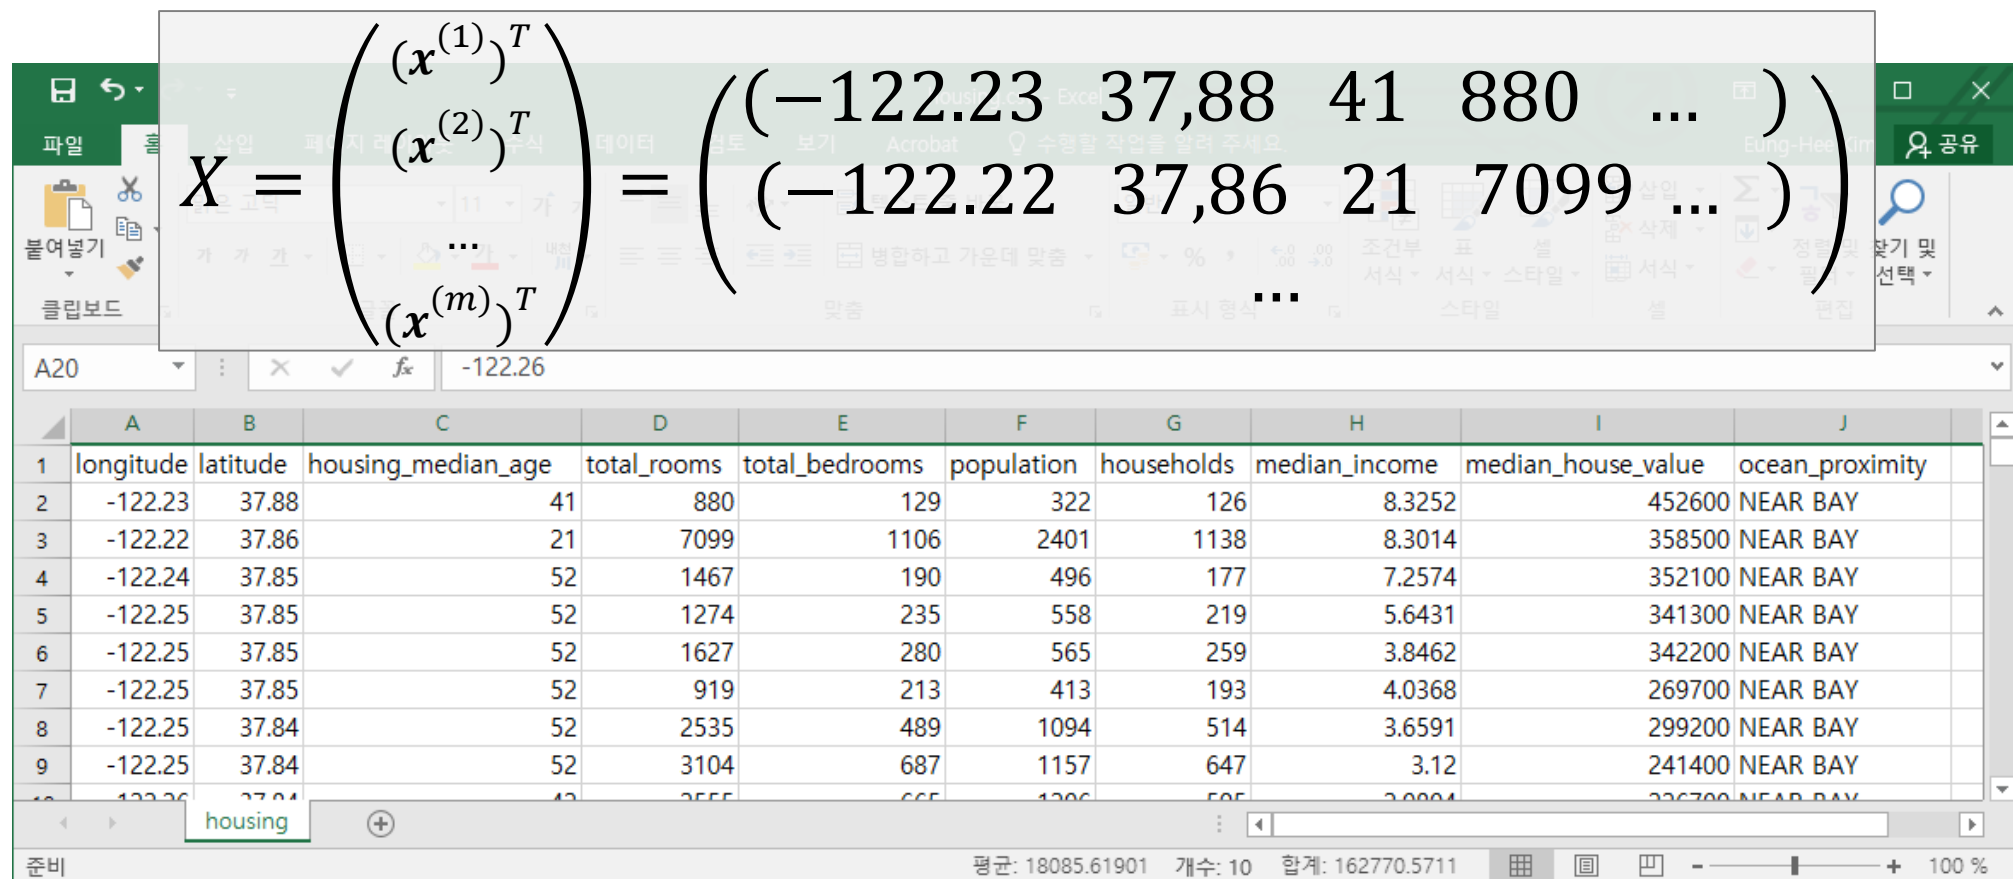


## 0.2 Select a Performance Measure

- Common machine learning notations
  - $m$ : the number of instances in the dataset
  - $\mathbf{x}^{(i)}$ : a vector of all the feature values (excluding the label) of the  $i^{th}$  instance
  - $y^{(i)}$ : the label (answer) of the  $i^{th}$  instance
  - $\mathbf{X}$ : a matrix containing all the feature values (excluding labels) of all instances
    - $i^{th}$  row of  $\mathbf{X}$  is the transpose of  $\mathbf{x}^{(i)}$ , denoted  $(\mathbf{x}^{(i)})^T$

## 0.2 Select a Performance Measure

- Common machine learning notations
  - $m$ : the number of instances in the dataset
  - $\mathbf{x}^{(i)}$ : a vector of all the feature values (excluding the label) of the  $i^{th}$  instance
  - $y^{(i)}$ : the label (answer) of the  $i^{th}$  instance
  - $\mathbf{X}$ : a matrix containing all the feature values (excluding labels) of all instances
    - $i^{th}$  row of  $\mathbf{X}$  is the transpose of  $\mathbf{x}^{(i)}$ , denoted  $(\mathbf{x}^{(i)})^T$



The image shows an Excel spreadsheet with a housing dataset. A formula overlay is present, defining the matrix  $\mathbf{X}$  as a matrix of feature vectors  $\mathbf{x}^{(i)}$  transposed. The spreadsheet data is as follows:

	A	B	C	D	E	F	G	H	I	J
1	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value	ocean_proximity
2	-122.23	37.88	41	880	129	322	126	8.3252	452600	NEAR BAY
3	-122.22	37.86	21	7099	1106	2401	1138	8.3014	358500	NEAR BAY
4	-122.24	37.85	52	1467	190	496	177	7.2574	352100	NEAR BAY
5	-122.25	37.85	52	1274	235	558	219	5.6431	341300	NEAR BAY
6	-122.25	37.85	52	1627	280	565	259	3.8462	342200	NEAR BAY
7	-122.25	37.85	52	919	213	413	193	4.0368	269700	NEAR BAY
8	-122.25	37.84	52	2535	489	1094	514	3.6591	299200	NEAR BAY
9	-122.25	37.84	52	3104	687	1157	647	3.12	241400	NEAR BAY
10	-122.26	37.84	42	2555	665	1206	505	2.9854	236700	NEAR BAY

## 0.2 Select a Performance Measure

- Common machine learning notations
  - $m$ : the number of instances in the dataset
  - $\mathbf{x}^{(i)}$ : a vector of all the feature values (excluding the label) of the  $i^{th}$  instance
  - $y^{(i)}$ : the label (answer) of the  $i^{th}$  instance
  - $\mathbf{X}$ : a matrix containing all the feature values (excluding labels) of all instances
    - $i^{th}$  row of  $\mathbf{X}$  is the transpose of  $\mathbf{x}^{(i)}$ , denoted  $(\mathbf{x}^{(i)})^T$
  - $h$ : your hypothesis (prediction function/model)
    - $h(\mathbf{x}^{(i)}) = \hat{y}^{(i)}$ : the predicted value by  $h$  for the  $i^{th}$  instance
    - $\hat{y}^{(i)} - y^{(i)}$ : the prediction error

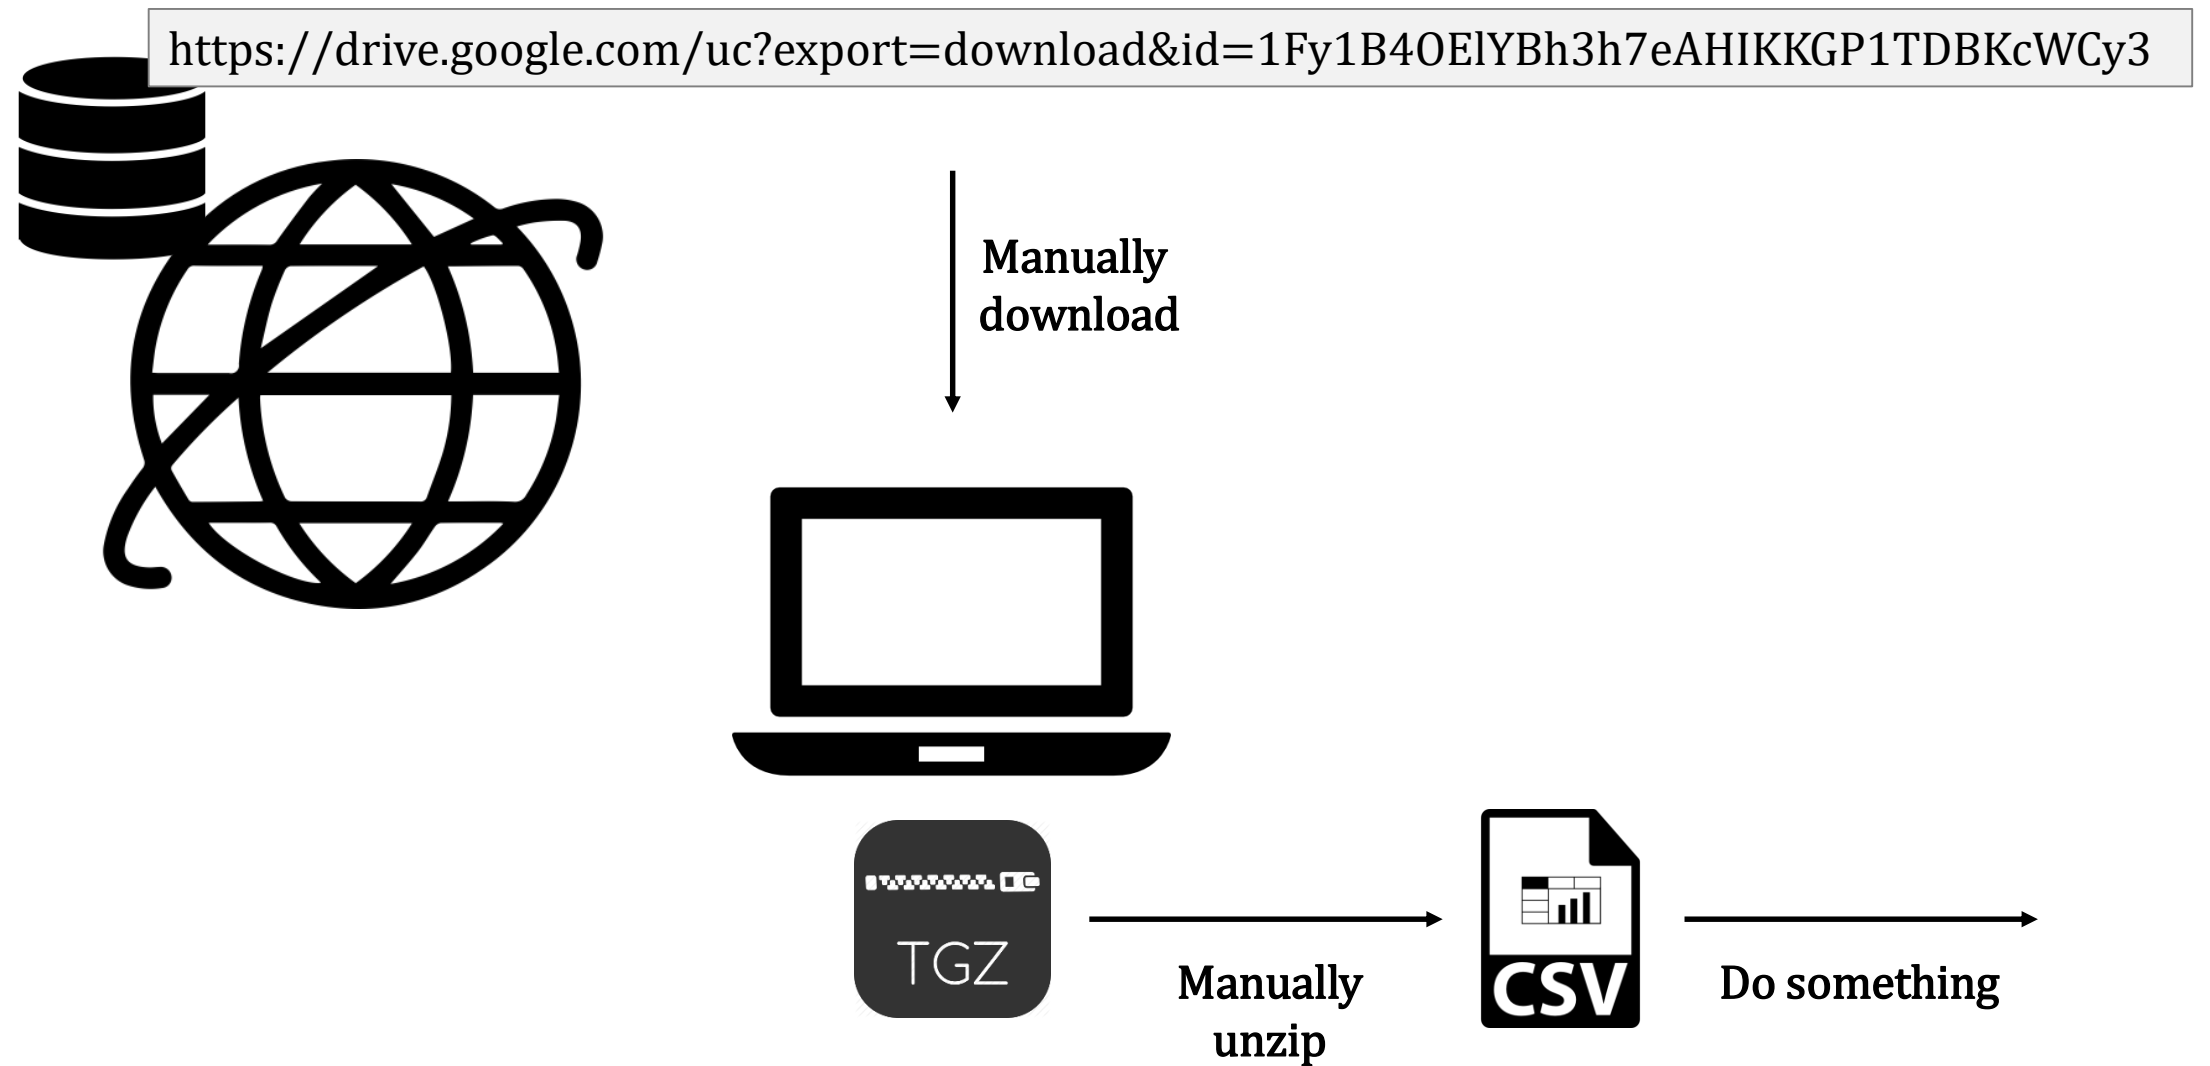
## 0.2 Select a Performance Measure

- Common machine learning notations
  - $m$ : the number of instances in the dataset
  - $\mathbf{x}^{(i)}$ : a vector of all the feature values (excluding the label) of the  $i^{th}$  instance
  - $y^{(i)}$ : the label (answer) of the  $i^{th}$  instance
  - $\mathbf{X}$ : a matrix containing all the feature values (excluding labels) of all instances
    - $i^{th}$  row of  $\mathbf{X}$  is the transpose of  $\mathbf{x}^{(i)}$ , denoted  $(\mathbf{x}^{(i)})^T$
  - $h$ : your hypothesis (prediction function/model)
    - $h(\mathbf{x}^{(i)}) = \hat{y}^{(i)}$ : the predicted value by  $h$  for the  $i^{th}$  instance
    - $\hat{y}^{(i)} - y^{(i)}$ : the prediction error
  - $\text{RMSE}(\mathbf{X}, h) = \sqrt{\frac{1}{m} \sum_{i=1}^m (h(\mathbf{x}^{(i)}) - y^{(i)})^2}$ 
    - Root Mean Square Error
  - $\text{MAE}(\mathbf{X}, h) = \frac{1}{m} \sum_{i=1}^m |h(\mathbf{x}^{(i)}) - y^{(i)}|$ 
    - Mean Absolute Error

# End-to-End Machine Learning Project


No.	Action	Package/library
0	Look at the big picture	—
1	Get the data	tarfile, urllib, pandas
2	Discover and visualize the data to gain insights	pandas, matplotlib
3	Prepare the data for Machine Learning algorithms	pandas, scikit-learn, numpy
4	Select a model and train it	
5	Fine-tune your model	
6	Present your solution	
7	Launch, monitor and maintain your system	joblib, flask

# California census data



# California census data

Frequently  
updated



<https://drive.google.com/uc?export=download&id=1Fy1B40ElyBh3h7eAHIKKGP1TDBKcWCy3>



Manually  
download




Manually  
unzip




Do something

# California census data

Frequently  
updated



<https://drive.google.com/uc?export=download&id=1Fy1B40ElyBh3h7eAHIKKGP1TDBKcWCy3>



Automatically  
download



Automatically  
unzip



Do something

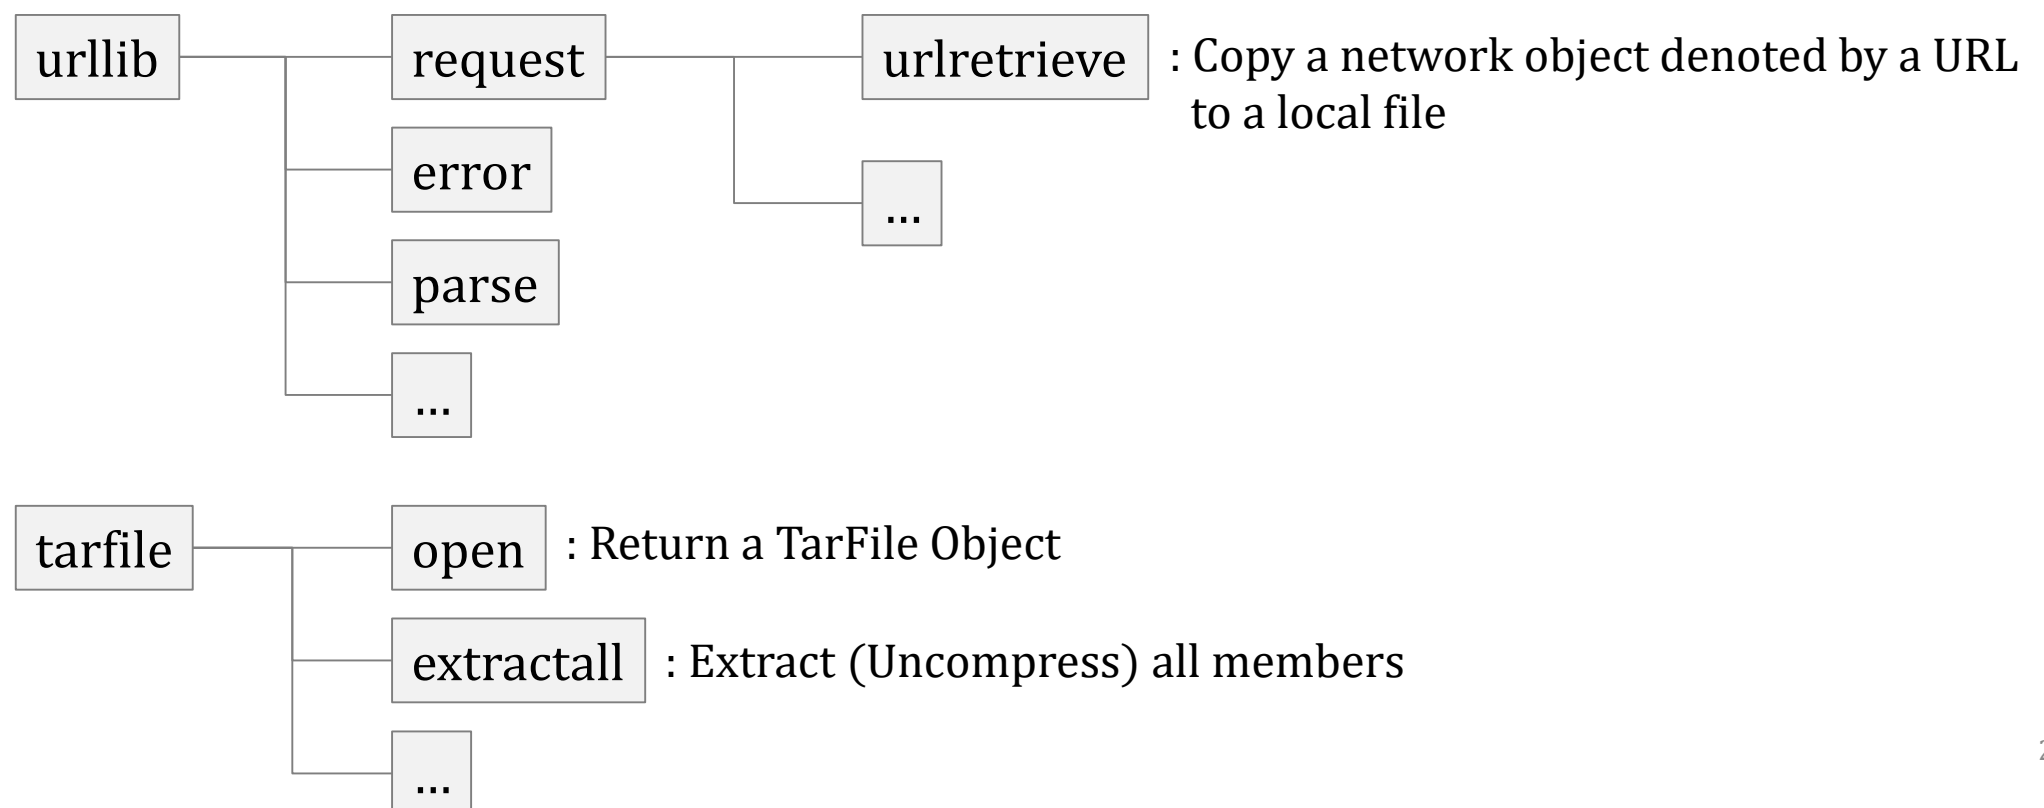


## 1.1 Download and uncompress data

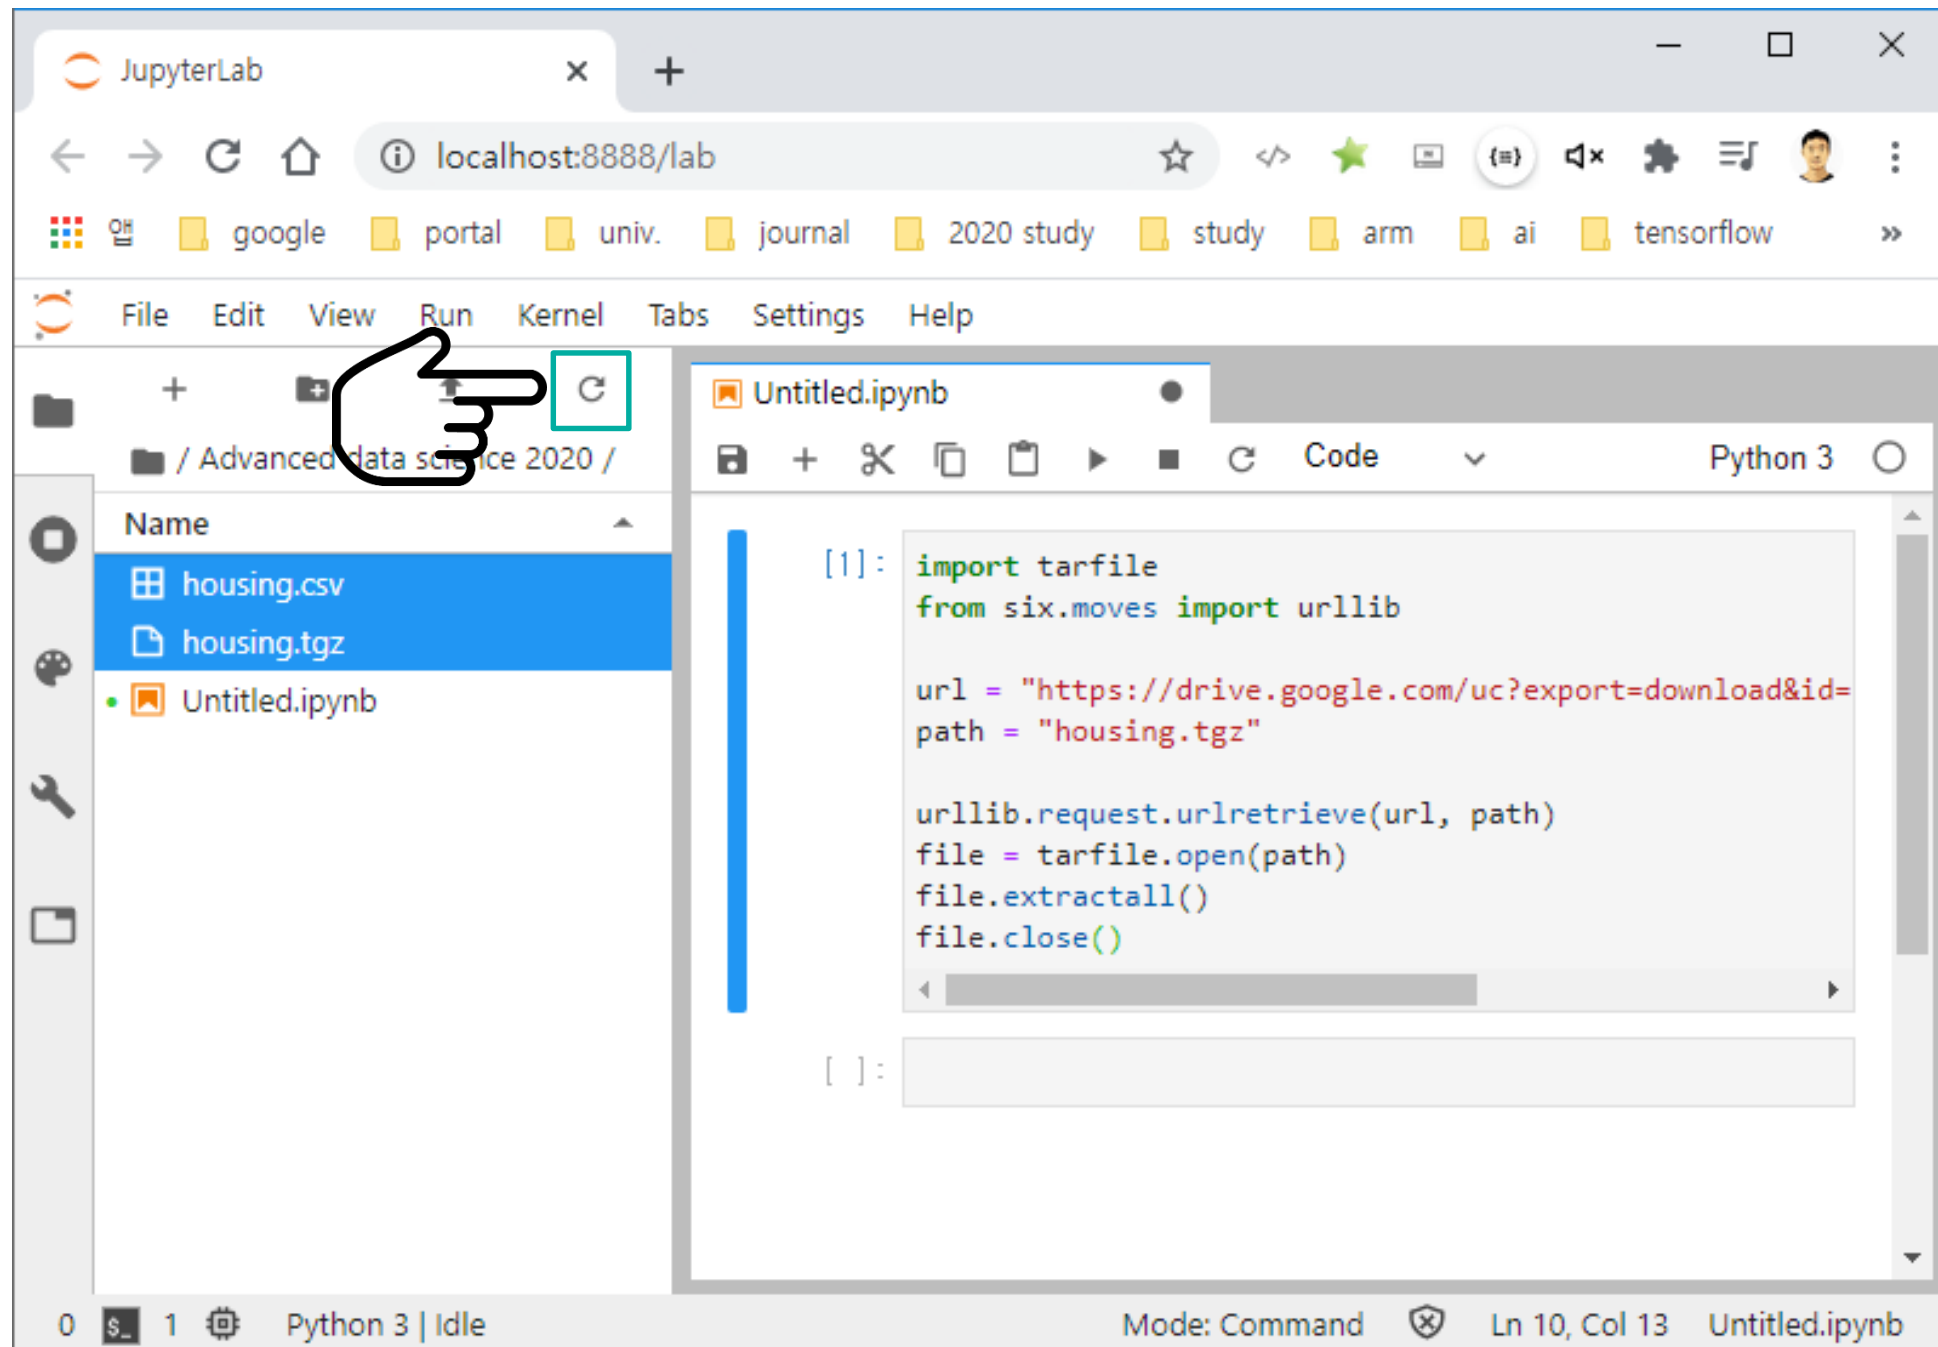
```
[1]: import tarfile
      from six.moves import urllib

      url = "https://drive.google.com/uc?export=download&id=1Fy1B40ElYBh3h7eAHIKKGP1TDBKcWCy3"
      path = "housing.tgz"

      urllib.request.urlretrieve(url, path)
      file = tarfile.open(path)
      file.extractall()
      file.close()
```



## 1.1 Download and uncompress data



## 1.2 Read data

```
import tarfile
from six.moves import urllib

url = "https://drive.google.com/uc?export=download&id=1Fy1B40E"
path = "housing.tgz"

urllib.request.urlretrieve(url, path)
file = tarfile.open(path)
file.extractall()
file.close()
```



```
import pandas as pd

data = pd.read_csv("housing.csv")
data.head()
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms
0	-122.23	37.88	41.0	880.0	129.0
1	-122.22	37.86	21.0	7099.0	1106.0
2	-122.24	37.85	52.0	1467.0	190.0
3	-122.25	37.85	52.0	1274.0	235.0
4	-122.25	37.85	52.0	1627.0	280.0

## 1.3 Take a Quick Look

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 20640 entries, 0 to 20639  
Data columns (total 10 columns):  
longitude                20640 non-null float64  
latitude                 20640 non-null float64  
housing_median_age       20640 non-null float64  
total_rooms              20640 non-null float64  
total_bedrooms           20433 non-null float64  
population               20640 non-null float64  
households               20640 non-null float64  
median_income             20640 non-null float64  
median_house_value       20640 non-null float64  
ocean_proximity          20640 non-null object  
dtypes: float64(9), object(1)  
memory usage: 1.6+ MB
```

## 1.3 Take a Quick Look

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 20640 entries, 0 to 20639  
Data columns (total 10 columns):  
longitude                20640 non-null float64  
latitude                 20640 non-null float64  
housing_median_age       20640 non-null float64  
total_rooms               20640 non-null float64  
total_bedrooms           20433 non-null float64  
population               20640 non-null float64  
households               20640 non-null float64  
median_income            20640 non-null float64  
median_house_value       20640 non-null float64  
ocean_proximity          20640 non-null object  
dtypes: float64(9), object(1)  
memory usage: 1.6+ MB
```

Missing values

Categorical values

## 1.3 Take a Quick Look

```
data["ocean_proximity"].value_counts()
```

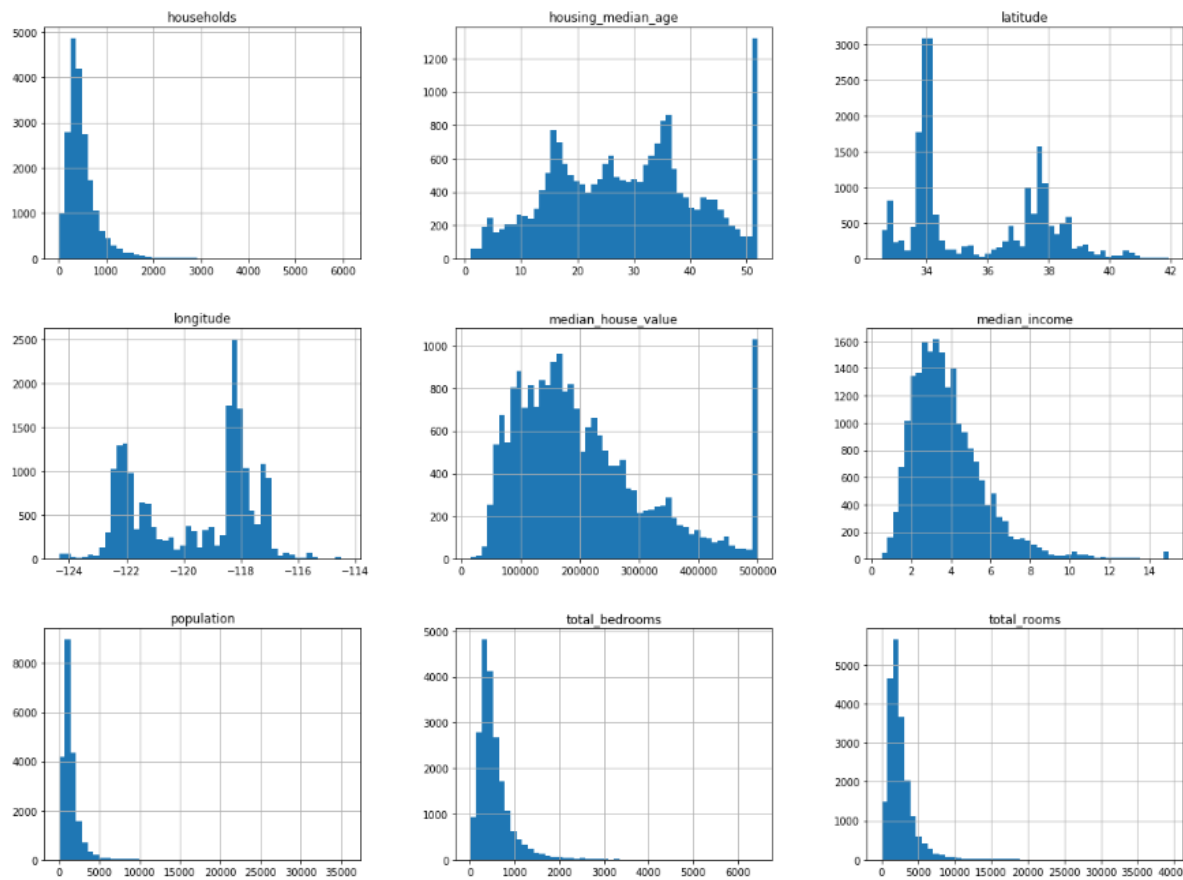
```
<1H OCEAN      9136  
INLAND         6551  
NEAR OCEAN     2658  
NEAR BAY       2290  
ISLAND          5  
Name: ocean_proximity, dtype: int64
```

## 1.3 Take a Quick Look

```
import matplotlib.pyplot as plt

data.hist(bins=50, figsize=(20, 15))

plt.show()
```



Right skewed

Difference scales

Preprocessed  
values

# End-to-End Machine Learning Project

No.	Action	Package/library
0	Look at the big picture	—
1	Get the data	tarfile, urllib, pandas
2	Discover and visualize the data to gain insights	pandas, matplotlib
3	Prepare the data for Machine Learning algorithms	pandas, scikit-learn, numpy
4	Select a model and train it	
5	Fine-tune your model	
6	Present your solution	
7	Launch, monitor and maintain your system	joblib, flask



## 2.1 Visualize Geographical Data

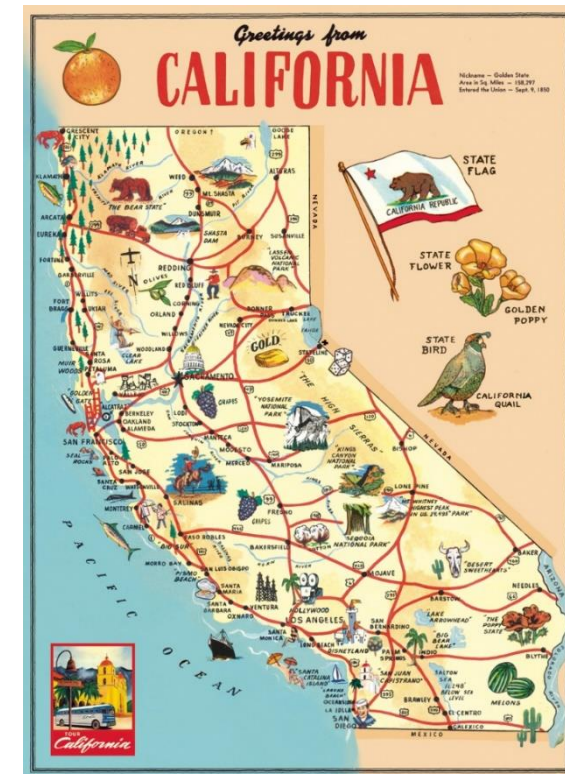



data  
(DataFrame)


	longitude	latitude	housing_median_age	total_rooms	total_bedrooms
0	-122.23	37.88	41.0	880.0	129.0
1	-122.22	37.86	21.0	7099.0	1106.0
2	-122.24	37.85	52.0	1467.0	190.0
3	-122.25	37.85	52.0	1274.0	235.0
4	-122.25	37.85	52.0	1627.0	280.0

```
data.plot(kind="scatter", x="longitude", y="latitude", alpha=0.1)
```

## 2.1 Visualize Geographical Data



## 2.1 Visualize Geographical Data




Size
population

Color
median house value

## 2.1 Visualize Geographical Data

```
data.plot(kind="scatter", x="longitude", y="latitude", alpha=0.3,  
          s=data["population"]/100, label="population",  
          c="median_house_value", cmap=plt.get_cmap("rainbow"),  
          colorbar=True, sharex=False)
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x2976a0c31c8>



## 2.2 Look for Correlation

```
corr_matrix = data.corr()  
corr_matrix["median_house_value"].sort_values(ascending=False)
```


```
median_house_value    1.000000  
median_income          0.688075  
total_rooms            0.134153  
housing_median_age     0.105623  
households             0.065843  
total_bedrooms         0.049686  
population            -0.024650  
longitude              -0.045967  
latitude               -0.144160  
Name: median_house_value, dtype: float64
```

- **Correlation:** how tightly the instances are clustered about a straight line.
  - Range:  $-1$  and  $+1$

## 2.2 Look for Correlation

```
data.plot(kind="scatter",  
          x="median_income",  
          y="median_house_value",  
          alpha=0.1)
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x1fe07479470>



## 2.3 Feature extraction

(by manual combinations)

Attribute (Feature)
longitude
latitude
housing_median_age
total_rooms
total_bedrooms
population
households
median_income
median_house_value
ocean_proximity



## 2.3 Feature extraction

(by manual combinations)

Attribute (Feature)
longitude
latitude
housing_median_age
total_rooms
total_bedrooms
population
households
median_income
median_house_value
ocean_proximity
$\text{rooms\_per\_household} = \frac{\text{total\_bedrooms}}{\text{households}}$
$\text{bedrooms\_per\_room} = \frac{\text{total\_bedrooms}}{\text{total\_rooms}}$



## 2.3 Feature extraction

(by manual combinations)

Original data		
...	...	...

+

Extended data	
rooms_per_household	bedrooms_per_room

```
data["rooms_per_household"] = data["total_rooms"]/data["households"]
data["bedrooms_per_room"] = data["total_bedrooms"]/data["total_rooms"]
data.head()
```

ne	median_house_value	ocean_proximity	rooms_per_household	bedrooms_per_room
52	452600.0	NEAR BAY	6.984127	0.146591
14	358500.0	NEAR BAY	6.238137	0.155797
74	352100.0	NEAR BAY	8.288136	0.129516
31	341300.0	NEAR BAY	5.817352	0.184458
62	342200.0	NEAR BAY	6.281853	0.172096



## 2.3 Feature extraction

(by manual combinations)

```
corr_matrix = data.corr()  
corr_matrix["median_house_value"].sort_values(ascending=False)
```

median_house_value	1.000000
median_income	0.688075
rooms_per_household	0.151948
total_rooms	0.134153
housing_median_age	0.105623
households	0.065843
total_bedrooms	0.049686
population	-0.024650
longitude	-0.045967
latitude	-0.144160
bedrooms_per_room	-0.255880

Name: median\_house\_value, dtype: float64

# End-to-End Machine Learning Project

No.	Action	Package/library
0	Look at the big picture	—
1	Get the data	tarfile, urllib, pandas
2	Discover and visualize the data to gain insights	pandas, matplotlib
3	Prepare the data for Machine Learning algorithms	pandas, scikit-learn
4	Select a model and train it	
5	Fine-tune your model	
6	Present your solution	
7	Launch, monitor and maintain your system	joblib, flask


## 3.1 Handling Categorical Attributes (Feature values)

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   longitude              20640 non-null  float64
1   latitude               20640 non-null  float64
2   housing_median_age     20640 non-null  float64
3   total_rooms            20640 non-null  float64
4   total_bedrooms         20433 non-null  float64
5   population             20640 non-null  float64
6   households              20640 non-null  float64
7   median_income          20640 non-null  float64
8   median_house_value     20640 non-null  float64
9   ocean_proximity        20640 non-null  object
10  rooms_per_household    20640 non-null  float64
11  bedrooms_per_room      20433 non-null  float64
dtypes: float64(11), object(1)
memory usage: 1.9+ MB
```

Categorical values

## Transformation #1



## Transformation #2

ocean_proximity
<1H OCEAN
NEAR OCEAN
INLAND
NEAR BAY
ISLAND



	ocean_proximity				
<1H OCEAN	1	0	0	0	0
NEAR OCEAN	0	1	0	0	0
INLAND	0	0	1	0	0
NEAR BAY	0	0	0	1	0
ISLAND	0	0	0	0	1


**One hot encoding**

### 3.1 Handling Categorical Attributes (Feature values)

```
from sklearn.preprocessing import OneHotEncoder

encoder = OneHotEncoder()
ocean = data[["ocean_proximity"]]
ocean = encoder.fit_transform(ocean)
```

ocean_proximity	
0	NEAR BAY
1	NEAR BAY
2	NEAR BAY
3	NEAR BAY
4	NEAR BAY
...	...
20635	INLAND
20636	INLAND
20637	INLAND
20638	INLAND
20639	INLAND



0	0	0	1	0
0	0	0	1	0
0	0	0	1	0
0	0	0	1	0
0	0	0	1	0
...				
0	1	0	0	0
0	1	0	0	0
0	1	0	0	0
0	1	0	0	0
0	1	0	0	0

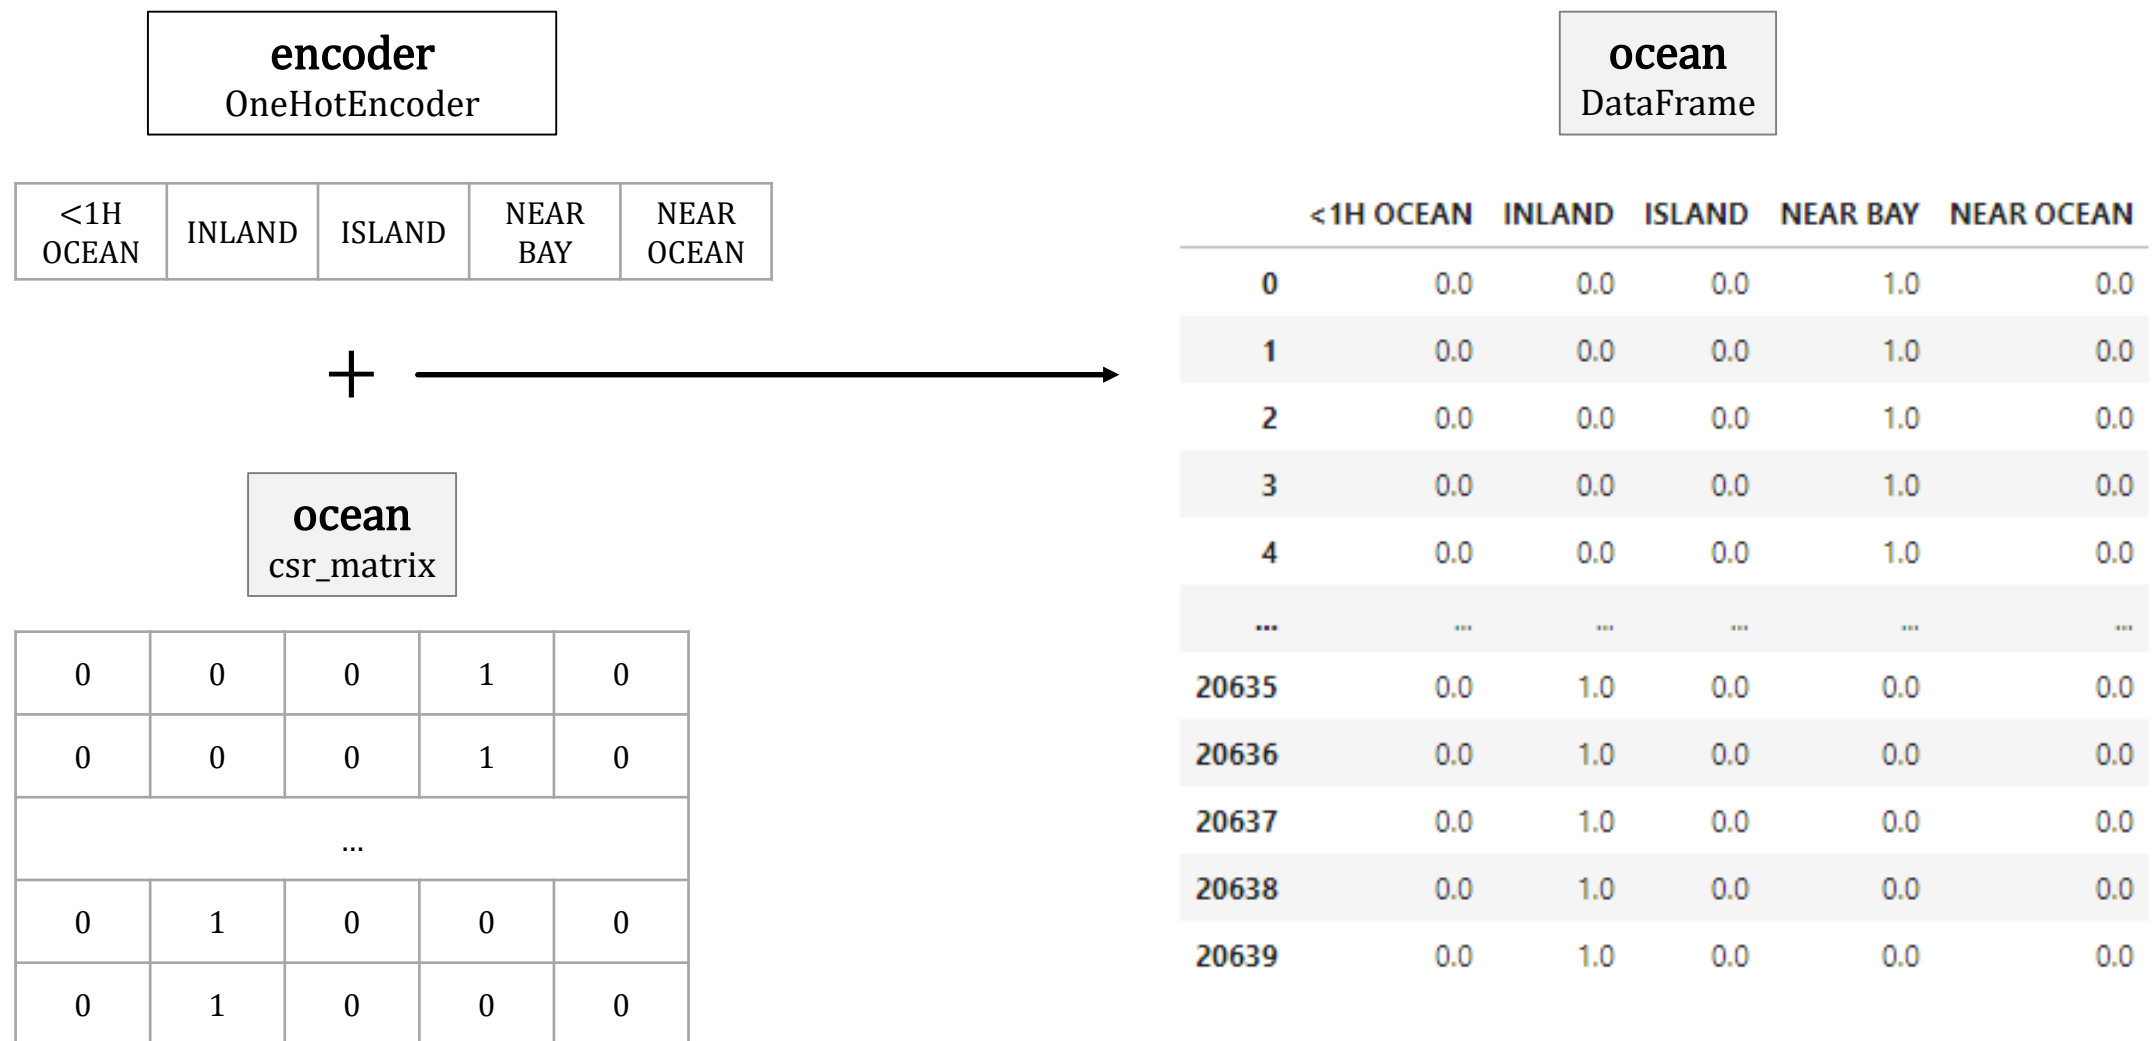
### 3.1 Handling Categorical Attributes (Feature values)

```
print(ocean.toarray())  
print(encoder.categories_)
```

```
[[0. 0. 0. 1. 0.]  
 [0. 0. 0. 1. 0.]  
 [0. 0. 0. 1. 0.]  
 ...  
 [0. 1. 0. 0. 0.]  
 [0. 1. 0. 0. 0.]  
 [0. 1. 0. 0. 0.]  
 [array(['<1H OCEAN', 'INLAND', 'ISLAND', 'NEAR BAY', 'NEAR OCEAN'],  
       dtype=object)]
```



### 3.1 Handling Categorical Attributes (Feature values)



### 3.1 Handling Categorical Attributes (Feature values)

```
ocean = pd.DataFrame(columns=encoder.categories_[0],  
                      data=ocean.toarray())
```

ocean

	<1H OCEAN	INLAND	ISLAND	NEAR BAY	NEAR OCEAN
0	0.0	0.0	0.0	1.0	0.0
1	0.0	0.0	0.0	1.0	0.0
2	0.0	0.0	0.0	1.0	0.0
3	0.0	0.0	0.0	1.0	0.0
4	0.0	0.0	0.0	1.0	0.0
...	...	...	...	...	...
20635	0.0	1.0	0.0	0.0	0.0
20636	0.0	1.0	0.0	0.0	0.0
20637	0.0	1.0	0.0	0.0	0.0
20638	0.0	1.0	0.0	0.0	0.0
20639	0.0	1.0	0.0	0.0	0.0

## 3.1 Handling Categorical Attributes (Feature values)

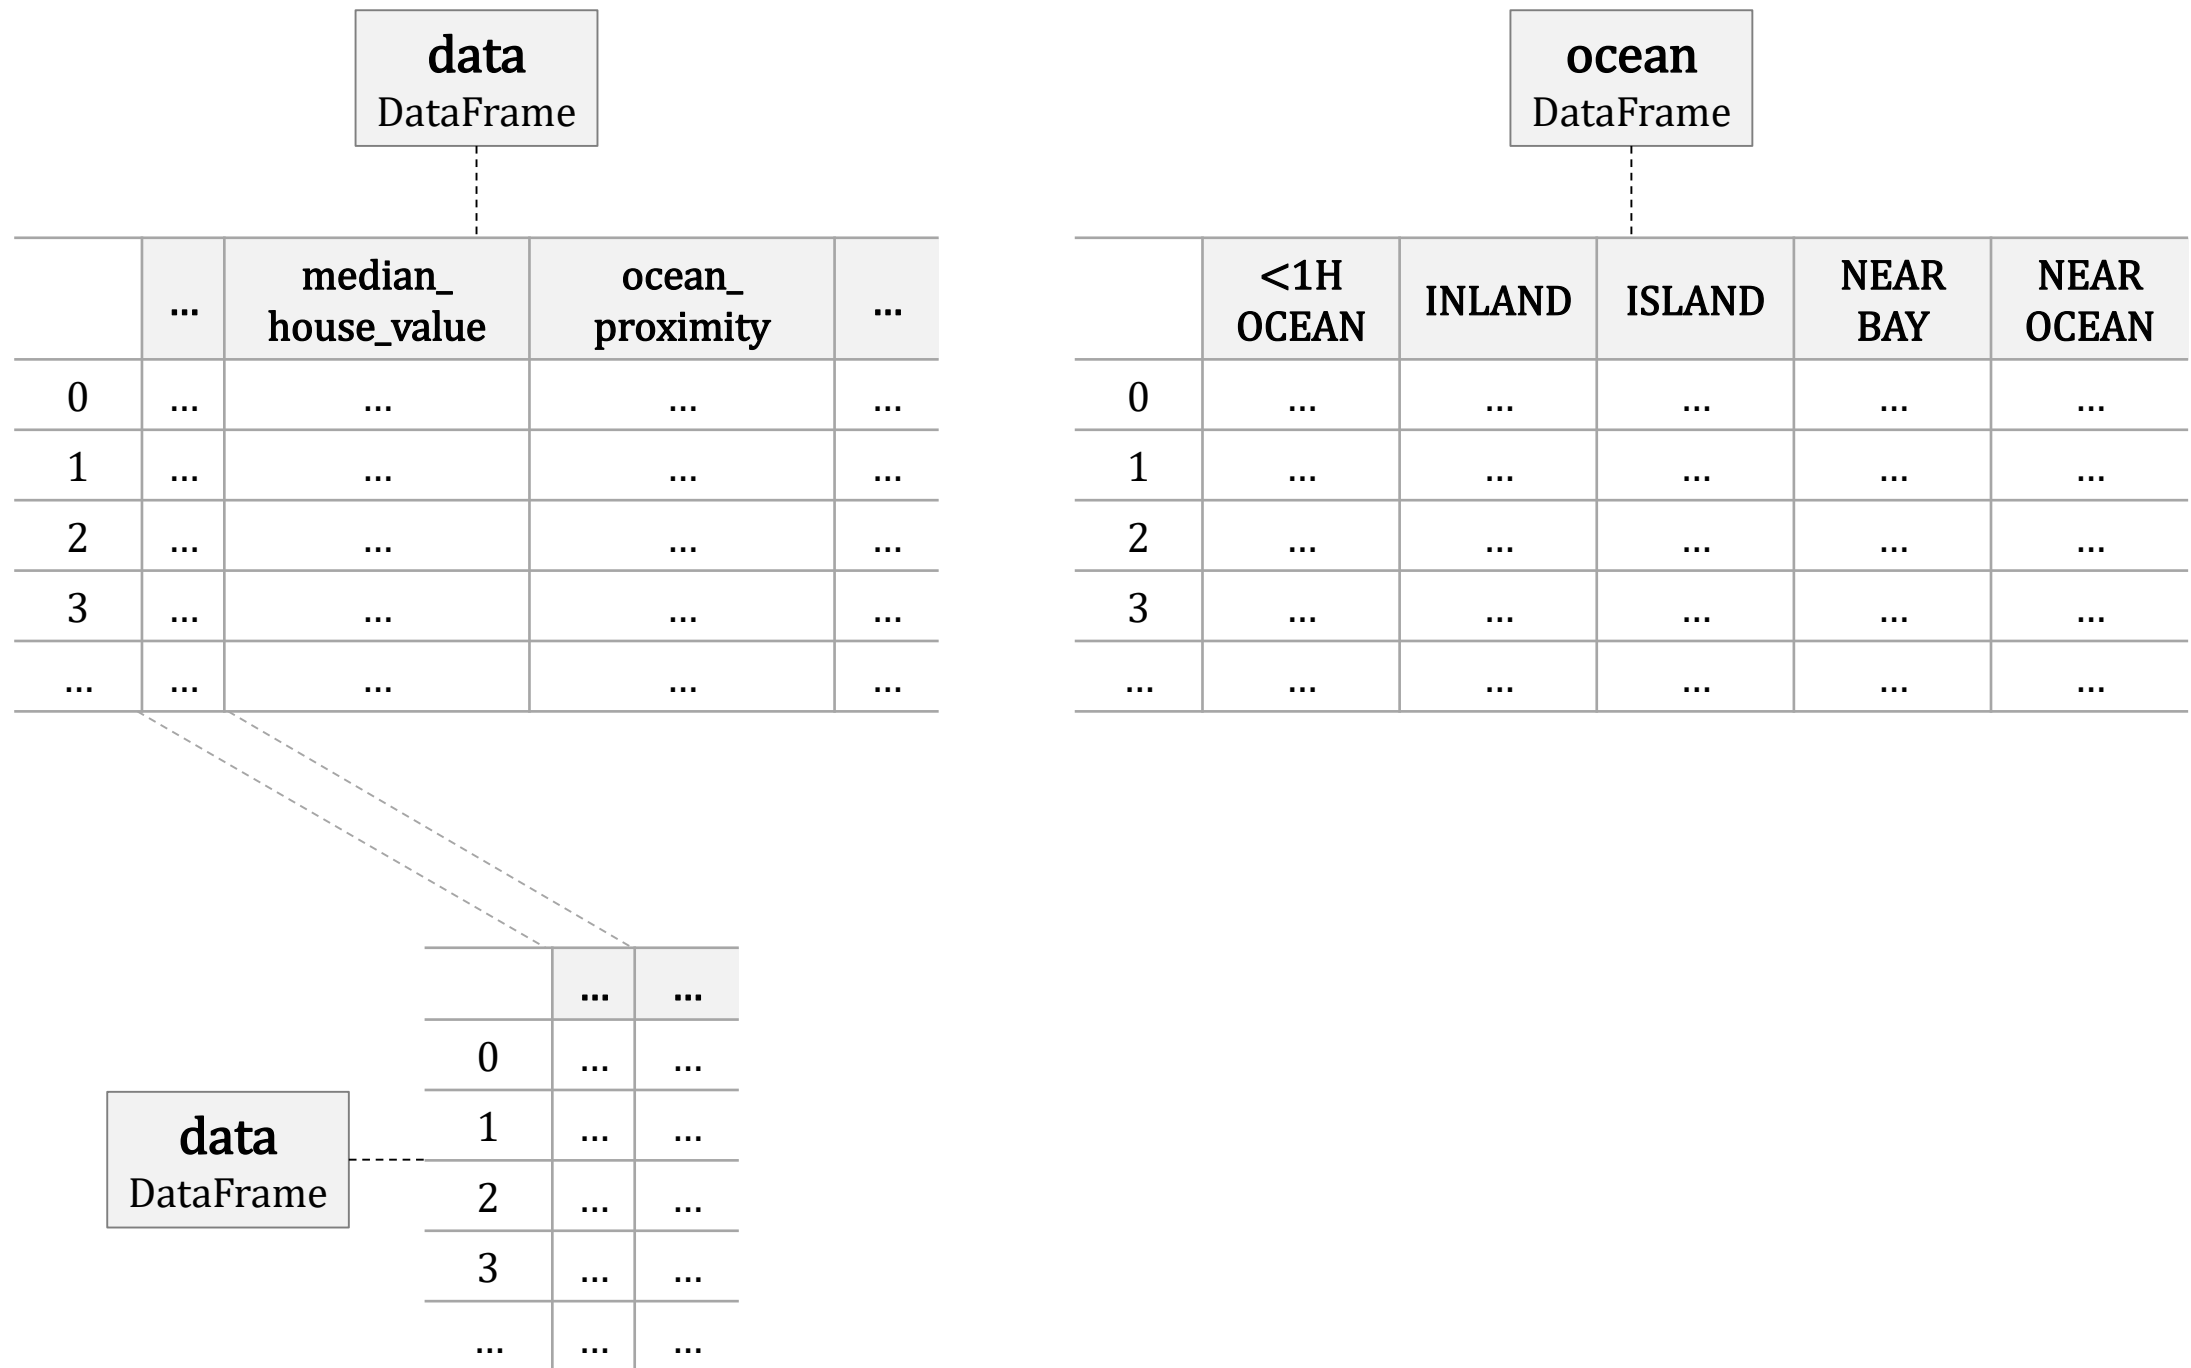
**data**  
DataFrame

	...	median_ house_value	ocean_ proximity	...
0	...	...	...	...
1	...	...	...	...
2	...	...	...	...
3	...	...	...	...
...	...	...	...	...

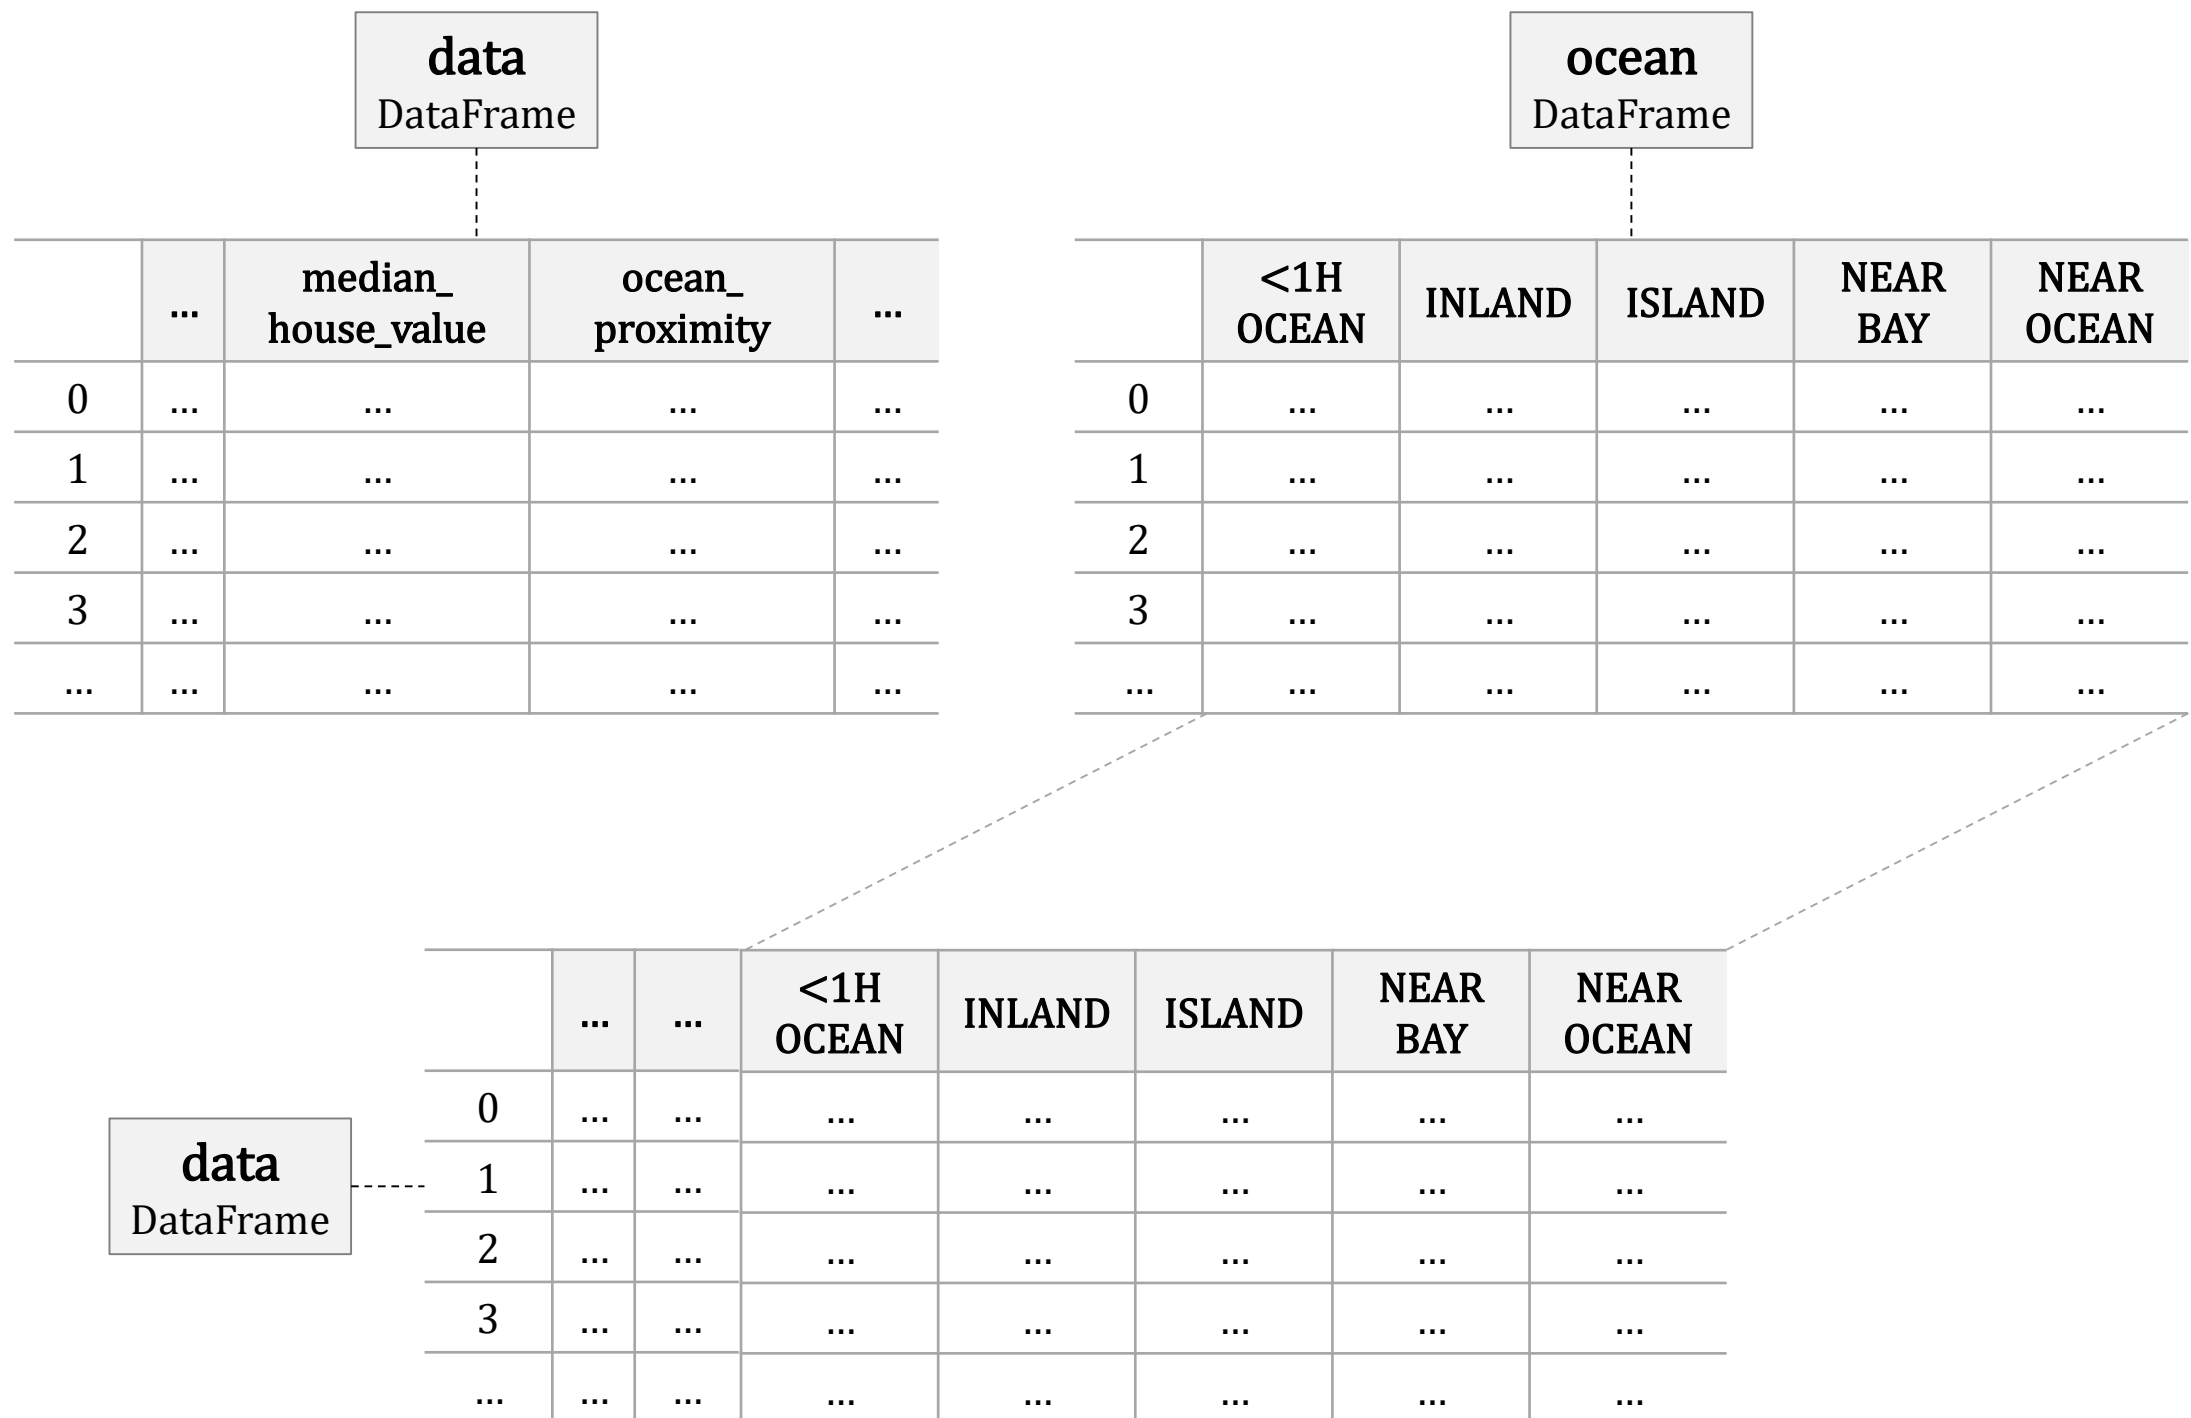
**ocean**  
DataFrame

	<1H OCEAN	INLAND	ISLAND	NEAR BAY	NEAR OCEAN
0	...	...	...	...	...
1	...	...	...	...	...
2	...	...	...	...	...
3	...	...	...	...	...
...	...	...	...	...	...

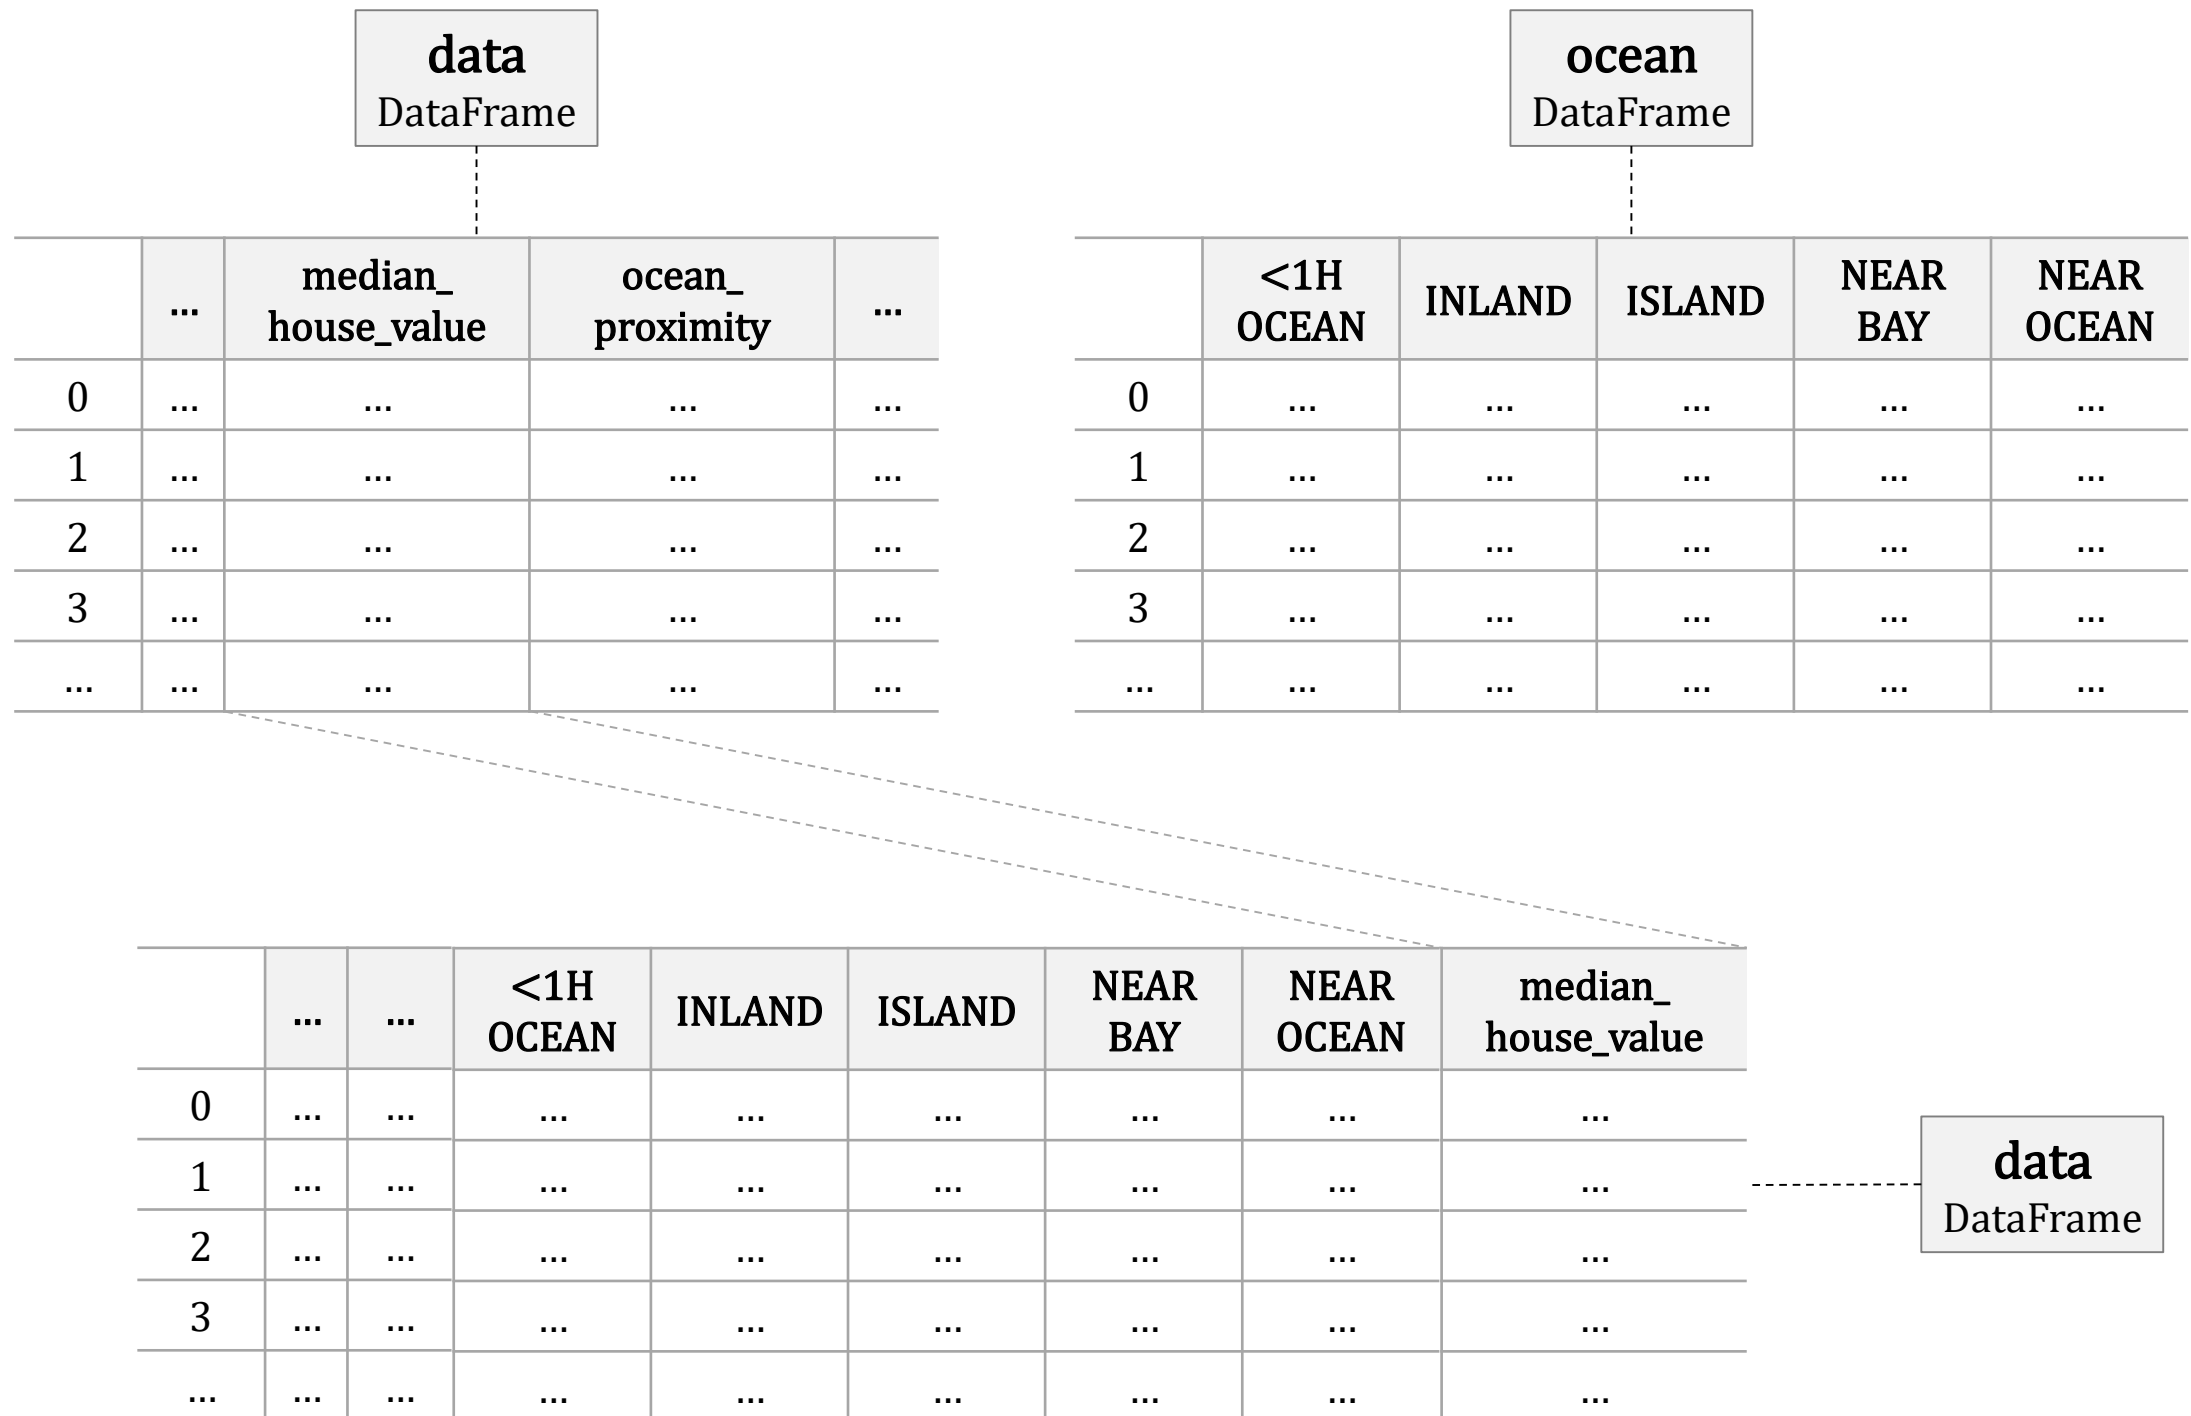
## 3.1 Handling Categorical Attributes (Feature values)



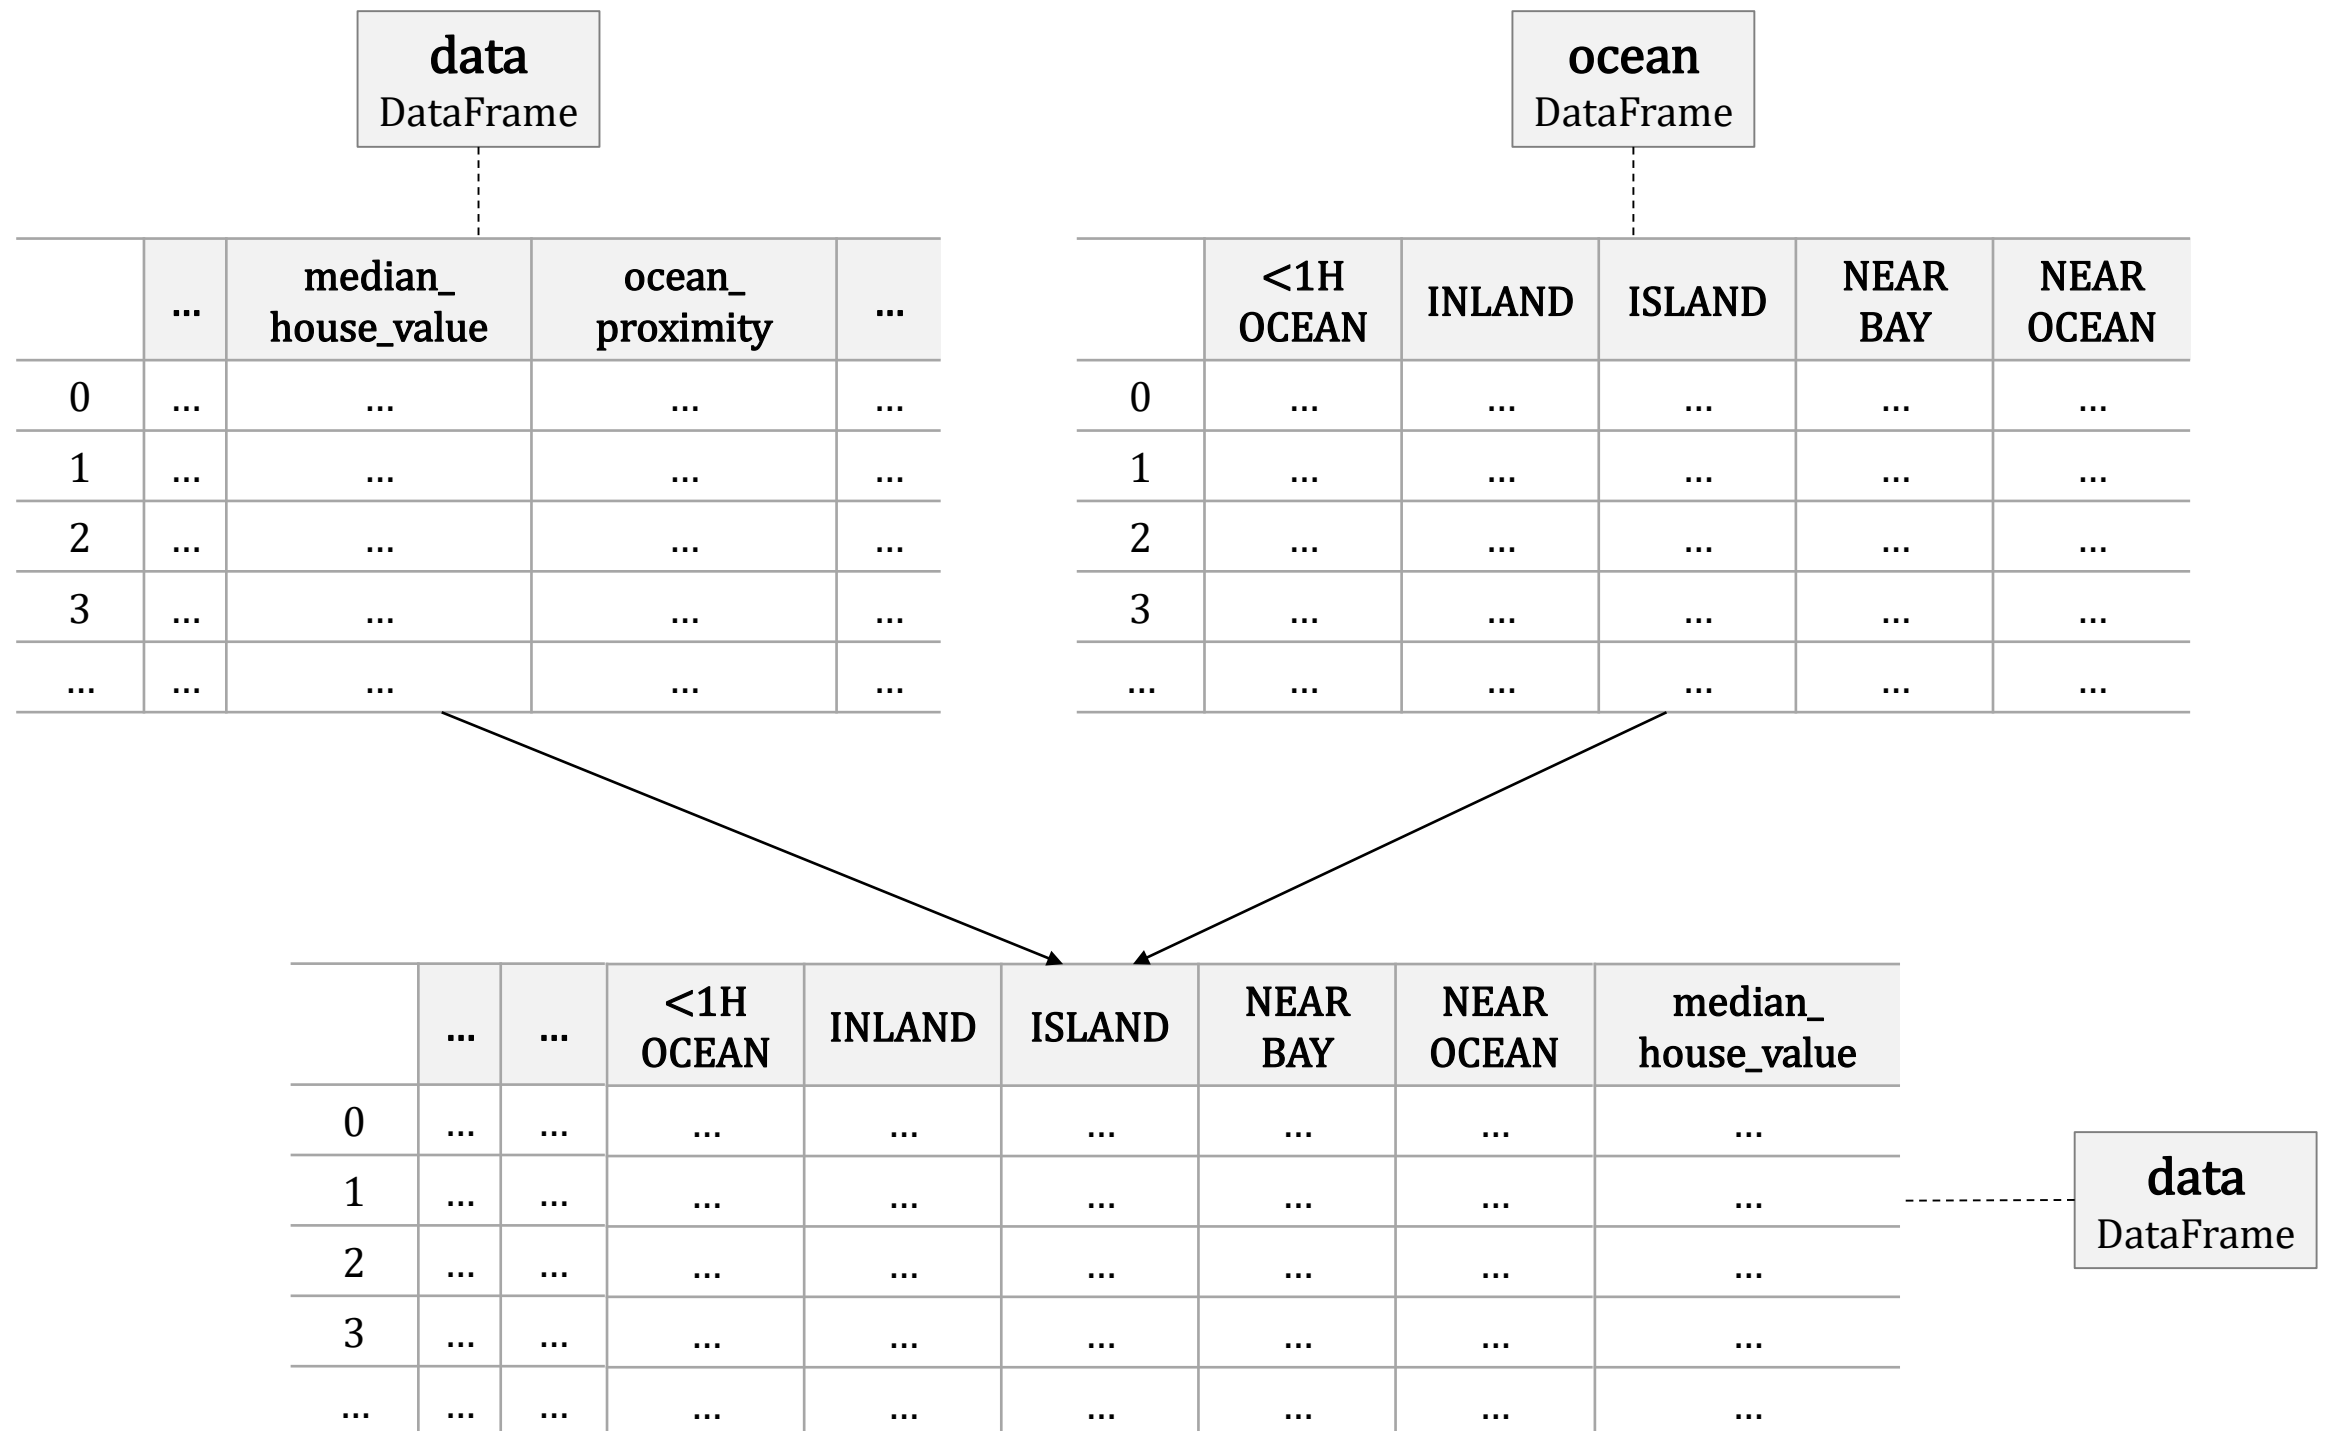
### 3.1 Handling Categorical Attributes (Feature values)



## 3.1 Handling Categorical Attributes (Feature values)



## 3.1 Handling Categorical Attributes (Feature values)



### 3.1 Handling Categorical Attributes (Feature values)

data DataFrame					ocean DataFrame					
	...	median_ house_value	ocean_ proximity	...		<1H OCEAN	INLAND	ISLAND	NEAR BAY	NEAR OCEAN
0	...	...	...	...	0	...	...	...	...	...
1	...	...	...	...	1	...	...	...	...	...
2	...	...	...	...	2	...	...	...	...	...
3	...	...	...	...	3	...	...	...	...	...
...	...	...	...	...	...	...	...	...	...	...

Slice → Drop → Drop → Concatenate → Concatenate

	...	...	<1H OCEAN	INLAND	ISLAND	NEAR BAY	NEAR OCEAN	median_ house_value	data DataFrame
0	...	...	...	...	...	...	...	...	
1	...	...	...	...	...	...	...	...	
2	...	...	...	...	...	...	...	...	
3	...	...	...	...	...	...	...	...	
...	...	...	...	...	...	...	...	...	



Slice → Drop → Drop → Concatenate → Concatenate

```
value = data[["median_house_value"]]  
value
```

median_house_value	
0	452600.0
1	358500.0
2	352100.0
3	341300.0
4	342200.0
...	...
20635	78100.0
20636	77100.0
20637	92300.0
20638	84700.0
20639	89400.0

Slice → Drop → Drop → Concatenate → Concatenate

```
data.drop(["median_house_value", "ocean_proximity"], axis=1, inplace=True)  
data
```


	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population
0	-122.23	37.88	41.0	880.0	129.0	322.0
1	-122.22	37.86	21.0	7099.0	1106.0	2401.0
2	-122.24	37.85	52.0	1467.0	190.0	496.0
3	-122.25	37.85	52.0	1274.0	235.0	558.0
4	-122.25	37.85	52.0	1627.0	280.0	565.0
...	...	...	...	...	...	...
20635	-121.09	39.48	25.0	1665.0	374.0	845.0
20636	-121.21	39.49	18.0	697.0	150.0	356.0
20637	-121.22	39.43	17.0	2254.0	485.0	1007.0
20638	-121.32	39.43	18.0	1860.0	409.0	741.0
20639	-121.24	39.37	16.0	2785.0	616.0	1387.0

Slice → Drop → Drop → Concatenate → Concatenate


```
data = pd.concat([data, ocean, value], axis=1)
data
```

bedrooms_per_room	<1H OCEAN	INLAND	ISLAND	NEAR BAY	NEAR OCEAN	median_house_value
0.146591	0.0	0.0	0.0	1.0	0.0	452600.0
0.155797	0.0	0.0	0.0	1.0	0.0	358500.0
0.129516	0.0	0.0	0.0	1.0	0.0	352100.0
0.184458	0.0	0.0	0.0	1.0	0.0	341300.0
0.172096	0.0	0.0	0.0	1.0	0.0	342200.0
...	...	...	...	...	...	...
0.224625	0.0	1.0	0.0	0.0	0.0	78100.0
0.215208	0.0	1.0	0.0	0.0	0.0	77100.0
0.215173	0.0	1.0	0.0	0.0	0.0	92300.0
0.219892	0.0	1.0	0.0	0.0	0.0	84700.0
0.221185	0.0	1.0	0.0	0.0	0.0	89400.0

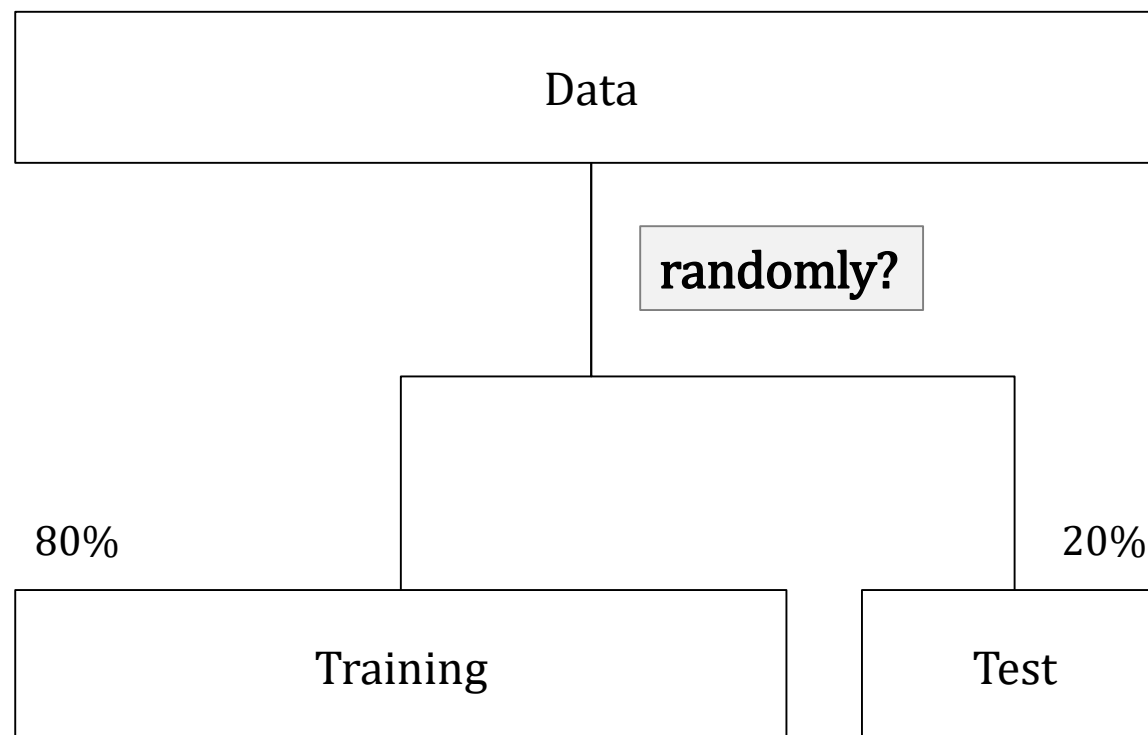
## 3.2 Create training and test set




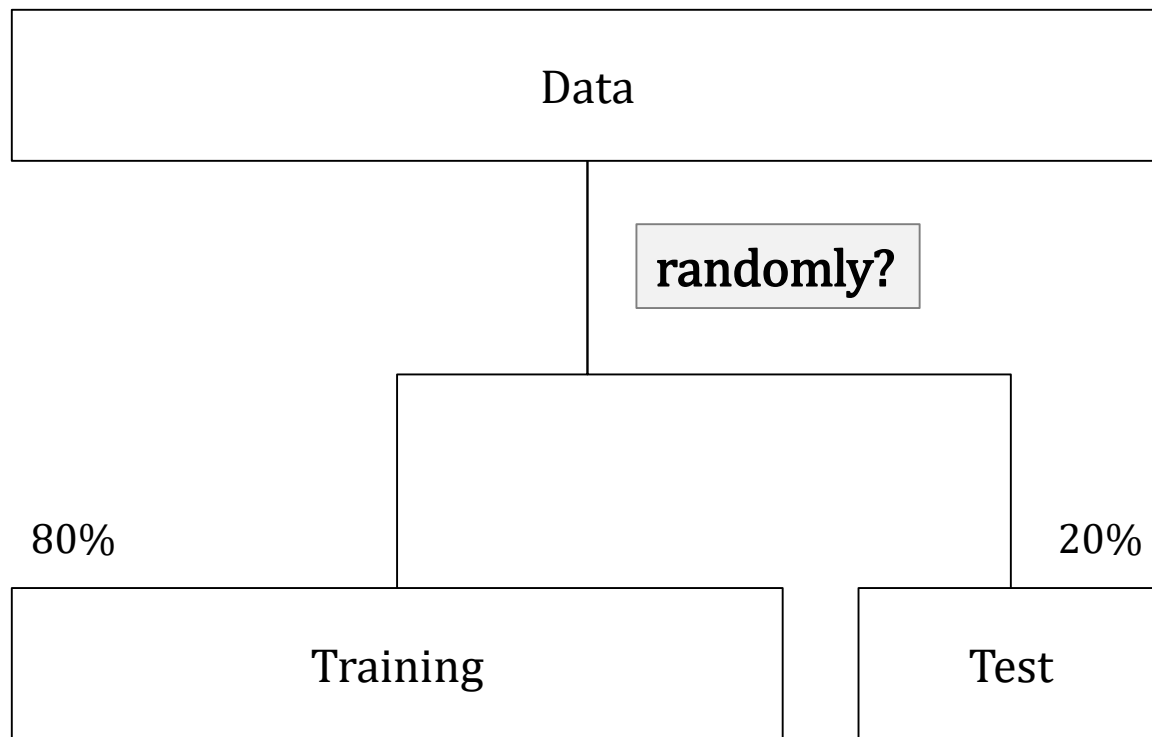
## 3.2 Create training and test set



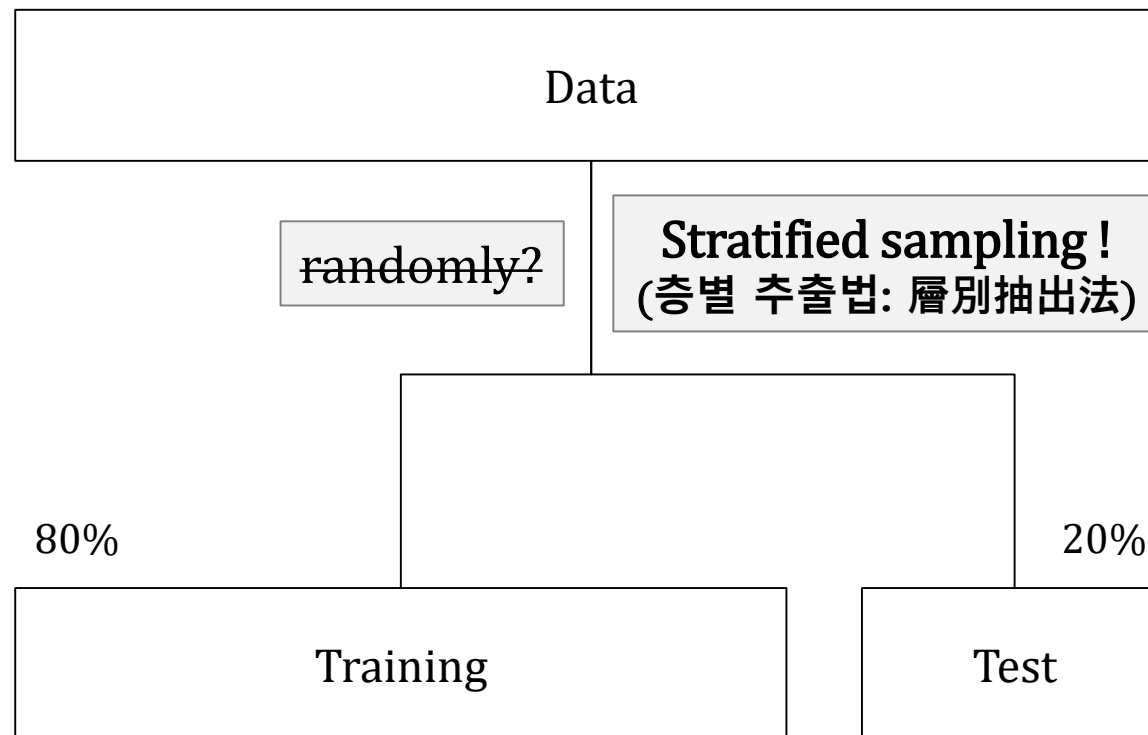
## 3.2 Create training and test set



## 3.2 Create training and test set



## 3.2 Create training and test set





## 3.2 Create training and test set


bedrooms_per_room	<1H OCEAN	INLAND	ISLAND	NEAR BAY	NEAR OCEAN	median_house_value
0.146591	0.0	0.0	0.0	1.0	0.0	452600.0
0.155797	0.0	0.0	0.0	1.0	0.0	358500.0
0.129516	0.0	0.0	0.0	1.0	0.0	352100.0
0.184458	0.0	0.0	0.0	1.0	0.0	341300.0
0.172096	0.0	0.0	0.0	1.0	0.0	342200.0
...	...	...	...	...	...	...
0.224625	0.0	1.0	0.0	0.0	0.0	78100.0
0.215208	0.0	1.0	0.0	0.0	0.0	77100.0
0.215173	0.0	1.0	0.0	0.0	0.0	92300.0
0.219892	0.0	1.0	0.0	0.0	0.0	84700.0
0.221185	0.0	1.0	0.0	0.0	0.0	89400.0

**Continuous**  
(not categorical)

## 3.2 Create training and test set

```
data["median_house_value"].plot(kind="hist", bins=50)
```


```
<matplotlib.axes._subplots.AxesSubplot at 0x29769c66608>
```



## 3.2 Create training and test set

```
data["median_house_value"].plot(kind="hist", bins=50)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x29769c66608>
```

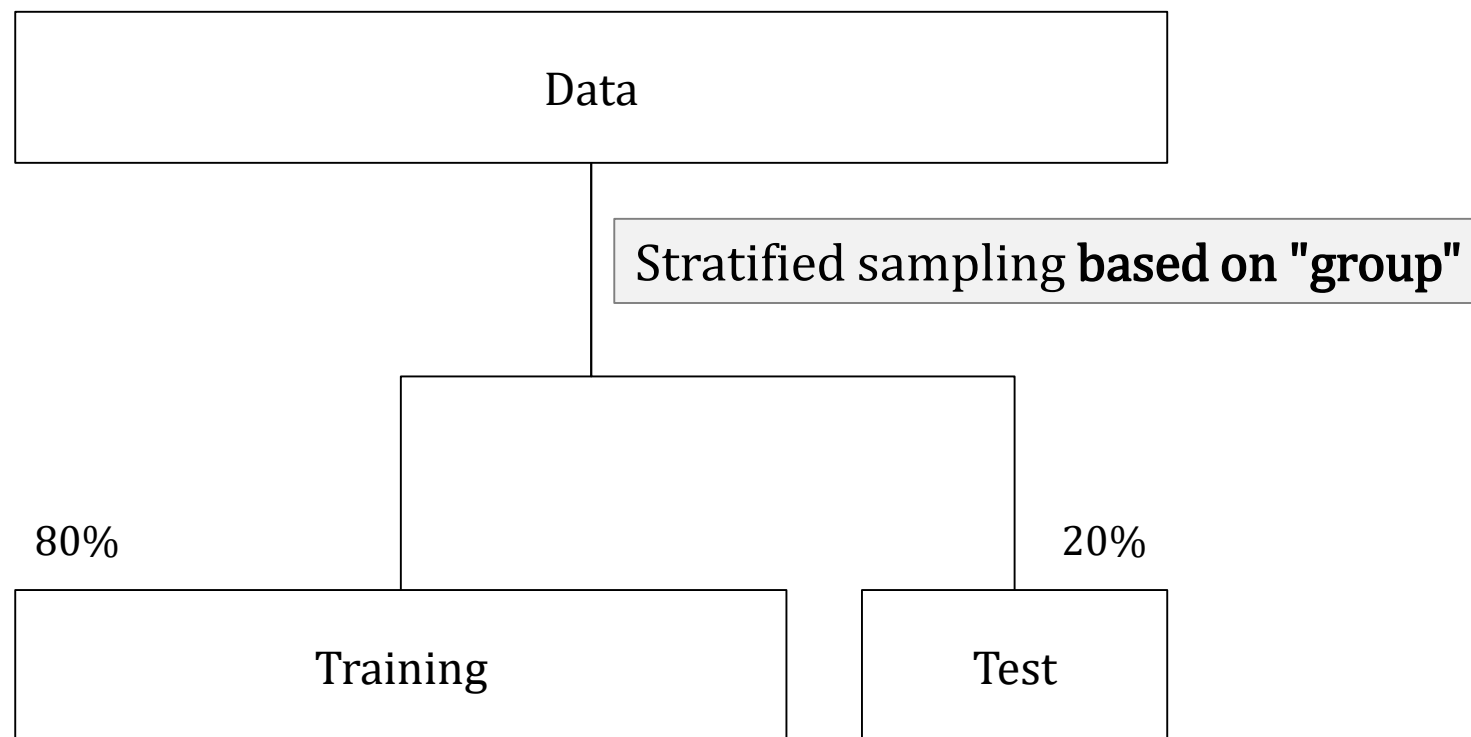


## 3.2 Create training and test set

```
data["group"] = pd.cut(x=data["median_house_value"],  
                       bins=[0, 100000, 200000, 300000, 400000, 500001],  
                       labels=["group_1", "group_2", "group_3", "group_4", "group_5"])  
data
```

sehold	bedrooms_per_room	<1H OCEAN	INLAND	ISLAND	NEAR BAY	NEAR OCEAN	median_house_value	group
384127	0.146591	0.0	0.0	0.0	1.0	0.0	452600.0	group_5
238137	0.155797	0.0	0.0	0.0	1.0	0.0	358500.0	group_4
288136	0.129516	0.0	0.0	0.0	1.0	0.0	352100.0	group_4
317352	0.184458	0.0	0.0	0.0	1.0	0.0	341300.0	group_4
281853	0.172096	0.0	0.0	0.0	1.0	0.0	342200.0	group_4
...	...	...	...	...	...	...	...	...
345455	0.224625	0.0	1.0	0.0	0.0	0.0	78100.0	group_1
114035	0.215208	0.0	1.0	0.0	0.0	0.0	77100.0	group_1
205543	0.215173	0.0	1.0	0.0	0.0	0.0	92300.0	group_1
329513	0.219892	0.0	1.0	0.0	0.0	0.0	84700.0	group_1
254717	0.221185	0.0	1.0	0.0	0.0	0.0	89400.0	group_1

## 3.2 Create training and test set



## 3.2 Create training and test set

```
from sklearn.model_selection import train_test_split

train, test = train_test_split(data,
                               test_size=0.2,
                               stratify=data["group"],
                               random_state=0)
```

```
train["group"].value_counts() / len(train)
```

```
group_2    0.400799
group_3    0.236131
group_1    0.177204
group_4    0.101381
group_5    0.084484
Name: group, dtype: float64
```

```
test["group"].value_counts() / len(test)
```

```
group_2    0.400921
group_3    0.235950
group_1    0.177326
group_4    0.101260
group_5    0.084545
Name: group, dtype: float64
```

## 3.2 Create training and test set

- We don't need the "group" feature any longer → Drop it

```
train = train.drop("group", axis=1)
train
```

s_per_room	<1H OCEAN	INLAND	ISLAND	NEAR BAY	NEAR OCEAN	median_house_value
0.193681	1.0	0.0	0.0	0.0	0.0	210200.0
0.291627	1.0	0.0	0.0	0.0	0.0	343800.0
0.236520	0.0	0.0	0.0	0.0	1.0	159400.0
0.186665	0.0	1.0	0.0	0.0	0.0	145000.0
0.222914	0.0	1.0	0.0	0.0	0.0	86700.0
...	...	...	...	...	...	...
0.222484	0.0	0.0	0.0	0.0	1.0	333800.0
0.177977	0.0	1.0	0.0	0.0	0.0	94900.0
0.274680	1.0	0.0	0.0	0.0	0.0	312500.0
0.120746	0.0	0.0	0.0	1.0	0.0	441400.0
0.922414	1.0	0.0	0.0	0.0	0.0	112500.0

## 3.2 Create training and test set

- We don't need the "group" feature any longer → Drop it

```
test = test.drop("group", axis=1)
test
```

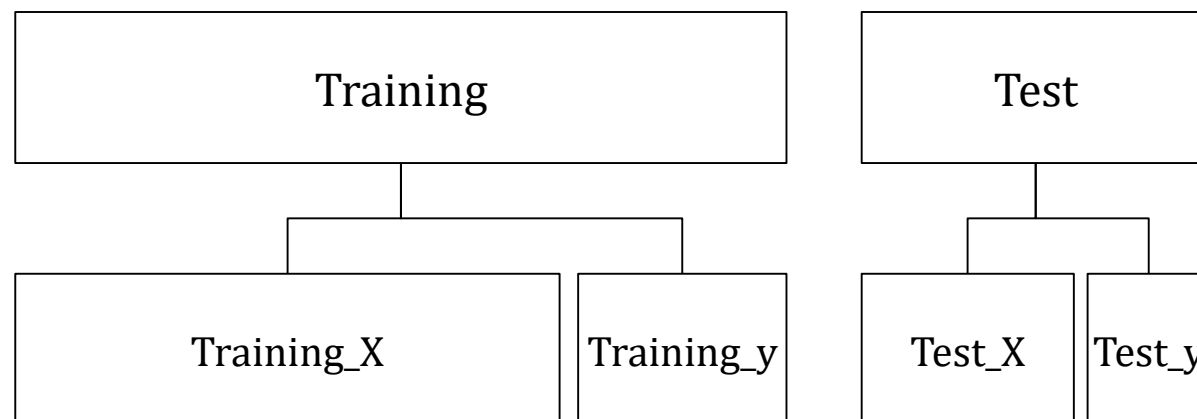
sq_per_room	<1H OCEAN	INLAND	ISLAND	NEAR BAY	NEAR OCEAN	median_house_value
0.187652	1.0	0.0	0.0	0.0	0.0	226800.0
0.196163	0.0	1.0	0.0	0.0	0.0	164000.0
NaN	1.0	0.0	0.0	0.0	0.0	122200.0
0.304348	0.0	1.0	0.0	0.0	0.0	110400.0
0.204733	0.0	1.0	0.0	0.0	0.0	150000.0
...	...	...	...	...	...	...
0.210106	0.0	1.0	0.0	0.0	0.0	39200.0
0.270597	0.0	1.0	0.0	0.0	0.0	141400.0
0.191644	0.0	1.0	0.0	0.0	0.0	87500.0
0.179072	1.0	0.0	0.0	0.0	0.0	223300.0
0.208522	0.0	1.0	0.0	0.0	0.0	106700.0



## 3.2 Create training and test set

- Common machine learning notations
  - $x$ : a vector (district) of all the feature values (excluding the label)
  - $y$ : the label (answer) of  $x$

$$x = \begin{pmatrix} -118.36 \\ 33.92 \\ 46 \\ 1,231 \\ \dots \end{pmatrix} \quad y = 226,800$$



## 3.2 Create training and test set

- For training data

```
train_y = train[["median_house_value"]]  
train_y
```

median_house_value	
6229	210200.0
5406	343800.0
14033	159400.0
2236	145000.0
1898	86700.0
...	...
17097	333800.0
12834	94900.0
8836	312500.0
1580	441400.0
4552	112500.0

```
train.drop("median_house_value", axis=1, inplace=True)  
train_X = train  
train_X
```

_household	bedrooms_per_room	<1H OCEAN	INLAND	ISLAND	NEAR BAY	NEAR OCEAN
5.495177	0.193681	1.0	0.0	0.0	0.0	0.0
3.791971	0.291627	1.0	0.0	0.0	0.0	0.0
4.388013	0.236520	0.0	0.0	0.0	0.0	1.0
6.156938	0.186665	0.0	1.0	0.0	0.0	0.0
5.474950	0.222914	0.0	1.0	0.0	0.0	0.0
...	...	...	...	...	...	...
5.005359	0.222484	0.0	0.0	0.0	0.0	1.0
5.628829	0.177977	0.0	1.0	0.0	0.0	0.0
3.879433	0.274680	1.0	0.0	0.0	0.0	0.0
8.106101	0.120746	0.0	0.0	0.0	1.0	0.0
1.260870	0.922414	1.0	0.0	0.0	0.0	0.0

## 3.2 Create training and test set

- For test data

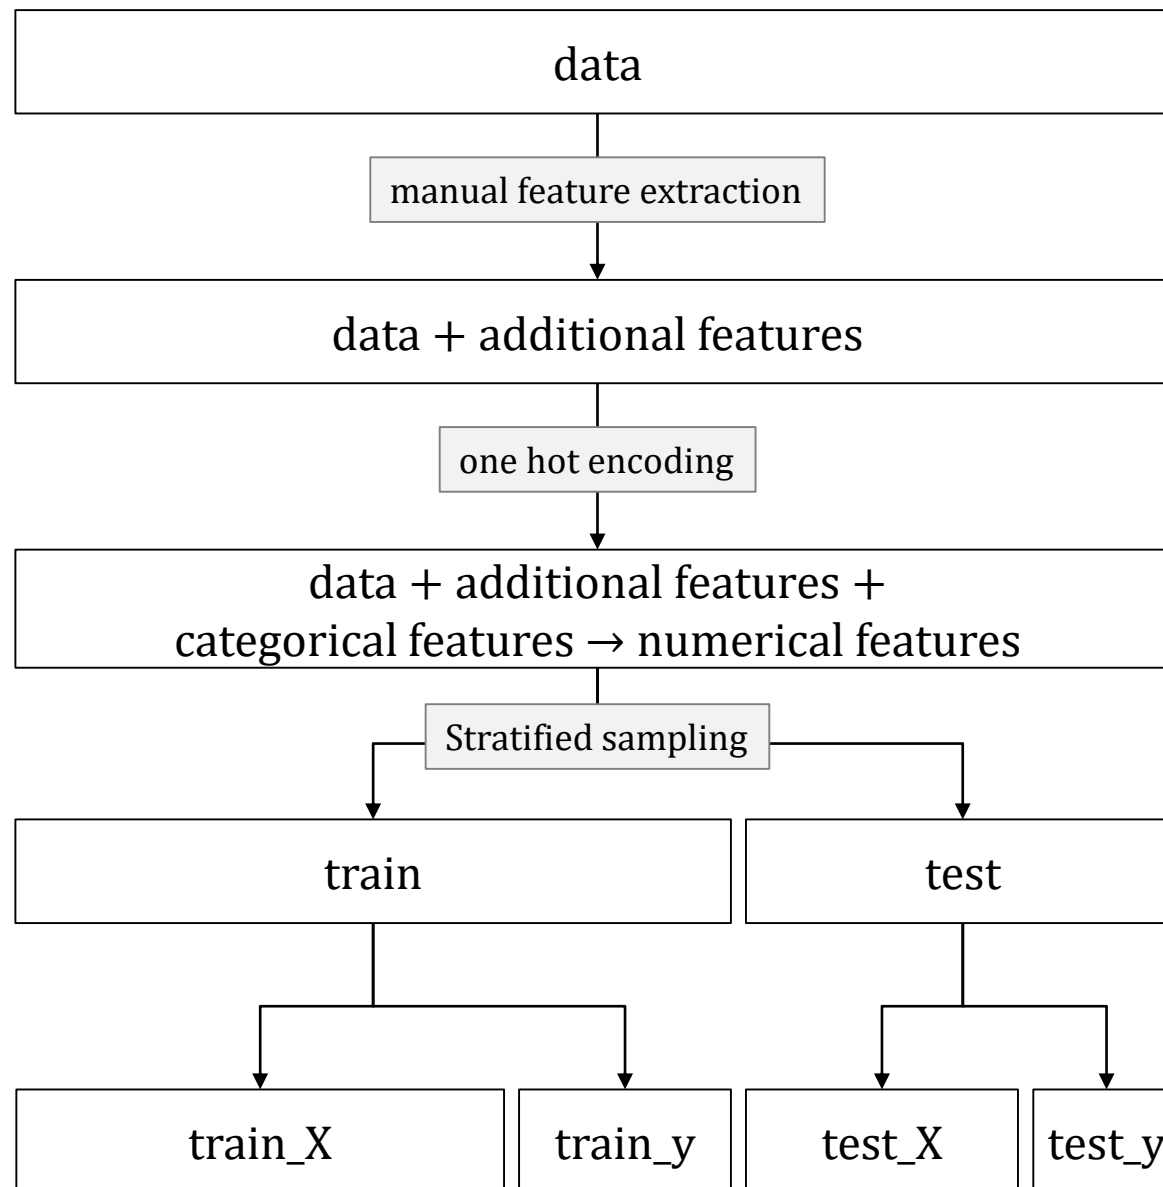
```
test_y = test[["median_house_value"]]  
test_y
```

median_house_value	
8429	226800.0
20117	164000.0
4767	122200.0
16192	110400.0
12909	150000.0
...	...
2475	39200.0
12704	141400.0
1096	87500.0
17715	223300.0
13779	106700.0

```
test.drop("median_house_value", axis=1, inplace=True)  
test_X = test  
test_X
```

bedrooms_per_room	<1H OCEAN	INLAND	ISLAND	NEAR BAY	NEAR OCEAN
0.187652	1.0	0.0	0.0	0.0	0.0
0.196163	0.0	1.0	0.0	0.0	0.0
NaN	1.0	0.0	0.0	0.0	0.0
0.304348	0.0	1.0	0.0	0.0	0.0
0.204733	0.0	1.0	0.0	0.0	0.0
...	...	...	...	...	...
0.210106	0.0	1.0	0.0	0.0	0.0
0.270597	0.0	1.0	0.0	0.0	0.0
0.191644	0.0	1.0	0.0	0.0	0.0
0.179072	1.0	0.0	0.0	0.0	0.0
0.208522	0.0	1.0	0.0	0.0	0.0

# Until now..



## 3.3 Data Cleaning

```
train_X.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 16512 entries, 6229 to 4552
```

```
Data columns (total 15 columns):
```

#	Column	Non-Null Count	Dtype
0	longitude	16512 non-null	float64
1	latitude	16512 non-null	float64
2	housing_median_age	16512 non-null	float64
3	total_rooms	16512 non-null	float64
4	total_bedrooms	16349 non-null	float64
5	population	16512 non-null	float64
6	households	16512 non-null	float64
7	median_income	16512 non-null	float64
8	rooms_per_household	16512 non-null	float64
9	bedrooms_per_room	16349 non-null	float64
10	<1H OCEAN	16512 non-null	float64
11	INLAND	16512 non-null	float64
12	ISLAND	16512 non-null	float64
13	NEAR BAY	16512 non-null	float64
14	NEAR OCEAN	16512 non-null	float64

```
dtypes: float64(15)
```

```
memory usage: 2.0 MB
```

**Missing values**  
(N/A: Not Available)

## 3.3 Data Cleaning

	...	total_bedrooms	...
0	...	...	...
1	...		...
2	...	...	...
3	...	...	...
...	...	...	...

1. Get rid of the corresponding districts

	...	total_bedrooms	...
0	...	...	...
1	...		...
2	...	...	...
3	...	...	...
...	...	...	...

2. Get rid of the whole attribute (feature)


	...	total_bedrooms	...
0	...	..	...
1	...		...
2	...	..	...
3	...	..	...
...	...	..	...

3. Set the values to some value

	...	total_bedrooms	...
0	...	...	...
1	...	some value	...
2	...	...	...
3	...	...	...
...	...	...	...

## 3.3 Data Cleaning

	...	total_bedrooms	...
0	...	...	...
1	...	zero/mean/ <b>median</b> /etc.	...
2	...	...	...
3	...	...	...
...	...	...	...



## 3.3 Data Cleaning

- For training data

```
median = train_X["total_bedrooms"].median()
train_X["total_bedrooms"].fillna(median, inplace=True)
train_X["bedrooms_per_room"].fillna(train_X["total_bedrooms"]/train_X["total_rooms"], inplace=True)
train_X.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 16512 entries, 6229 to 4552
```

```
Data columns (total 15 columns):
```

#	Column	Non-Null Count	Dtype
0	longitude	16512 non-null	float64
1	latitude	16512 non-null	float64
2	housing_median_age	16512 non-null	float64
3	total_rooms	16512 non-null	float64
4	total_bedrooms	16512 non-null	float64
5	population	16512 non-null	float64
6	households	16512 non-null	float64
7	median_income	16512 non-null	float64
8	rooms_per_household	16512 non-null	float64
9	bedrooms_per_room	16512 non-null	float64
10	<1H OCEAN	16512 non-null	float64
11	INLAND	16512 non-null	float64
12	ISLAND	16512 non-null	float64
13	NEAR BAY	16512 non-null	float64
14	NEAR OCEAN	16512 non-null	float64

```
dtypes: float64(15)
```



## 3.3 Data Cleaning

- For test data

```
test_X["total_bedrooms"].fillna(median, inplace=True)
test_X["bedrooms_per_room"].fillna(test_X["total_bedrooms"]/test_X["total_rooms"], inplace=True)
test_X.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 4128 entries, 8429 to 13779
Data columns (total 15 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   longitude             4128 non-null   float64
 1   latitude              4128 non-null   float64
 2   housing_median_age    4128 non-null   float64
 3   total_rooms            4128 non-null   float64
 4   total_bedrooms        4128 non-null   float64
 5   population             4128 non-null   float64
 6   households             4128 non-null   float64
 7   median_income          4128 non-null   float64
 8   rooms_per_household    4128 non-null   float64
 9   bedrooms_per_room      4128 non-null   float64
10   <1H OCEAN              4128 non-null   float64
11   INLAND                 4128 non-null   float64
12   ISLAND                 4128 non-null   float64
13   NEAR BAY               4128 non-null   float64
14   NEAR OCEAN             4128 non-null   float64
dtypes: float64(15)
```

# End-to-End Machine Learning Project



	No.	Action	Package/library
Data preprocessing for machine learning	0	Look at the big picture	—
	1	Get the data	tarfile, urllib, pandas
	2	Discover and visualize the data to gain insights	pandas, matplotlib
	3	Prepare the data for Machine Learning algorithms	pandas, scikit-learn, numpy
	4	Select a model and train it	
	5	Fine-tune your model	
	6	Present your solution	
	7	Launch, monitor and maintain your system	joblib, flask

# End-to-End Machine Learning Project

No.	Action	Package/library
0	Look at the big picture	—
1	Get the data	tarfile, urllib, pandas
2	Discover and visualize the data to gain insights	pandas, matplotlib
3	Prepare the data for Machine Learning algorithms	pandas, scikit-learn, numpy
4	Select a model and train it	
5	Fine-tune your model	
6	Present your solution	
7	Launch, monitor and maintain your system	joblib, flask

Three ML methods we gonna use today..

**Linear  
regression**

**Decision tree  
regression**

**Random forest  
regression**

## 4.1 Training and Evaluating on the Training Set

Linear  
regression

```
from sklearn.linear_model import LinearRegression

linear = LinearRegression()
linear.fit(train_X, train_y)
```

```
LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

```
from sklearn.metrics import mean_squared_error
import numpy as np

predictions = linear.predict(train_X)
mse = mean_squared_error(train_y, predictions)
rmse = np.sqrt(mse)
answer_mean = train_y["median_house_value"].mean()
print(str(rmse/answer_mean*100) + "%")
```

```
32.92999142716399%
```

## 4.1 Training and Evaluating on the Training Set

Decision tree  
regression

```
from sklearn.tree import DecisionTreeRegressor
```

```
tree = DecisionTreeRegressor()  
tree.fit(train_X, train_y)
```

```
DecisionTreeRegressor(ccp_alpha=0.0, criterion='mse', m  
ax_depth=None,  
max_features=None, max_leaf_nodes  
=None,  
min_impurity_decrease=0.0, min_im  
purity_split=None,  
min_samples_leaf=1, min_samples_s  
plit=2,  
min_weight_fraction_leaf=0.0, pre  
sort='deprecated',  
random_state=None, splitter='bes  
t')
```

```
predictions = tree.predict(train_X)  
mse = mean_squared_error(train_y, predictions)  
rmse = np.sqrt(mse)  
answer_mean = train_y["median_house_value"].mean()  
print(str(rmse/answer_mean*100) + "%")
```

0.0%

## 4.1 Training and Evaluating on the Training Set

**Linear  
regression**

**Under fitting**

**Decision tree  
regression**

**Over fitting**

**Random forest  
regression**

## 4.2 Evaluating on the Test Set

```
predictions = linear.predict(test_X)
mse = mean_squared_error(test_y, predictions)
rmse = np.sqrt(mse)
answer_mean = test_y["median_house_value"].mean()
print(str(rmse/answer_mean*100) + "%")
```

33.387643759619735%

```
predictions = tree.predict(test_X)
mse = mean_squared_error(test_y, predictions)
rmse = np.sqrt(mse)
answer_mean = test_y["median_house_value"].mean()
print(str(rmse/answer_mean*100) + "%")
```

33.33372756965135%



## 4.3 Evaluating on the Training and Test Set using Random forest regression

Linear  
regression

Decision tree  
regression

Random forest  
regression

## 4.3 Evaluating on the Training and Test Set using Random forest regression

```
from sklearn.ensemble import RandomForestRegressor
```

```
forest = RandomForestRegressor(n_estimators=5, random_state=0)  
forest.fit(train_X, train_y["median_house_value"].ravel())
```

```
RandomForestRegressor(bootstrap=True, ccp_alpha=0.0, criterion='mse',  
                       max_depth=None, max_features='auto', max_leaf_node  
s=None,  
                       max_samples=None, min_impurity_decrease=0.0,  
                       min_impurity_split=None, min_samples_leaf=1,  
                       min_samples_split=2, min_weight_fraction_leaf=0.0,  
                       n_estimators=5, n_jobs=None, oob_score=False,  
                       random_state=0, verbose=0, warm_start=False)
```

## 4.3 Evaluating on the Training and Test Set using Random forest regression

```
predictions = forest.predict(train_X)
mse = mean_squared_error(train_y, predictions)
rmse = np.sqrt(mse)
answer_mean = train_y["median_house_value"].mean()
print(str(rmse/answer_mean*100) + "%")
```

12.12961140719765%

```
predictions = forest.predict(test_X)
mse = mean_squared_error(test_y, predictions)
rmse = np.sqrt(mse)
answer_mean = test_y["median_house_value"].mean()
print(str(rmse/answer_mean*100) + "%")
```

26.17531878445498%


# Comparison

	Linear regression	Decision tree regression	Random forest regression
RMSE on training set	32.9%	0%	12.1%
RMSE on test set	33.3%	33.3%	26.1%

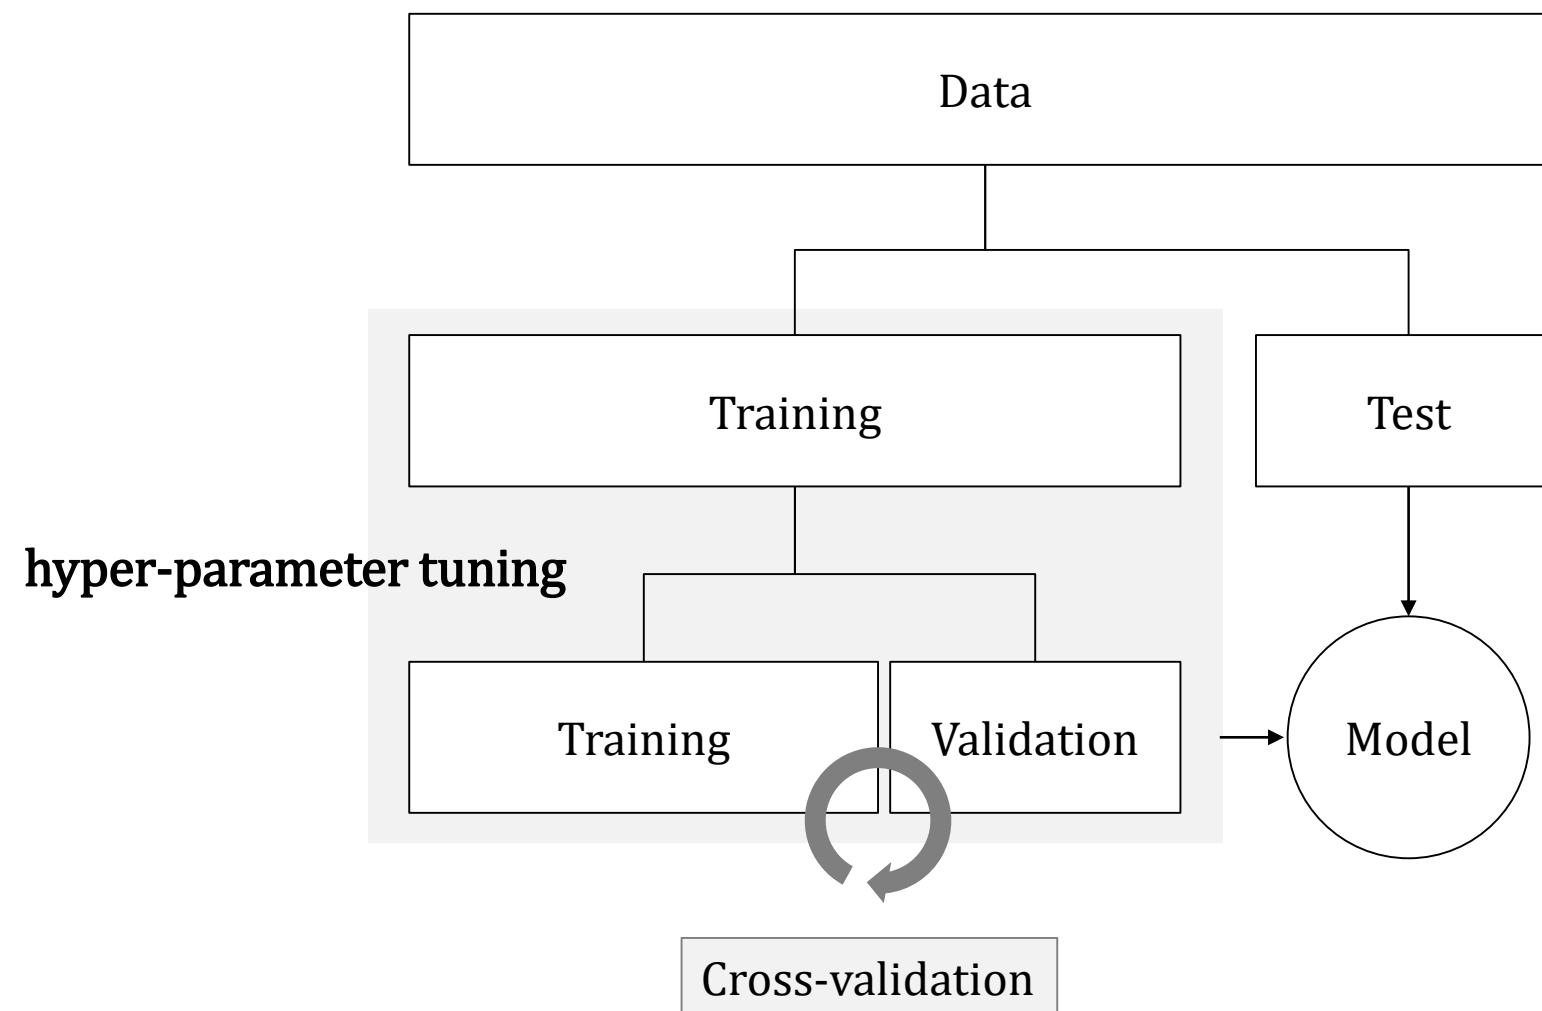
# End-to-End Machine Learning Project

No.	Action	Package/library
0	Look at the big picture	—
1	Get the data	tarfile, urllib, pandas
2	Discover and visualize the data to gain insights	pandas, matplotlib
3	Prepare the data for Machine Learning algorithms	pandas, scikit-learn, numpy
4	Select a model and train it	
5	Fine-tune your model	
6	Present your solution	
7	Launch, monitor and maintain your system	joblib, flask

# The ideal scenario



# The ideal scenario



## Grid Search with Cross Validation

```
from sklearn.model_selection import GridSearchCV

param = {"n_estimators": [5, 10, 30],
         "max_depth": [100, 200]}
forest = RandomForestRegressor(random_state=0)
search = GridSearchCV(forest, param,
                      cv=5, scoring="neg_mean_squared_error")
search.fit(train_X, train_y["median_house_value"].ravel())
```



## Your fine-tuned model

```
search.best_params_
```

```
{'max_depth': 100, 'n_estimators': 30}
```

```
forest = RandomForestRegressor(max_depth=100,  
                              n_estimators=30,  
                              random_state=0)  
forest.fit(train_X, train_y["median_house_value"].ravel())
```

## Your fine-tuned model

```
predictions = forest.predict(train_X)
mse = mean_squared_error(train_y, predictions)
rmse = np.sqrt(mse)
answer_mean = train_y["median_house_value"].mean()
print(str(rmse/answer_mean*100) + "%")
```

9.421054753574424%

```
predictions = forest.predict(test_X)
mse = mean_squared_error(test_y, predictions)
rmse = np.sqrt(mse)
answer_mean = test_y["median_house_value"].mean()
print(str(rmse/answer_mean*100) + "%")
```

23.886510037964545%

## Your fine-tune model

	Linear regression	Decision tree regression	Random forest regression	Fine-tuned model
RMSE on training set	32.9%	0%	12.1%	9.4%
RMSE on test set	33.3%	33.3%	26.1%	23.8%

# End-to-End Machine Learning Project

No.	Action	Package/library
0	Look at the big picture	—
1	Get the data	tarfile, urllib, pandas
2	Discover and visualize the data to gain insights	pandas, matplotlib
3	Prepare the data for Machine Learning algorithms	pandas, scikit-learn, numpy
4	Select a model and train it	
5	Fine-tune your model	
6	Present your solution	
7	Launch, monitor and maintain your system	joblib, flask

## Feature importance

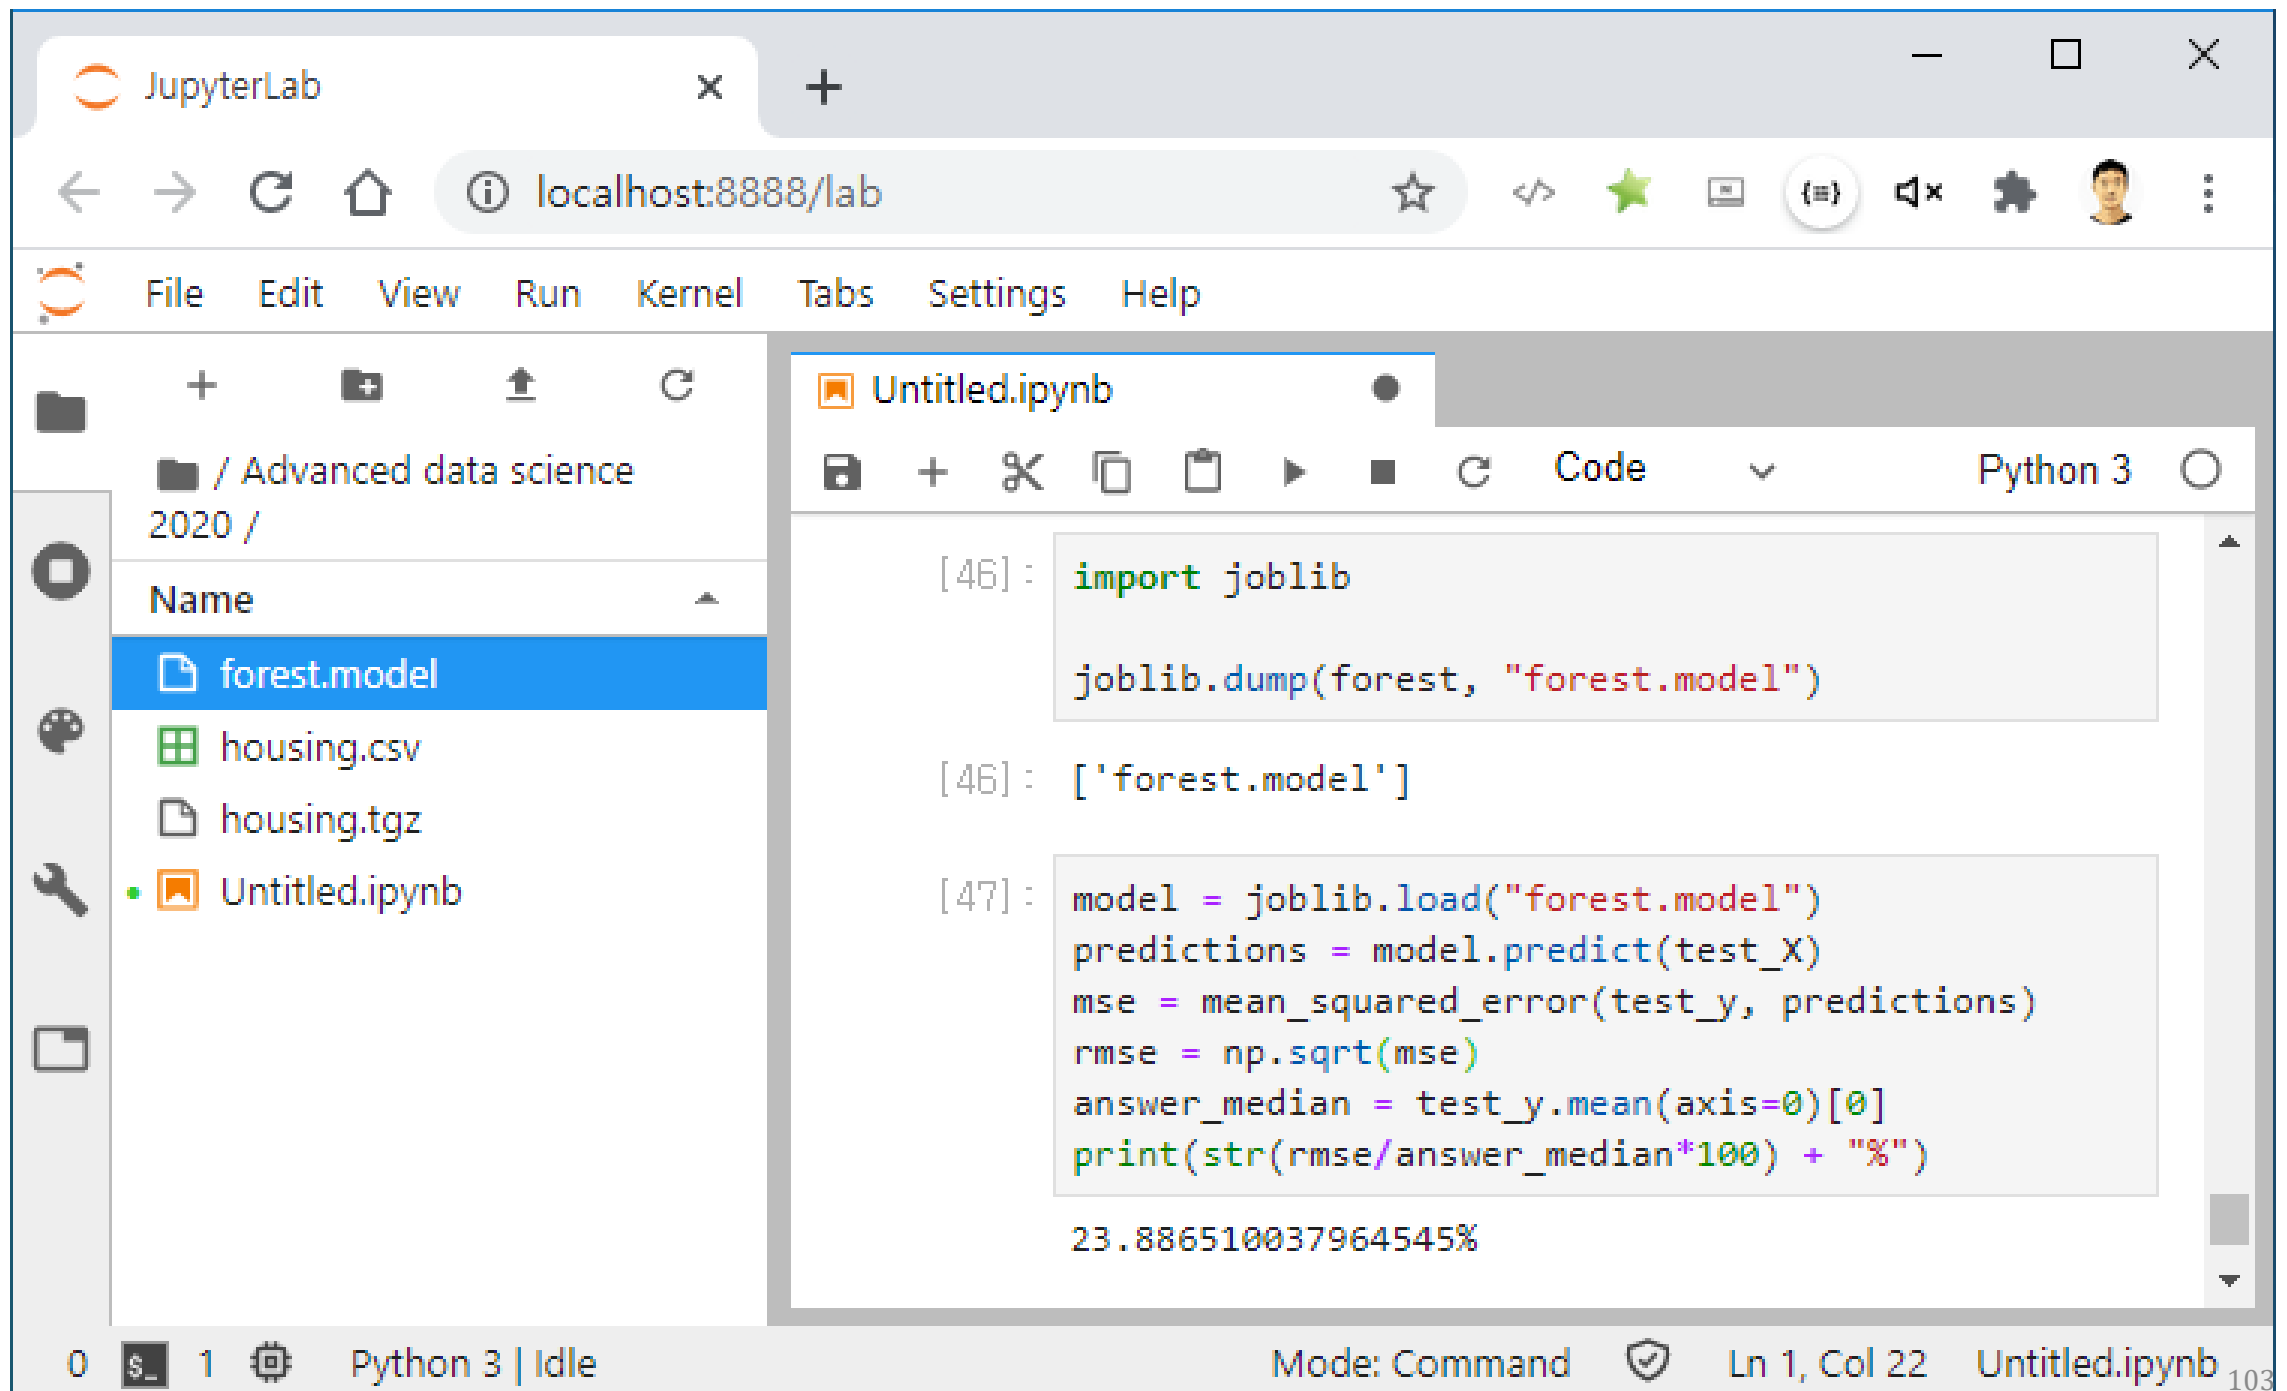
```
importances = forest.feature_importances_  
features = train_X.columns.values  
  
sorted(zip(importances, features), reverse=True)
```

```
[(0.47763306801959393, 'median_income'),  
 (0.14667781030892643, 'INLAND'),  
 (0.09675169156214641, 'longitude'),  
 (0.08728313888933595, 'latitude'),  
 (0.047588373717733196, 'housing_median_age'),  
 (0.03244868607116767, 'bedrooms_per_room'),  
 (0.028850969908527545, 'rooms_per_household'),  
 (0.024088508967163785, 'population'),  
 (0.01855807032696892, 'total_rooms'),  
 (0.015667964018780612, 'total_bedrooms'),  
 (0.01431359816365993, 'households'),  
 (0.006327717669091209, 'NEAR OCEAN'),  
 (0.0026947420610949524, '<1H OCEAN'),  
 (0.001037627384628986, 'NEAR BAY'),  
 (7.8032931180394e-05, 'ISLAND')]
```

# End-to-End Machine Learning Project

No.	Action	Package/library
0	Look at the big picture	—
1	Get the data	tarfile, urllib, pandas
2	Discover and visualize the data to gain insights	pandas, matplotlib
3	Prepare the data for Machine Learning algorithms	pandas, scikit-learn, numpy
4	Select a model and train it	
5	Fine-tune your model	
6	Present your solution	
7	Launch, monitor and maintain your system	joblib, flask

# Save & load your model



The screenshot displays the JupyterLab web interface in a browser window. The address bar shows `localhost:8888/lab`. The left sidebar contains a file explorer for the directory `/ Advanced data science 2020 /`, listing `forest.model`, `housing.csv`, `housing.tgz`, and `Untitled.ipynb`. The `forest.model` file is selected. The main area shows the `Untitled.ipynb` notebook with two code cells. The first cell (index 46) imports `joblib` and saves a `forest` model to `forest.model`. The second cell (index 47) loads the model, makes predictions on `test_X`, calculates the mean squared error (MSE), and prints the root mean squared error (RMSE) as a percentage of the median answer. The output of the second cell is `23.886510037964545%`. The bottom status bar indicates the current mode is 'Command' and the file is `Untitled.ipynb`.

JupyterLab

localhost:8888/lab

File Edit View Run Kernel Tabs Settings Help

/ Advanced data science 2020 /

Name

- forest.model
- housing.csv
- housing.tgz
- Untitled.ipynb

Untitled.ipynb

Code Python 3

```
[46]: import joblib

      joblib.dump(forest, "forest.model")

[46]: ['forest.model']


[47]: model = joblib.load("forest.model")
      predictions = model.predict(test_X)
      mse = mean_squared_error(test_y, predictions)
      rmse = np.sqrt(mse)
      answer_median = test_y.mean(axis=0)[0]
      print(str(rmse/answer_median*100) + "%")

      23.886510037964545%
```

0 Python 3 | Idle Mode: Command Ln 1, Col 22 Untitled.ipynb


## 기계 학습 모델 구축 → 많은 자원 소모


- 데이터 수집 및 분석
- 적합한 기계 학습 방법론 선택
  - Linear regression, decision tree, random forest, perceptron, neural network, etc.
- 다양한 hyper parameter 설정
  - Num. of layers, activation function, learning rate, etc.





## 노력의 산물





프로그래밍 언어의 장벽

지적 재산권(보안)

## **APIs & REST APIs**

# Application Programming Interface

# Application Programming Interface


- A way to let **software components** to **talk to each other**



API 경험이 있나요?

# API 경험이 있나요?

- Swing (Java)



JFrame
+ EXIT_ON_CLOSE: int - rootPane: JRootPane ...
+ JFrame() ... + setTitle(String): void + setSize(int, int): void + setIconImage(Image): void + setVisible(boolean): void ...

```
public static void main(String args)
{
    JFrame frame = new JFrame();
    frame.setTitle("Cacluator");
    frame.setVisible(true);
}
```

So an **API** could be anything in any form.


The only thing that it has to be is that

it has to be a **way to communicate with a software component.**




# REST API

- HTTP 프로토콜 기반의 API



# REST API

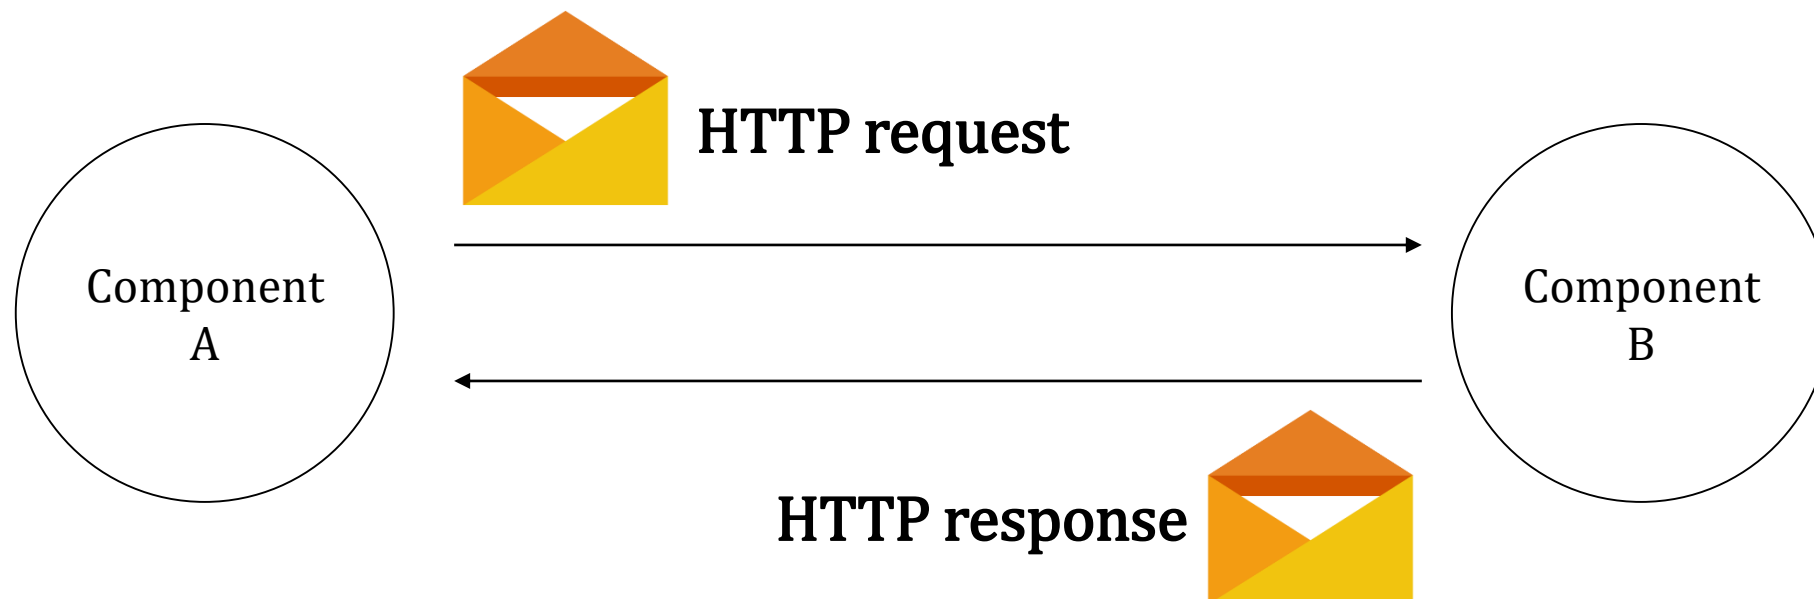
- HTTP 프로토콜 기반의 API



Client language	Server language
Python	Python
Python	Java
Java	Python
...	...

# REST API


- HTTP 프로토콜 기반의 API



CRUD	HTTP verb
Create	POST
Read	GET
Update	PUT/PATCH
Delete	DELTE


# REST API

- HTTP 프로토콜 기반의 API




CRUD	HTTP verb
Create	POST
Read	GET
Update	PUT/PATCH
Delete	DELTE

# New notebook



# New notebook



# Sample api

The screenshot shows the JupyterLab web interface in a browser window. The address bar displays `localhost:8888/lab`. The left sidebar shows a file explorer with the directory `/ Advanced data science 2020 /` containing files `forest.model`, `housing.csv`, `housing.tgz`, and two notebooks: `Untitled.ipynb` and `Untitled1.ipynb` (which is selected). The main area displays the `Untitled1.ipynb` notebook in Command mode, showing a single code cell with the following Python code:

```
[ ]: from flask import Flask

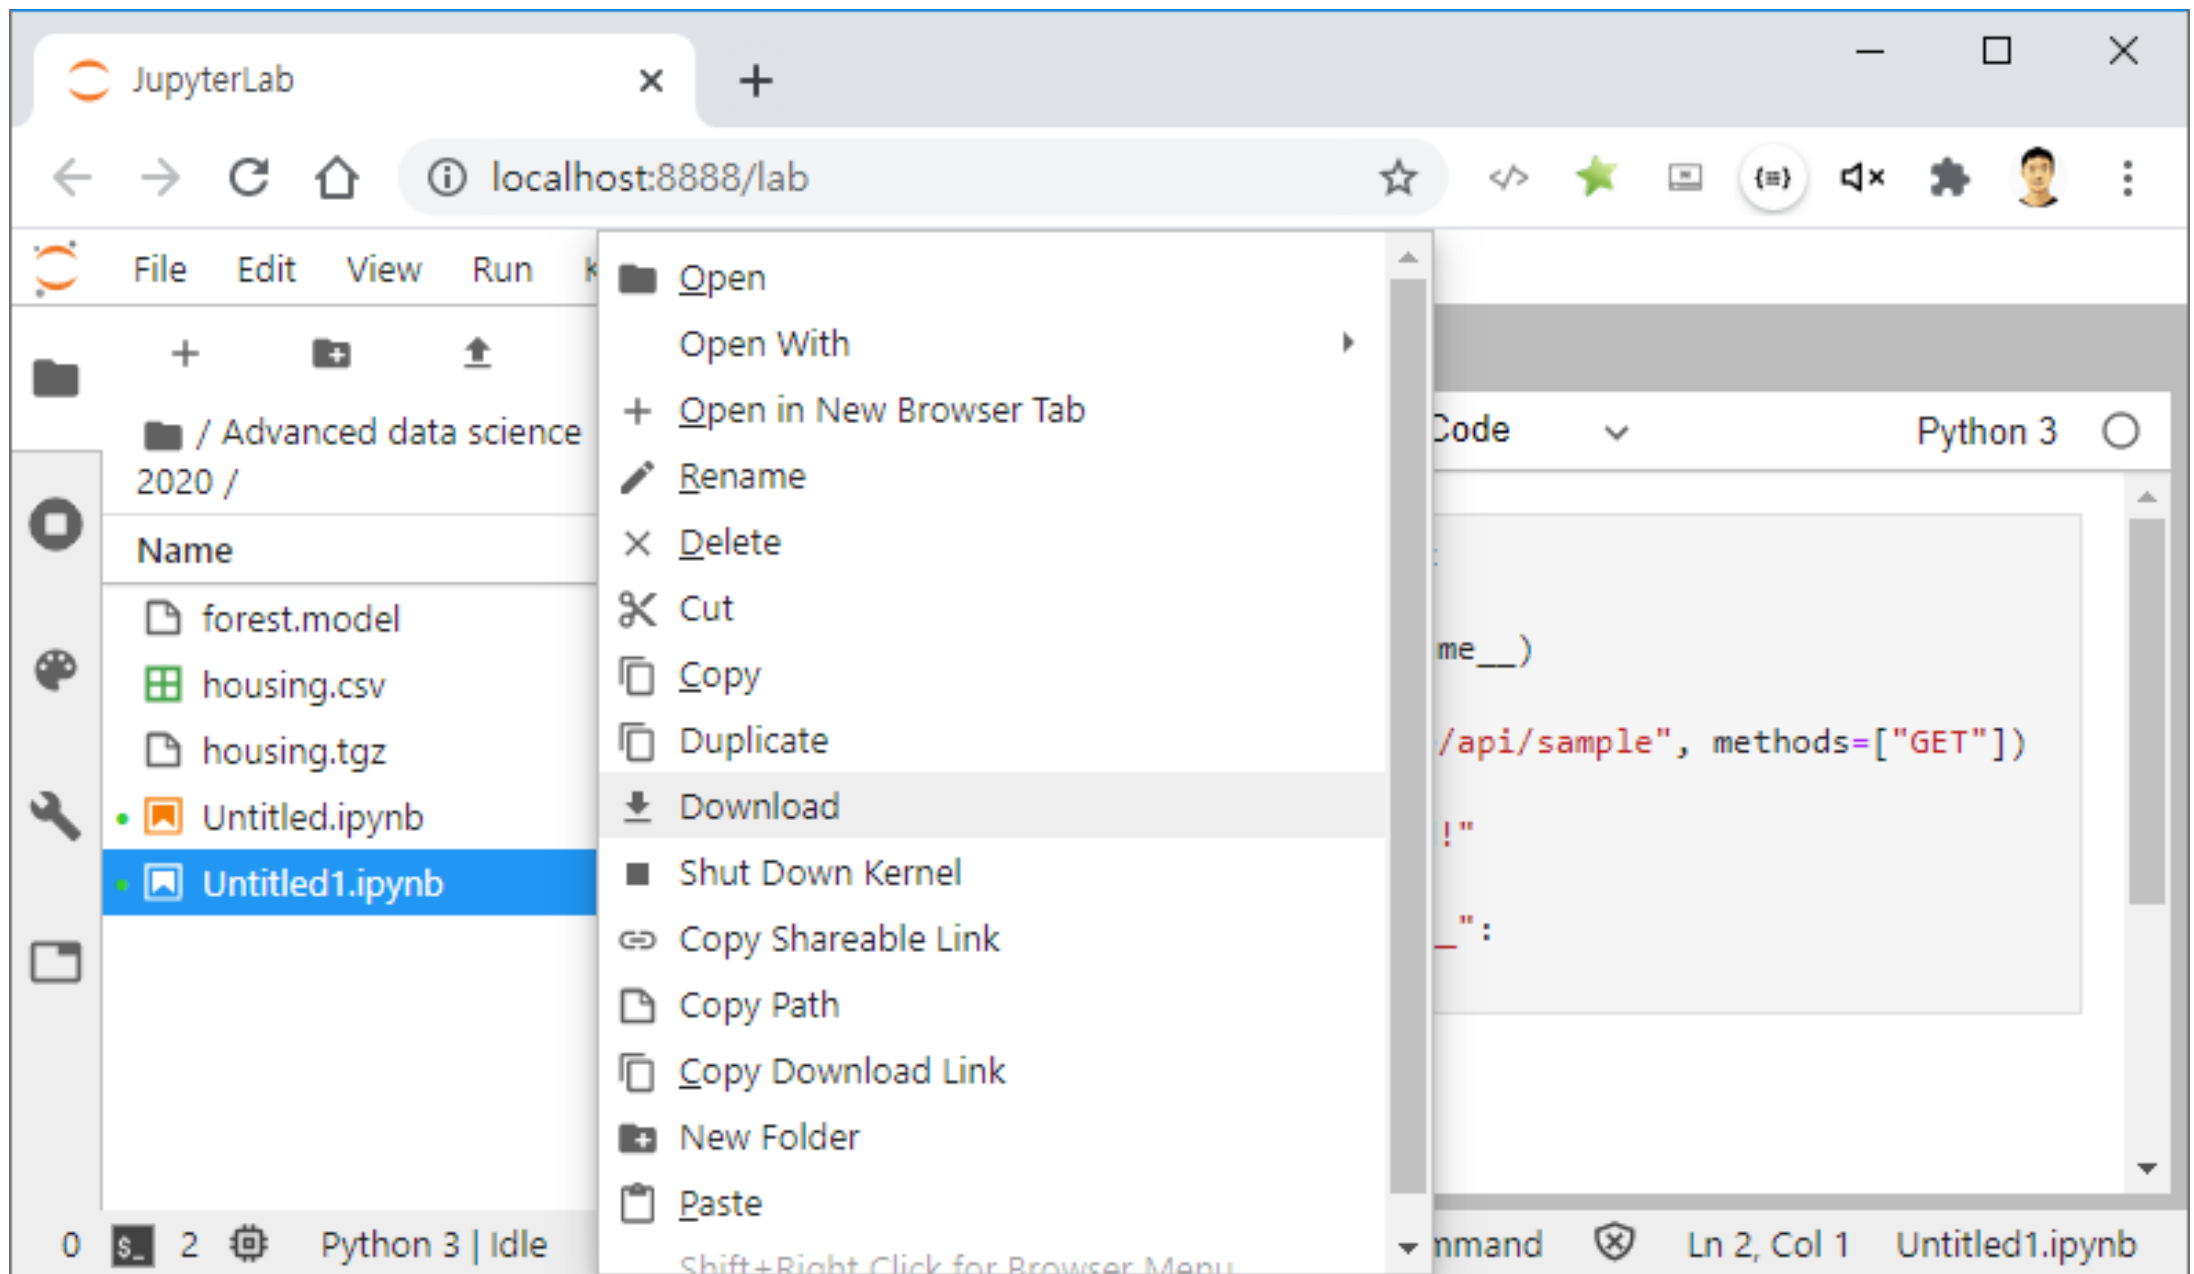
sample_api = Flask(__name__)

@sample_api.route("/gse/api/sample", methods=["GET"])
def hello_world():
    return "Hello World!"

if __name__ == "__main__":
    sample_api.run()
```

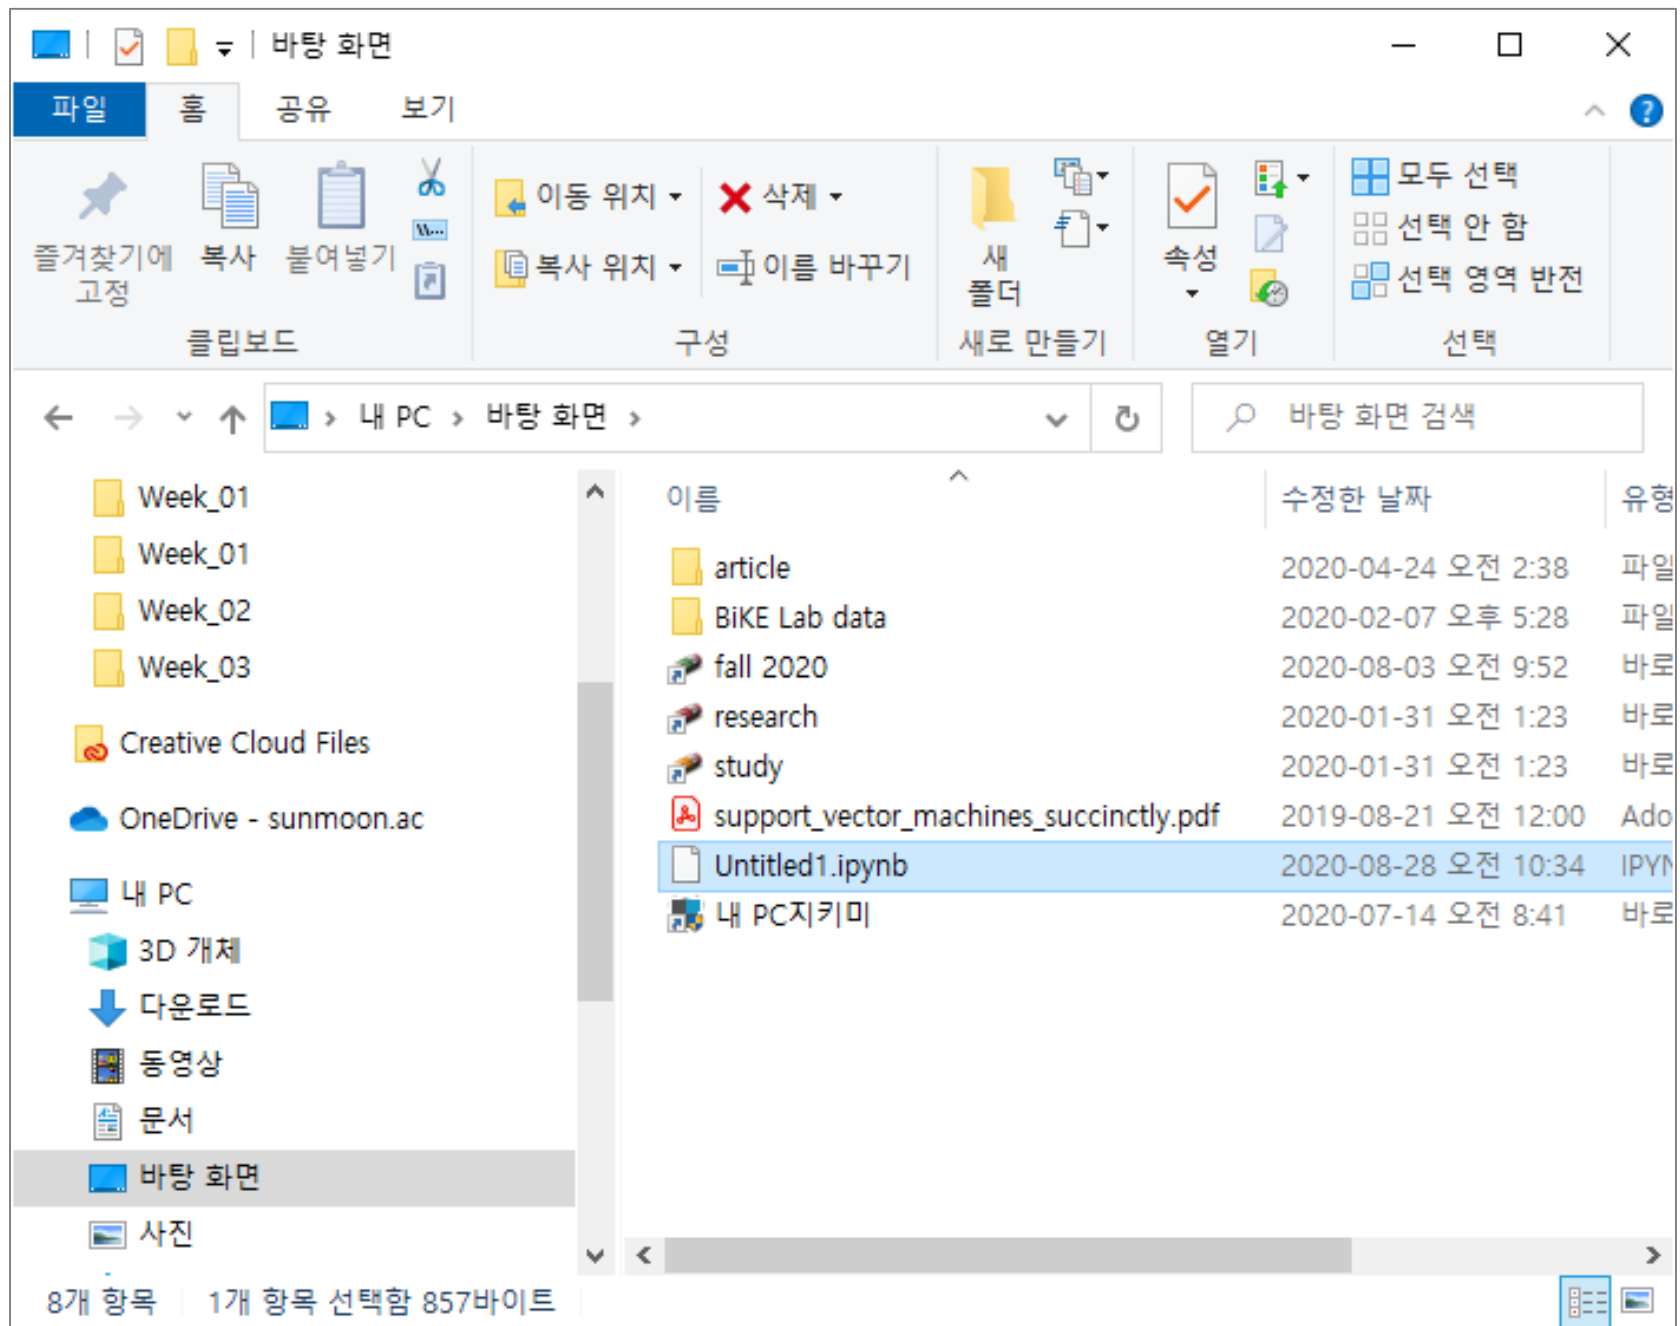
The status bar at the bottom indicates the current cell is at `Ln 2, Col 1` in `Untitled1.ipynb`, and the kernel is `Python 3 | Idle`.

## Download > Move to "Desktop"


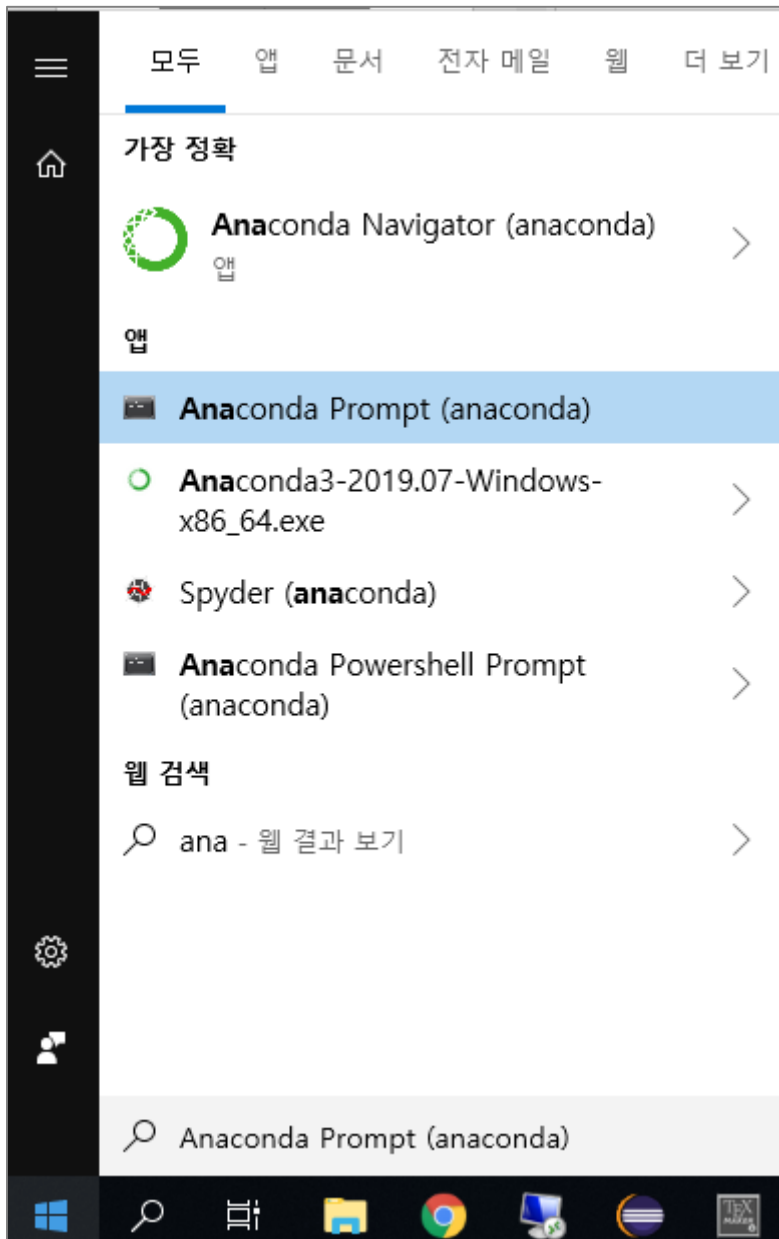




## Download > Move to "Desktop"



## Anaconda prompt 실행 및 바탕화면으로 이동



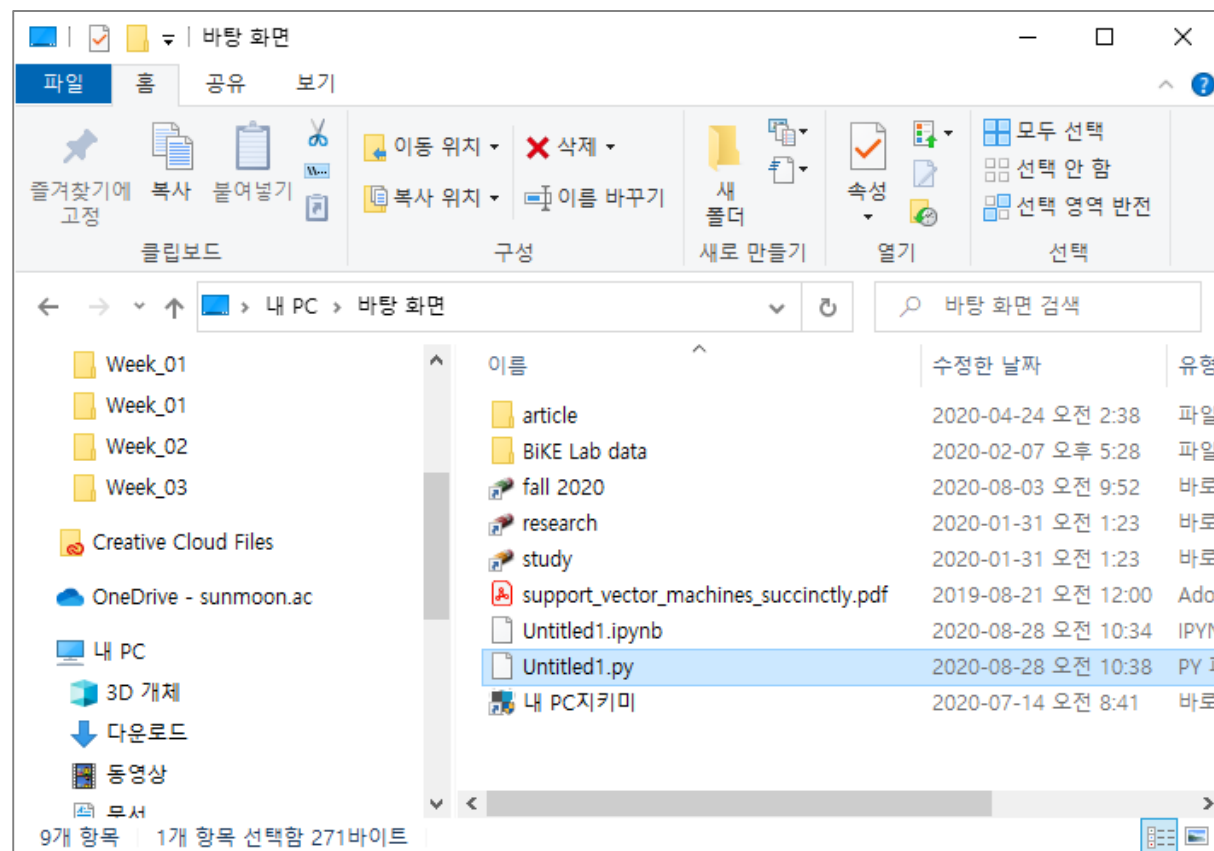
# Jupyter (.ipynb) to Python (.py)

```
Anaconda Prompt (anaconda3)

(base) C:\Users\Eung-HeeKim>cd Desktop

(base) C:\Users\Eung-HeeKim\Desktop>jupyter nbconvert --to script Untitled1.ipynb
[NbConvertApp] Converting notebook Untitled1.ipynb to script
[NbConvertApp] Writing 254 bytes to Untitled1.py

(base) C:\Users\Eung-HeeKim\Desktop>
```

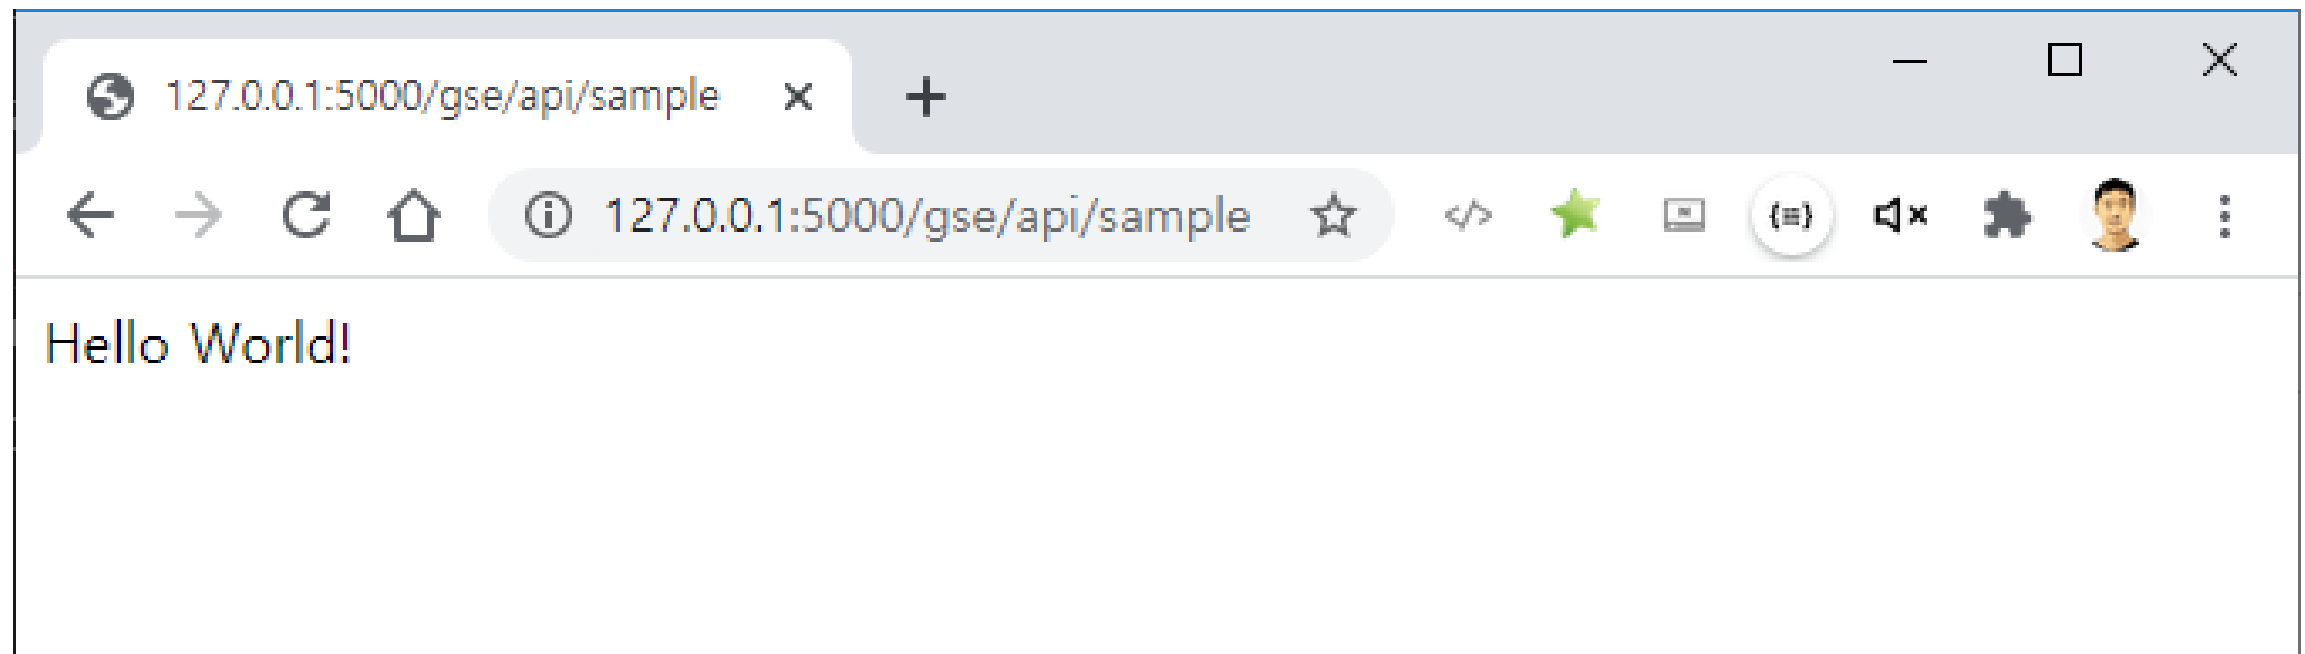


## Untitled1.py 실행


```
Anaconda Prompt (anaconda3) - python Untitled1.py

(base) C:\Users\Eung-HeeKim>cd Desktop
(base) C:\Users\Eung-HeeKim\Desktop>jupyter nbconvert --to script Untitled1.ipynb
[NbConvertApp] Converting notebook Untitled1.ipynb to script
[NbConvertApp] Writing 254 bytes to Untitled1.py
(base) C:\Users\Eung-HeeKim\Desktop>python Untitled1.py
* Serving Flask app "Untitled1" (lazy loading)
* Environment: production
  WARNING: This is a development server. Do not use it in a production deployment.
  Use a production WSGI server instead.
* Debug mode: off
* Running on http://127.0.0.1:5000/ (Press CTRL+C to quit)
```

## REST API 호출



## Untitled1.py 수정



```
*Untitled1.py - Windows 메모장
파일(F)  편집(E)  서식(O)  보기(V)  도움말(H)

#!/usr/bin/env python
# coding: utf-8

# In[ ]:

from flask import Flask

sample_api = Flask(__name__)

@sample_api.route("/gse/api/sample", methods=["GET"])
def hello_world():
    answer = "<html> <body> <h1>Hi</h1> </body> </html>"
    return answer

if __name__ == "__main__":
    sample_api.run()
```

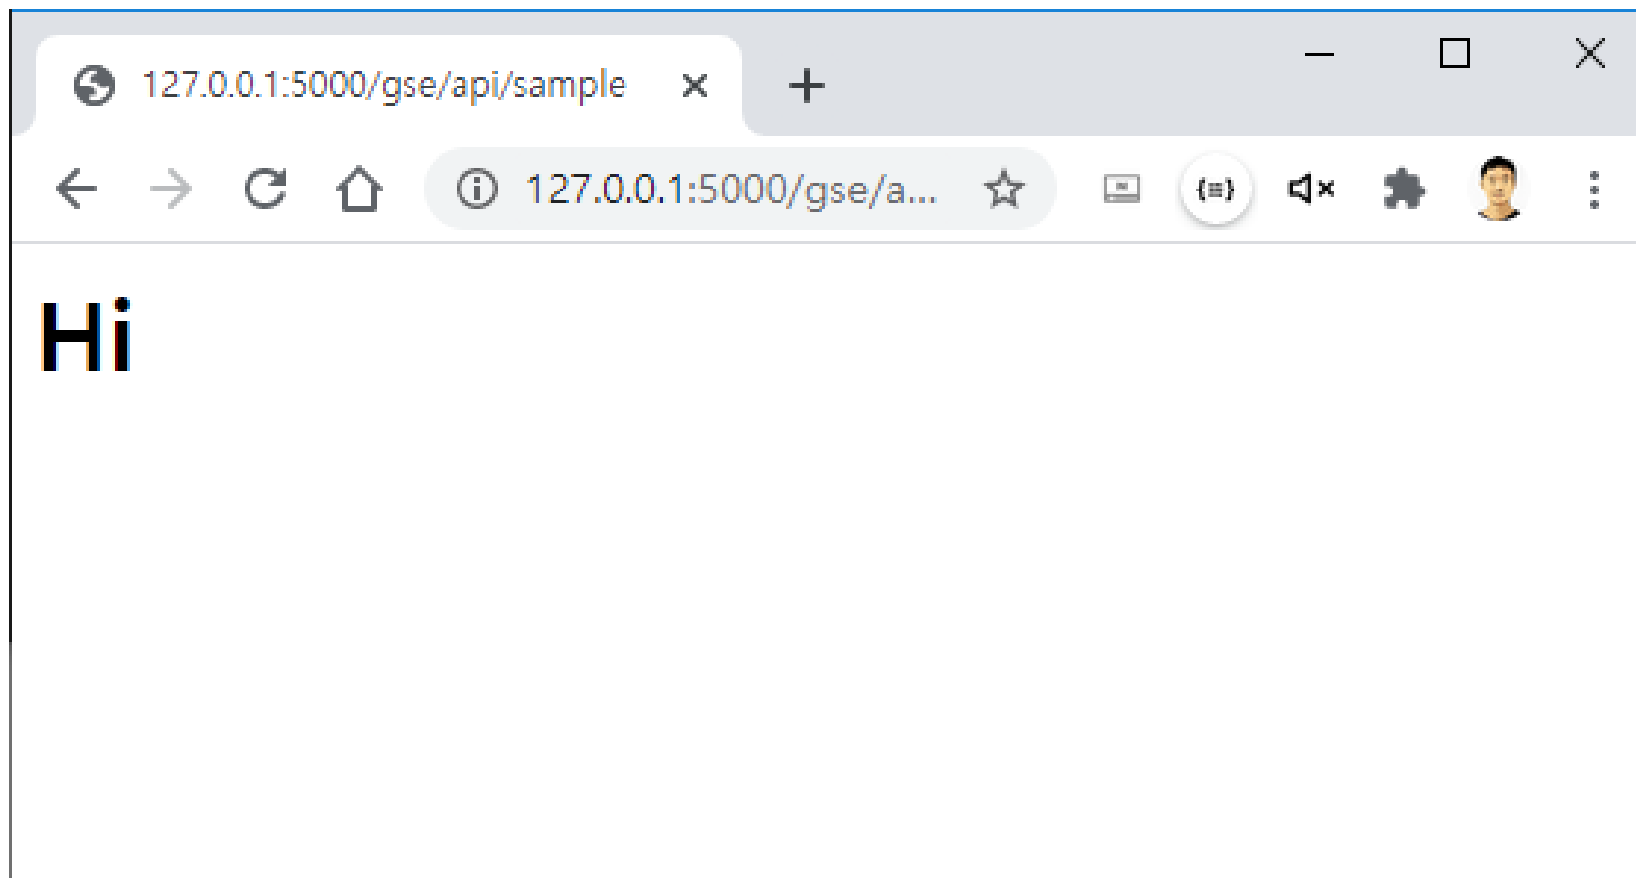
Ln 7, Col 24    100%    Windows (CRLF)    UTF-8

Untitled1.py 재실행  
CTRL+C → python Untitled1.py

```
Anaconda Prompt (anaconda3) - python Untitled1.py
[NbConvertApp] Writing 254 bytes to Untitled1.py
(base) C:\Users\Eung-HeeKim\Desktop>python Untitled1.py
* Serving Flask app "Untitled1" (lazy loading)
* Environment: production
  WARNING: This is a development server. Do not use it in a production deployment.
  Use a production WSGI server instead.
* Debug mode: off
* Running on http://127.0.0.1:5000/ (Press CTRL+C to quit)

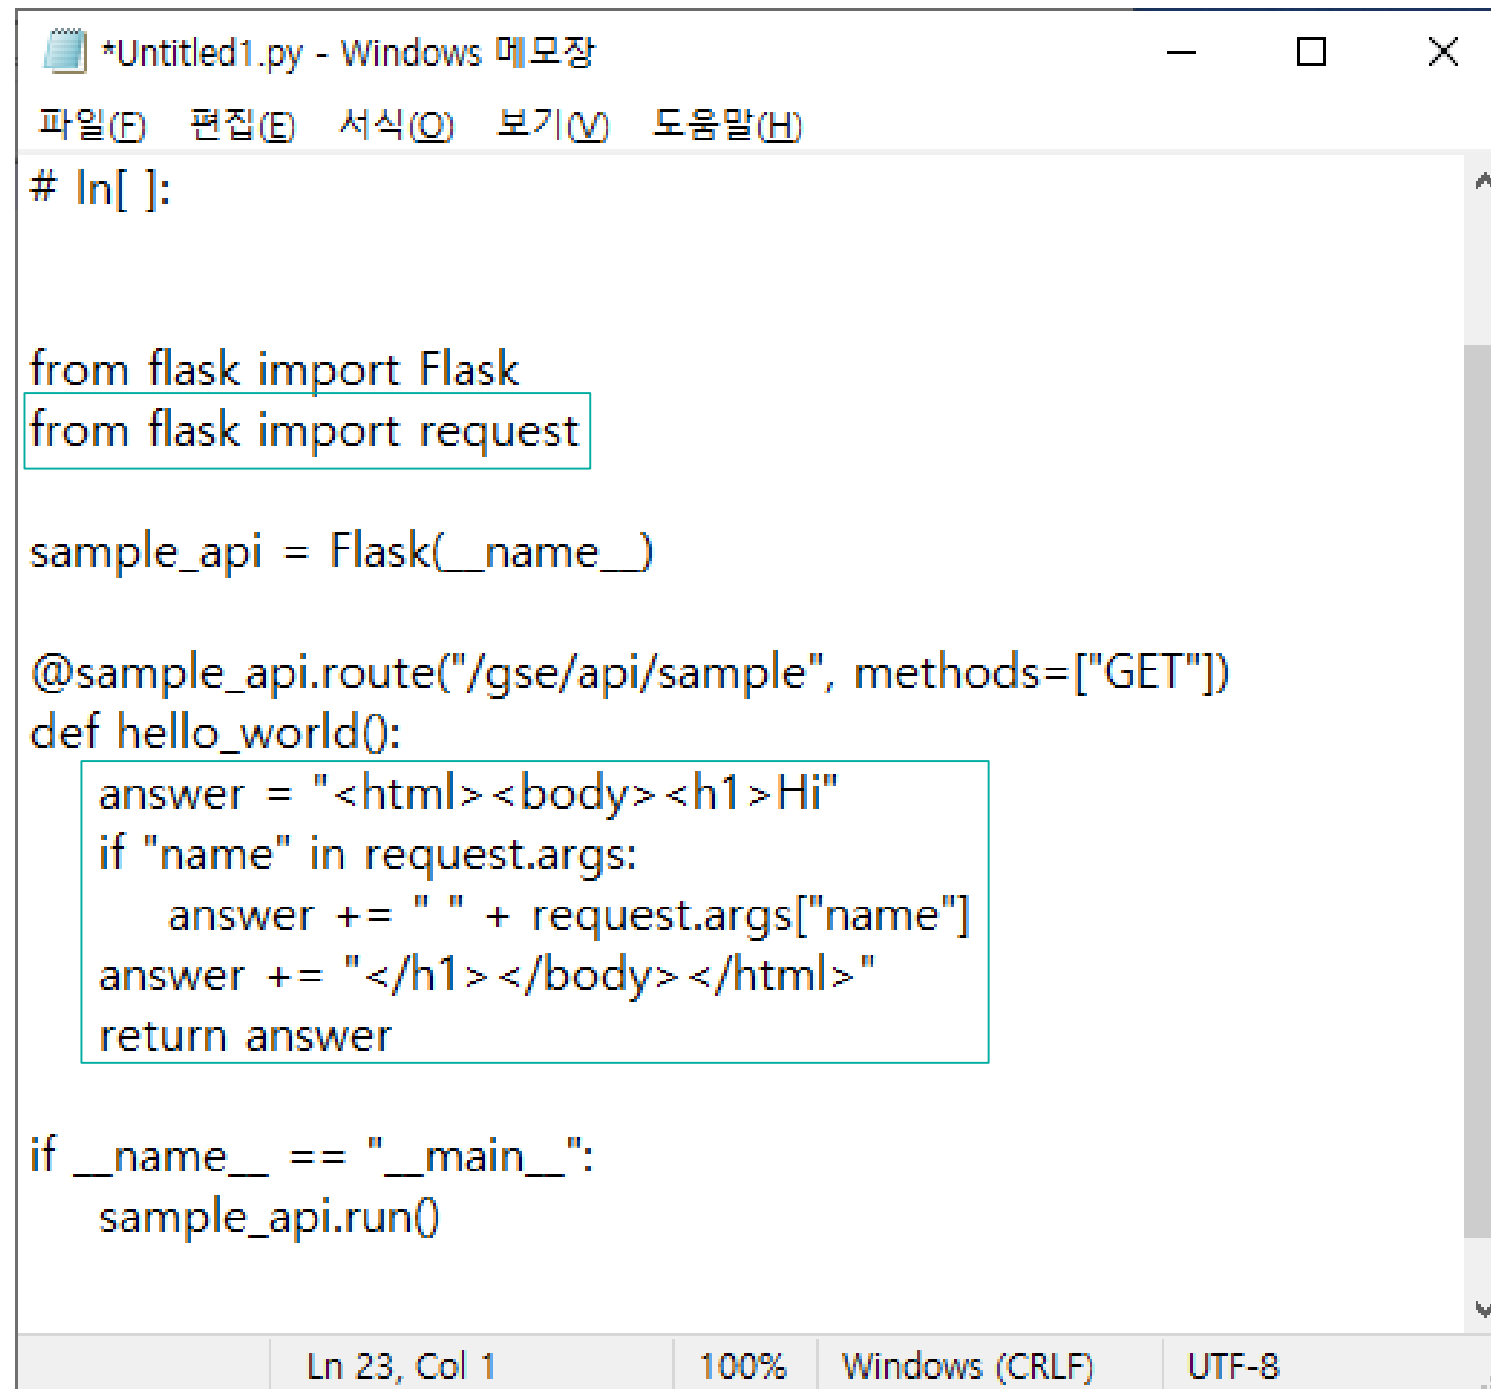
(base) C:\Users\Eung-HeeKim\Desktop>python Untitled1.py
* Serving Flask app "Untitled1" (lazy loading)
* Environment: production
  WARNING: This is a development server. Do not use it in a production deployment.
  Use a production WSGI server instead.
* Debug mode: off
* Running on http://127.0.0.1:5000/ (Press CTRL+C to quit)
```

## REST API 재호출





## 매개 변수(Parameter) 처리하기



```
*Untitled1.py - Windows 메모장
파일(E)  편집(E)  서식(O)  보기(V)  도움말(H)

# In[ ]:

from flask import Flask
from flask import request

sample_api = Flask(__name__)

@sample_api.route("/gse/api/sample", methods=["GET"])
def hello_world():
    answer = "<html><body><h1>Hi"
    if "name" in request.args:
        answer += " " + request.args["name"]
    answer += "</h1></body></html>"
    return answer

if __name__ == "__main__":
    sample_api.run()
```

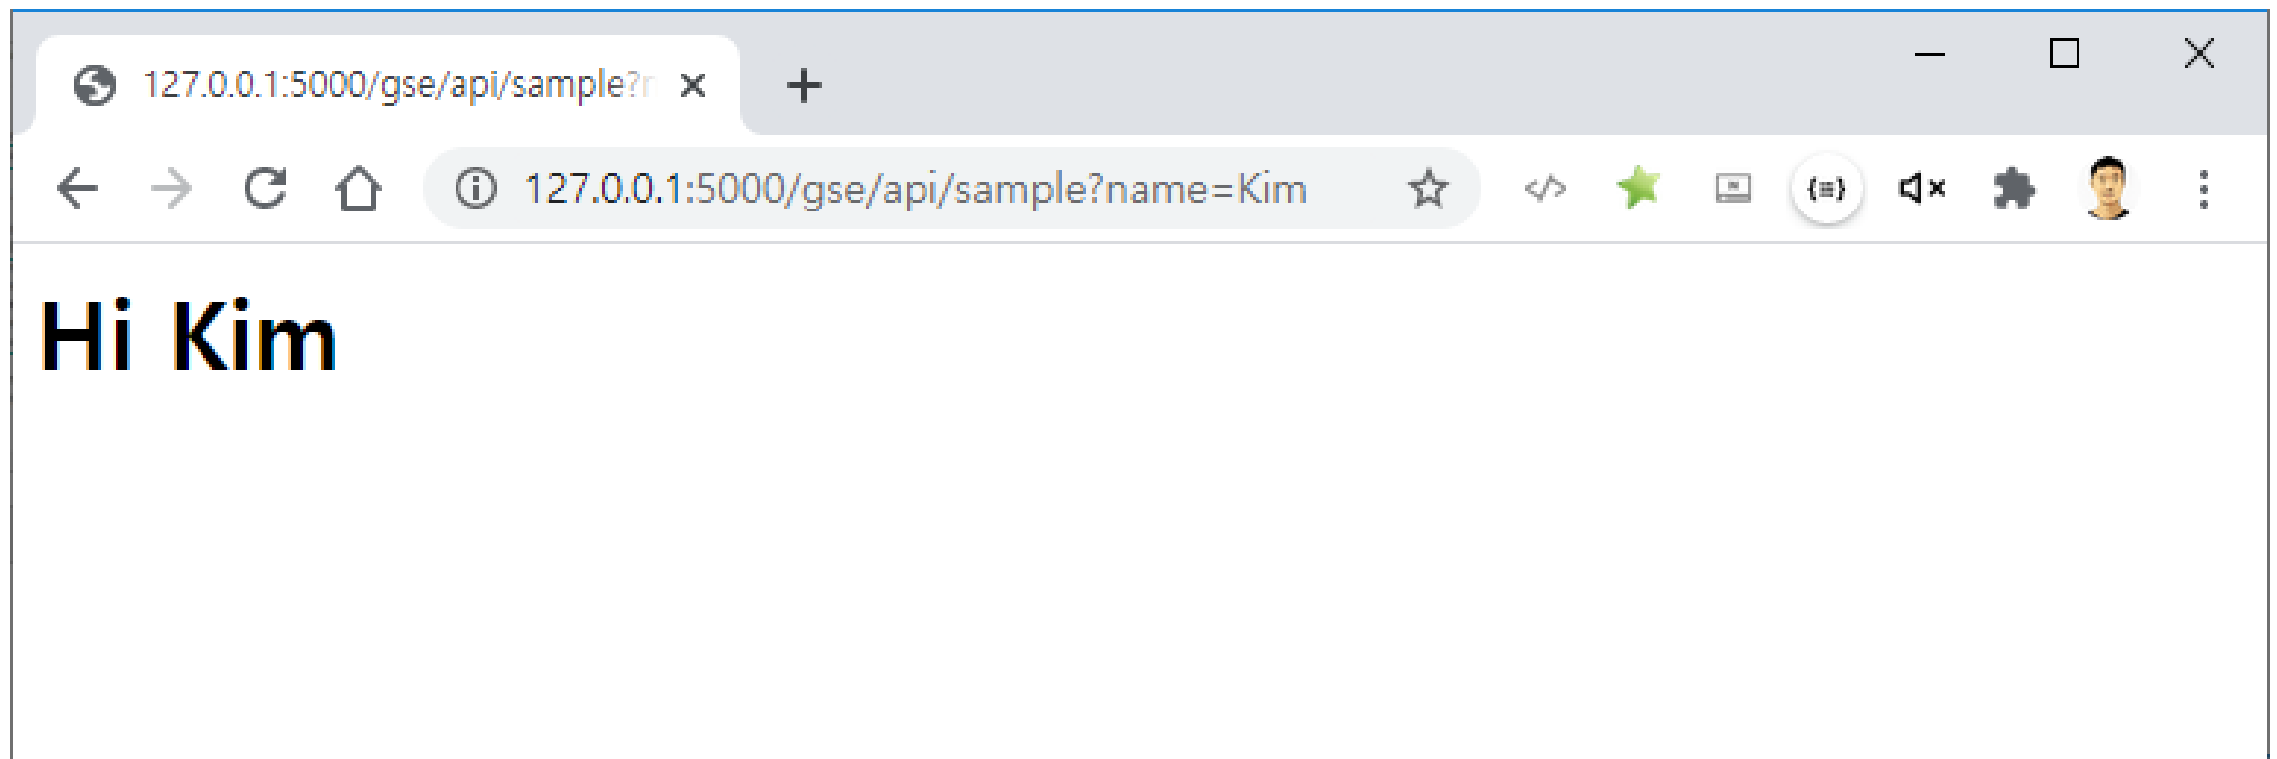
Ln 23, Col 1    100%    Windows (CRLF)    UTF-8

## Untitled.py 재실행

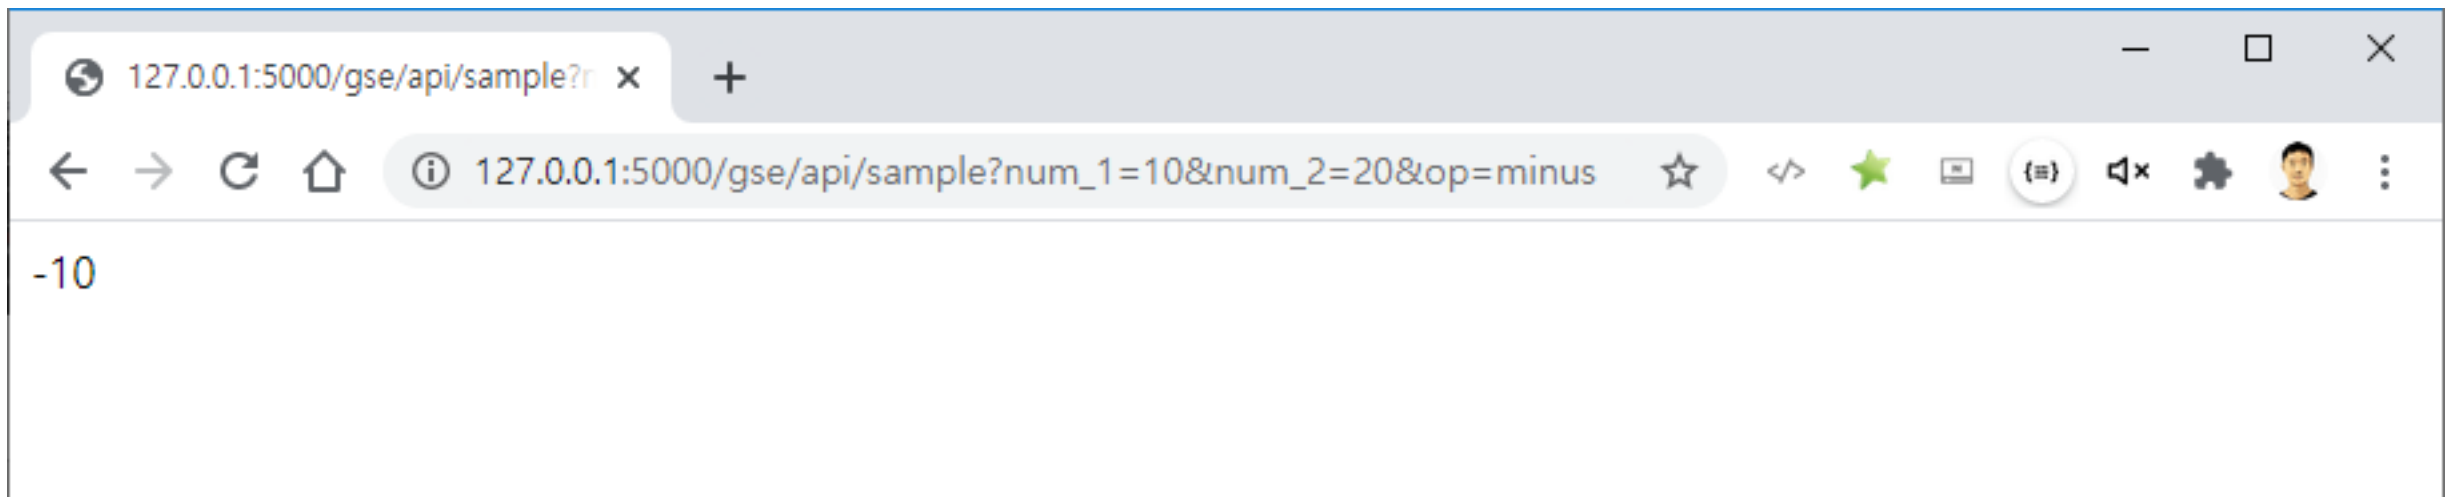

```
Anaconda Prompt (anaconda3) - python Untitled1.py
[NbConvertApp] Writing 254 bytes to Untitled1.py
(base) C:\Users\Eung-HeeKim\Desktop>python Untitled1.py
* Serving Flask app "Untitled1" (lazy loading)
* Environment: production
  WARNING: This is a development server. Do not use it in a production deployment.
  Use a production WSGI server instead.
* Debug mode: off
* Running on http://127.0.0.1:5000/ (Press CTRL+C to quit)

(base) C:\Users\Eung-HeeKim\Desktop>python Untitled1.py
* Serving Flask app "Untitled1" (lazy loading)
* Environment: production
  WARNING: This is a development server. Do not use it in a production deployment.
  Use a production WSGI server instead.
* Debug mode: off
* Running on http://127.0.0.1:5000/ (Press CTRL+C to quit)
```

## REST API 재호출



## 덧셈 & 뺄셈 REST API 만들기



## 덧셈 & 뺄셈 REST API 만들기

```
@sample_api.route('/gse/api/sample/', methods=['GET'])
def hello_world( ):
    result = 0
    num_1 = int(request.args["num_1"])
    num_2 = int(request.args["num_2"])
    op = request.args["op"]
    if op == "plus":
        result = num_1 + num_2
    elif op == "minus":
        result = num_1 - num_2
    return str(result)
```

# REST API 호출 in Python

The screenshot displays a JupyterLab environment. The top browser window shows the URL `localhost:8888/lab`. The left sidebar contains a file explorer for the directory `/ Advanced data science 2020 /`, listing files like `forest.model`, `housing.csv`, `housing.tgz`, and several `Untitled.ipynb` files. The main workspace shows a notebook titled `Untitled2.ipynb` in `Code` mode, using `Python 3`. The code in the notebook is as follows:

```
[1]: import requests

url = "http://127.0.0.1:5000/gse/api/sample"

params = {"num_1": 10, "num_2": 20, "op": "minus"}

response = requests.get(url, params)

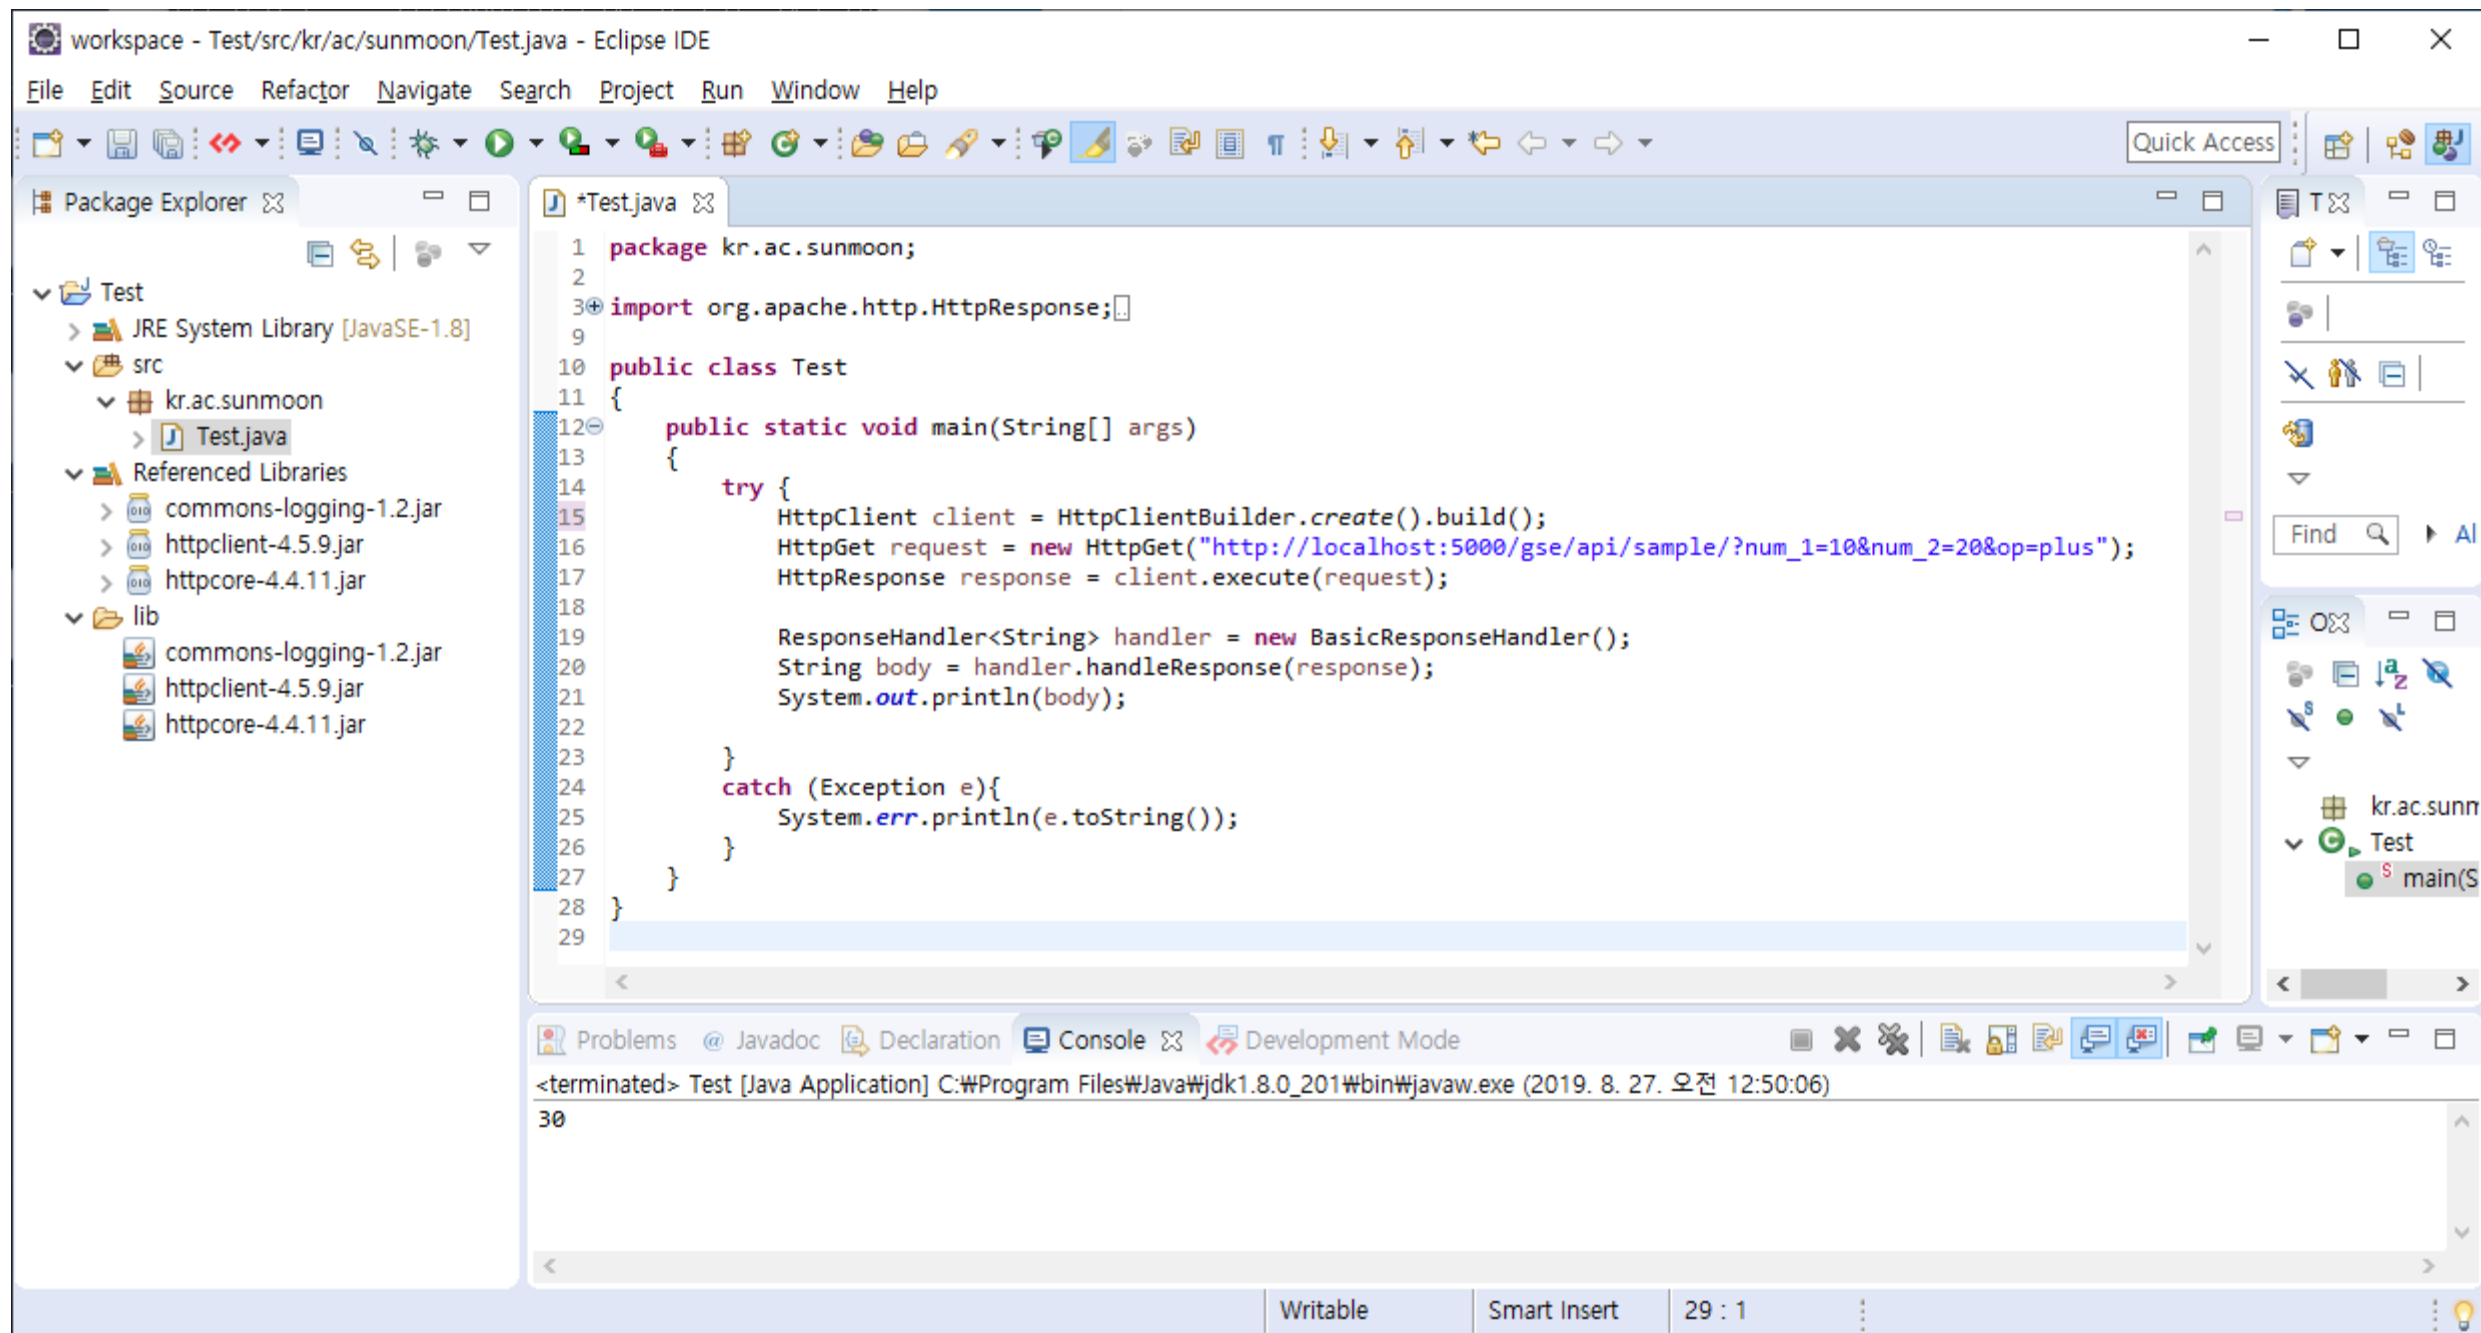
print(response.text)
```

The output of the code is `-10`. The bottom status bar indicates the current mode is `Edit`, the cursor is at `Ln 1, Col 1`, and the active file is `Untitled2.ipynb`.

## REST API 호출 in Java

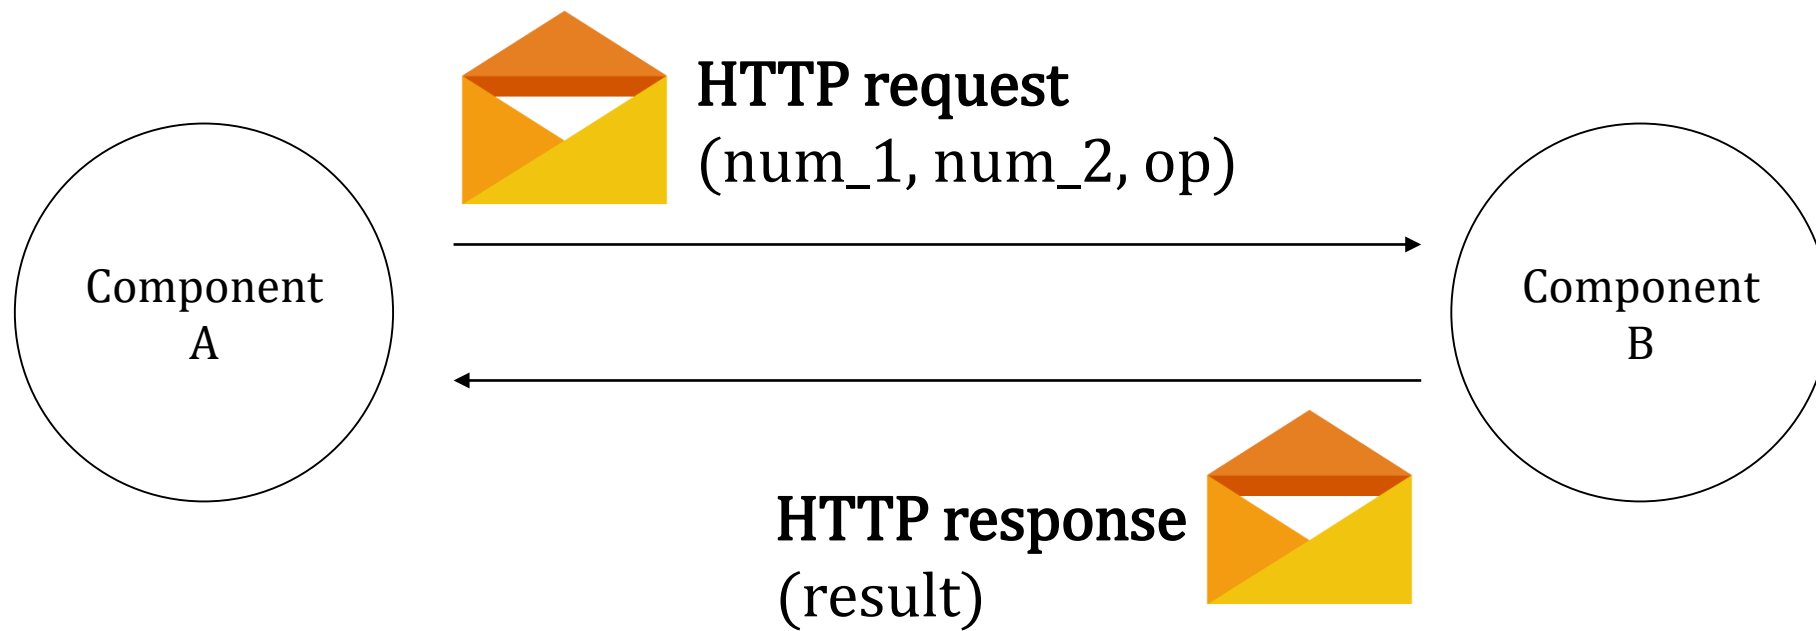
- Eclipse 실행
- Java project 생성
- kr.ac.sunmoon 패키지 생성
- lib 폴더 생성
- e-강의동 > 데이터사이언스응용 > 02주차 강의자료
  - 3개의 jar 파일 다운로드 및 lib 폴더에 복사 → build path에 추가
  - Test.java 파일 다운로드 및 kr.ac.sunmoon 패키지에 복사 → 실행

# REST API 호출 in Java

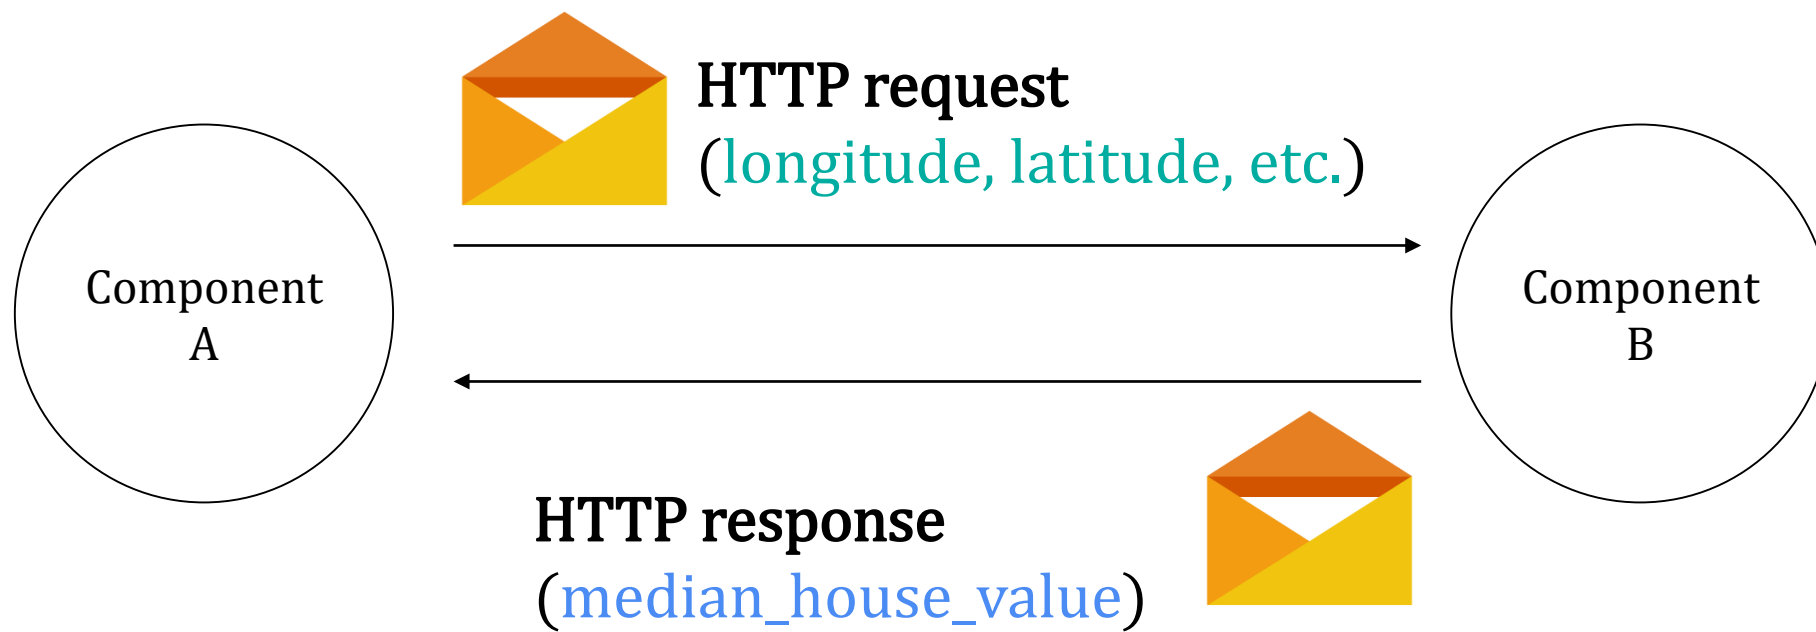




What if..



What if..



Thank you